

IEMS5709 Fall 2016

Advanced Topics in Information Processing:

Big Data Processing Systems and

Information Processing

Spark Streaming, SparkSQL, MLlib,

GraphX and Beyond

Prof. Wing C. Lau

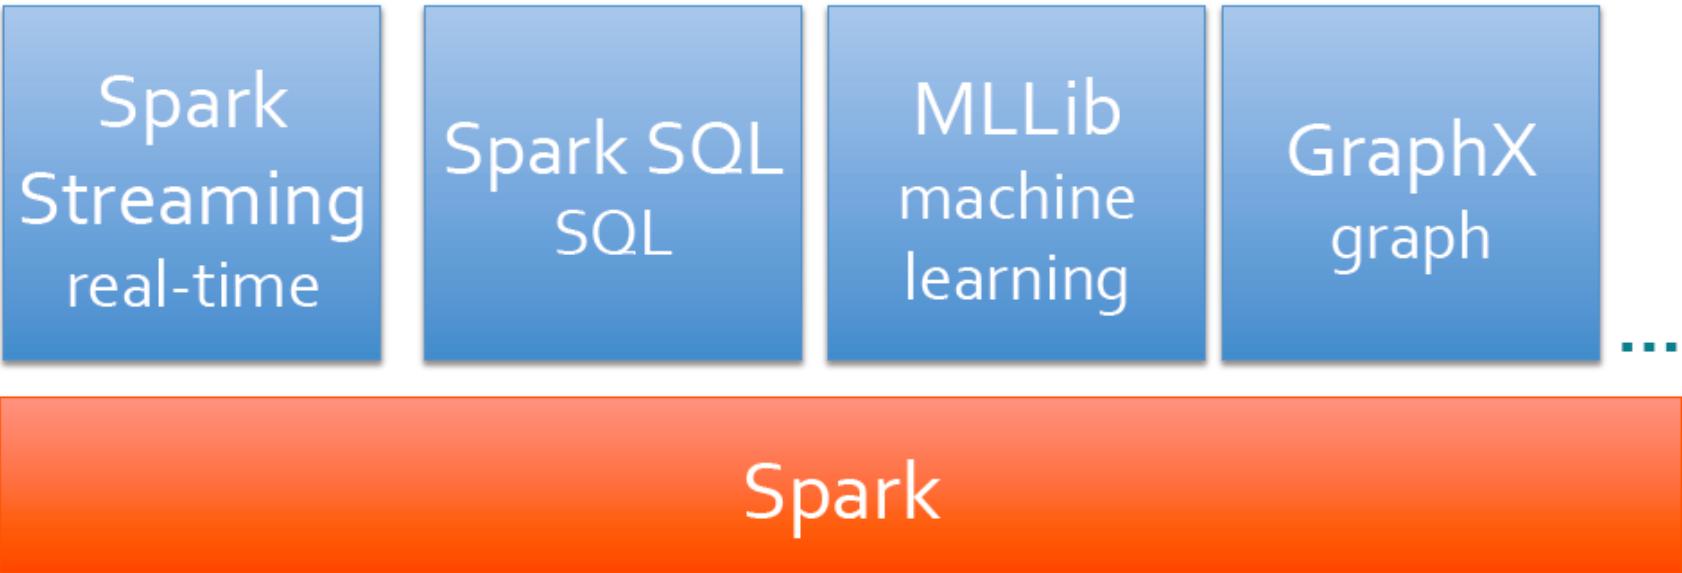
Department of Information Engineering

wclau@ie.cuhk.edu.hk

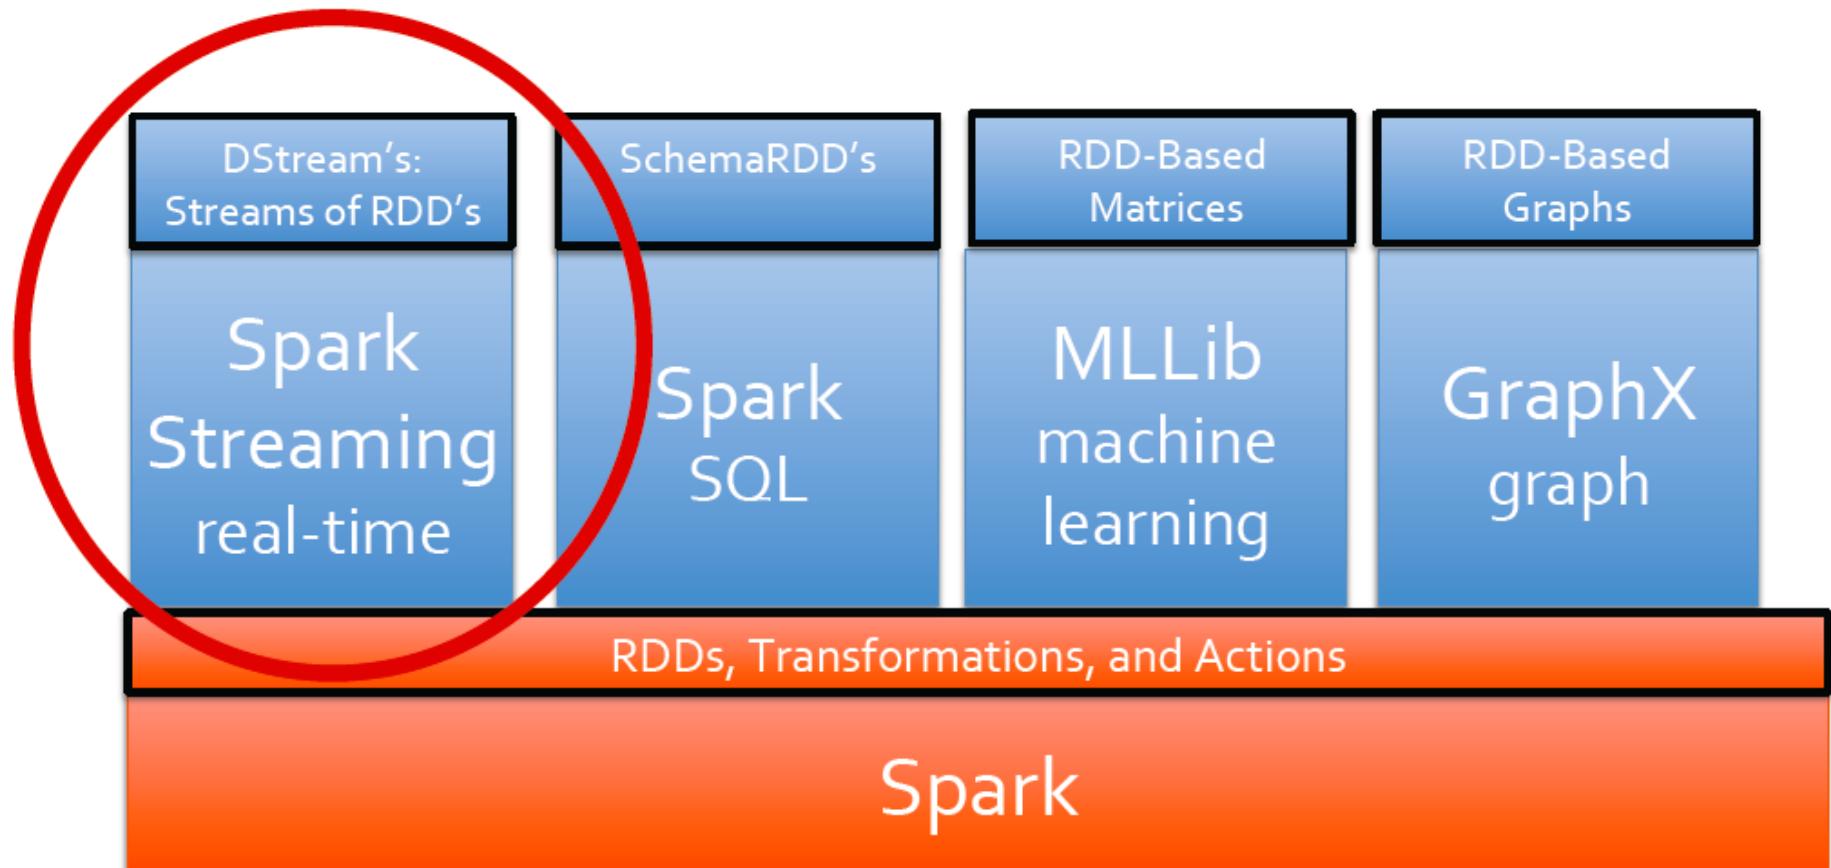
Acknowledgements

- Slides in this chapter are adapted from the following sources:
 - Matei Zaharia, “Spark 2.0,” Spark Summit East Keynote, Feb 2016.
 - Reynold Xin, “The Future of Real-Time in Spark,” Spark Summit East Keynote, Feb 2016.
 - Michael Armbrust, “Structuring Spark: SQL, DataFrames, DataSets, and Streaming,” Spark Summit East Keynote, Feb 2016.
 - Ankur Dave, “GraphFrames: Graph Queries in Spark SQL,” Spark Summit East, Feb 2016.
 - Michael Armbrust, “Spark DataFrames: Simple and Fast Analytics on Structured Data,” Spark Summit Amsterdam, Oct 2015.
 - Michael Armbrust et al, “Spark SQL: Relational Data Processing in Spark,” ACM SIGMOD 2015 Talk.
 - Michael Armbrust, “Spark SQL Deep Dive,” Melbourne Spark Meetup, June 2015.
 - Reynold Xin, “Spark,” Stanford CS347 Guest Lecture, May 2015.
 - Joseph K. Bradley, “Spark DataFrames and ML Pipelines,” MLconf Seattle, May 2015.
 - Ameet Talwalkar, “MLlib: Spark’s Machine Learning Library,” AMPCamps 5 Talk, Nov. 2014.
 - Shivaram Venkataraman, Zongheng Yang, “SparkR: Enabling Interactive Data Science at Scale,” AMPCamps 5 Talk, Nov. 2014.
 - Tathagata Das, “Spark Streaming: Large-scale near-real-time stream processing,” O’Reilly Strata Conference talk, 2013.
 - Joseph Gonzalez et al, “GraphX: Graph Analytics on Spark,” talk at AMPCAMP 3, 2013.
- All copyrights belong to the original authors of the materials.

Generality of RDDs in Spark



Generality of RDDs in Spark



Motivation for Spark Streaming

- Many Important Applications must process Large Data Streams at second-scale latencies
 - Site Statistics, Intrusion Detection, Online ML, Fraud Detection
- To build and scale these applications require:
 - Integration: with Offline Analytic Stack
 - Fault-tolerance: to handle Crashes and Stragglers
 - Efficiency: low cost beyond base processing
 - Work with distributed collections as you would with local ones

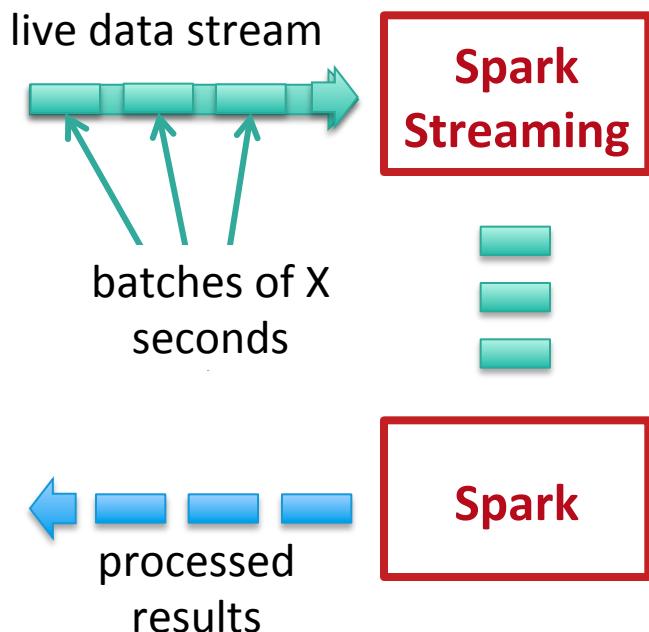
Spark Streaming Overview

- Low Latency, High-throughput and Fault Tolerant
- Discretized Stream (DStream): Micro-batches of RDDs
 - Operations are similar to RDD
 - Lineage for Fault-Tolerance
- Leverage Core Components from Spark
 - RDD data model and API
 - Data Partitioning and Shuffles
 - Task Scheduling
 - Monitoring/ Instrumentation
 - Scheduling and Resource Allocation
- Support Flume, Kafka, Twitter, Kinesis, etc for Data Ingestion
- Long-running Spark Applications

Discretized Stream Processing

Run a streaming computation as a **series of very small, deterministic batch jobs**

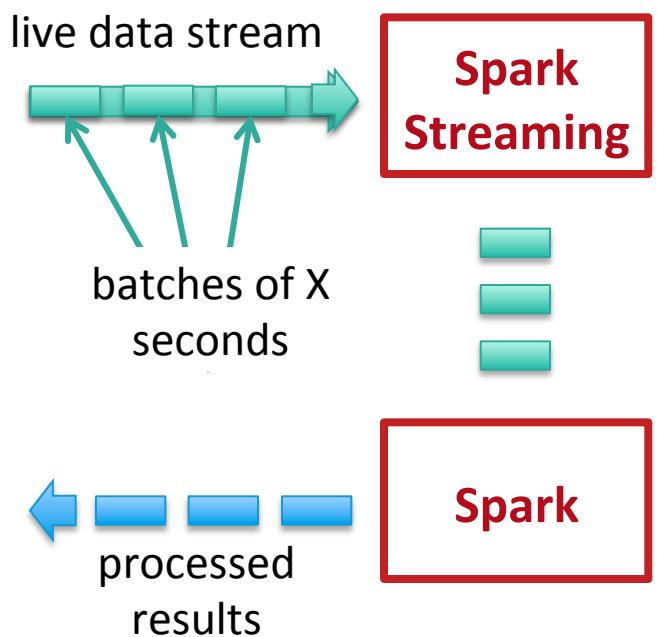
- Chop up the live stream into batches of X seconds
- Spark treats each batch of data as RDDs and processes them using RDD operations
- Finally, the processed results of the RDD operations are returned in batches



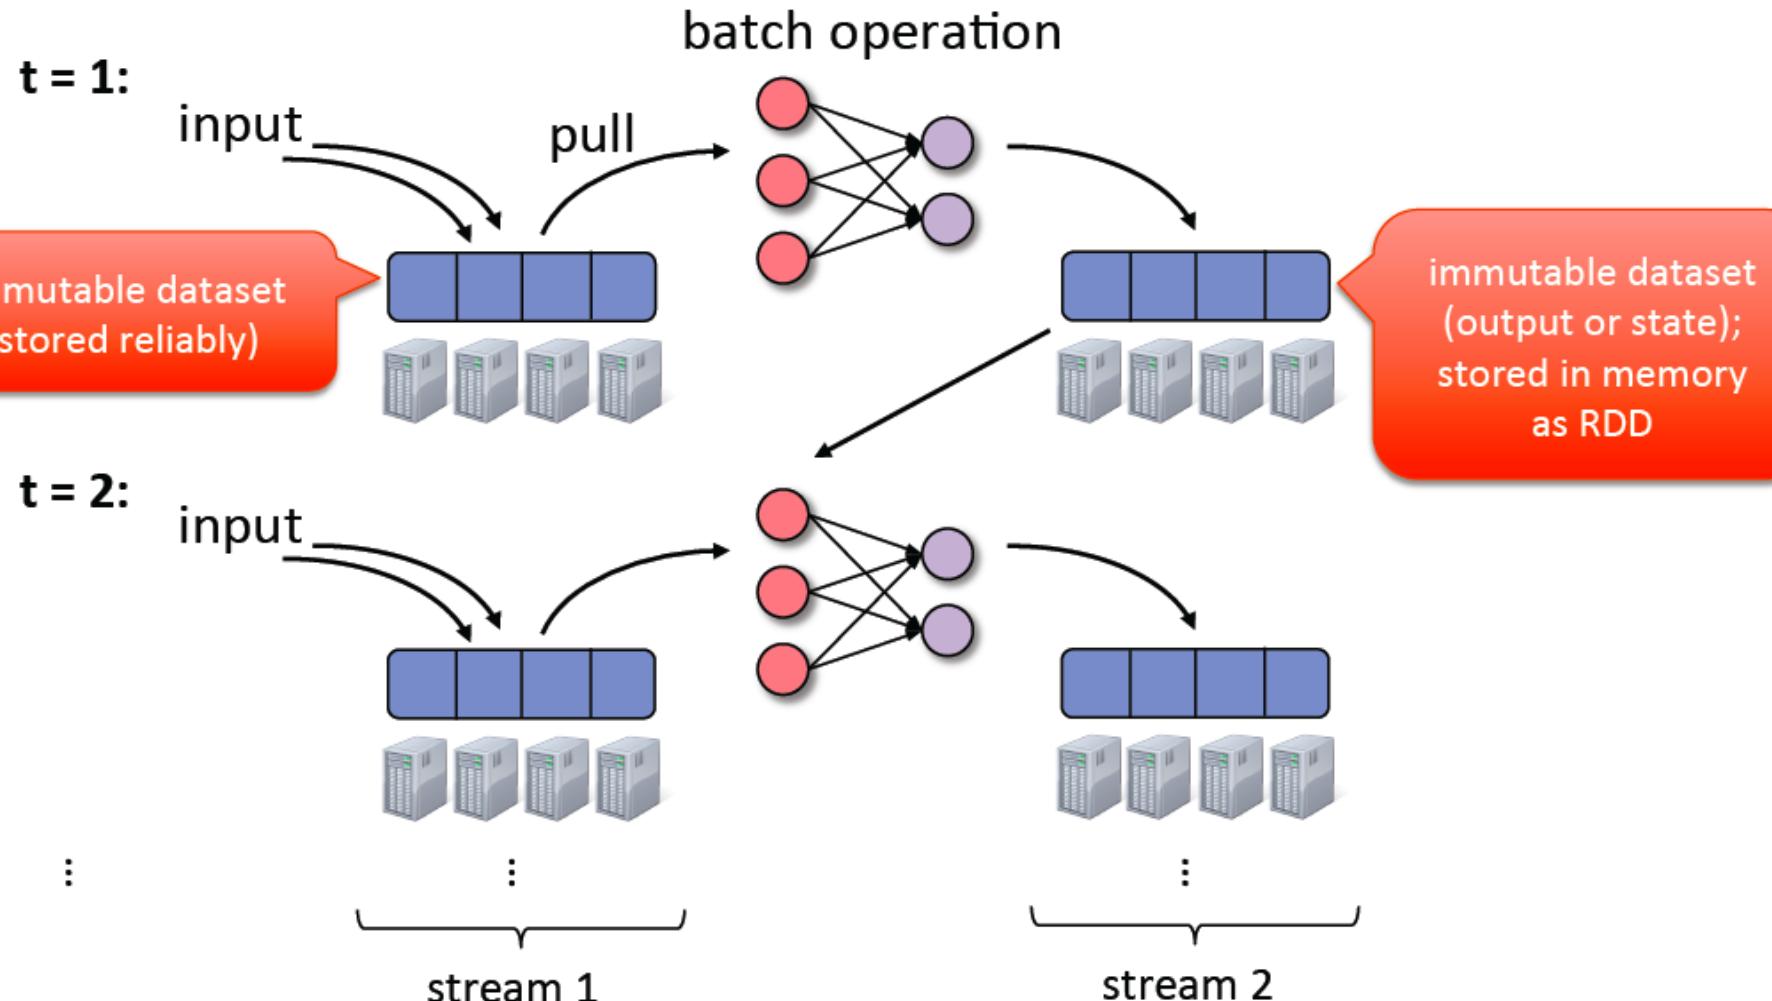
Discretized Stream Processing

Run a streaming computation as a **series of very small, deterministic batch jobs**

- Batch sizes as low as $\frac{1}{2}$ second, latency ~ 1 second
- Potential for combining batch processing and streaming processing in the same system



Discretized Stream Processing (Micro-Batching)



Programming Interface

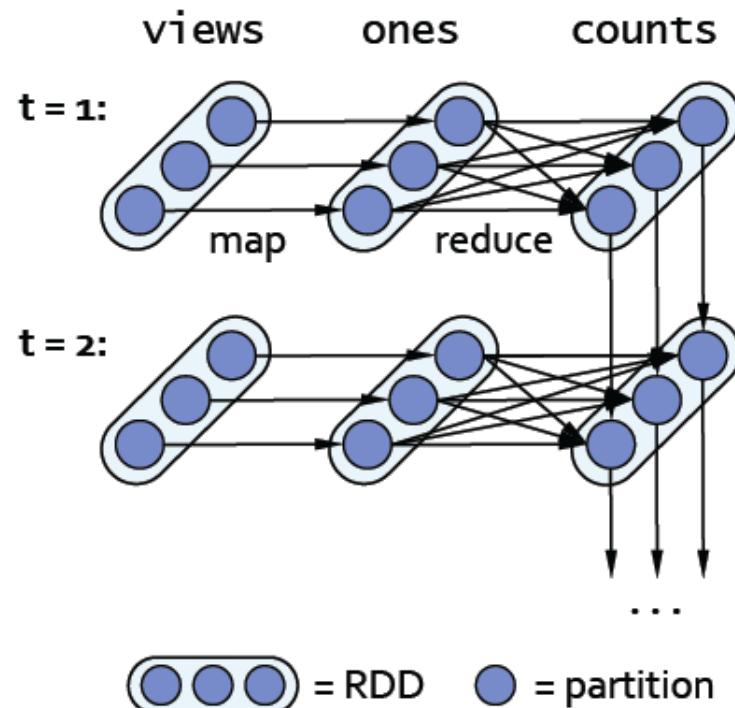
Simple functional API

```
views = readstream("http://...", "1s")
ones = views.map(ev => (ev.url, 1))
counts = ones.runningReduce(_ + _)
```

Interoperates with RDDs

```
// Join stream with static RDD
counts.join(historicCounts).map(...)

// Ad-hoc queries on stream state
counts.slice("21:00", "21:05").topK(10)
```



Spark Streaming API

- **Transformations** – modify data from one DStream to another
 - Standard RDD operations – map, filter, distinct, countByValue, reduceByKey, join, ...
 - Stateful, Sliding Window-based Operations
 - window, updateStateByKey, countByValueAndWindow
 - Window Size & Slide Interval
- **Output Operations** – send data to external entity
 - saveAsHadoopFiles – saves to HDFS
 - foreach – do anything with each batch of results
- Checkpointing
- Register DStream as a SQL table

Example 1: Get HashTags from Twitter

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)
```

DStream: a sequence of distributed datasets (RDDs)
representing a distributed stream of data



stored in memory as an RDD
(immutable, distributed dataset)

Example 1: Get HashTags from Twitter

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)
val hashTags = tweets.flatMap (status => getTags(status))
```

new DStream

transformation: modify data in one DStream to create another DStream

tweets DStream

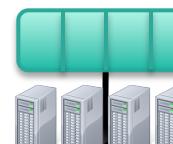
batch @ t

batch @ t+1

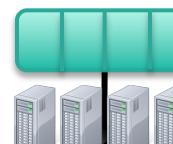
batch @ t+2



hashTags Dstream
[#cat, #dog, ...]



flatMap



flatMap



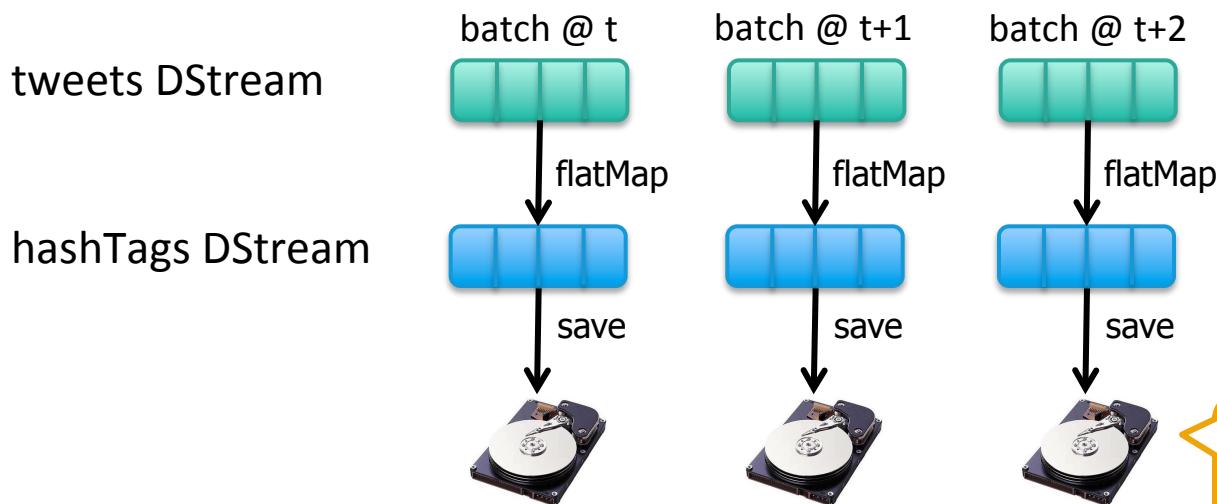
flatMap

new RDDs created
for every batch

Example 1: Get HashTags from Twitter

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)
val hashTags = tweets.flatMap (status => getTags(status))
hashTags.saveAsHadoopFiles("hdfs://...")
```

output operation: to push data to external storage

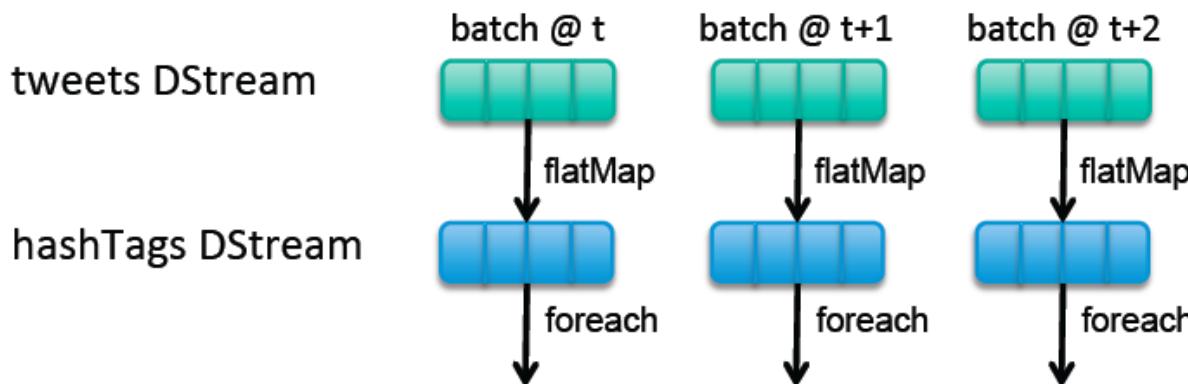


every batch saved to HDFS

Example 1: Get HashTags from Twitter

```
val tweets = TwitterUtils.createStream(ssc, None)
val hashTags = tweets.flatMap(status => getTags(status))
hashTags.foreachRDD(hashTagRDD => { ... })
```

foreach: do whatever you want with the processed data



Write to a database, update analytics
UI, do whatever you want

Example 1 in Java vs. Scala

Scala

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)

val hashTags = tweets.flatMap (status => getTags(status))

hashTags.saveAsHadoopFiles("hdfs://...")
```

Java

```
JavaDStream<Status> tweets = ssc.twitterStream(<Twitter username>,
<Twitter password>)

JavaDstream<String> hashTags = tweets.flatMap(new Function<...> { })  
                                Function object

hashTags.saveAsHadoopFiles("hdfs://...")
```

Spark program vs Spark Streaming program

Spark Streaming program on Twitter stream

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)
val hashTags = tweets.flatMap (status => getTags(status))
hashTags.saveAsHadoopFiles("hdfs://...")
```

Spark program on Twitter log file

```
val tweets = sc.hadoopFile("hdfs://...")
val hashTags = tweets.flatMap (status => getTags(status))
hashTags.saveAsHadoopFile("hdfs://...")
```

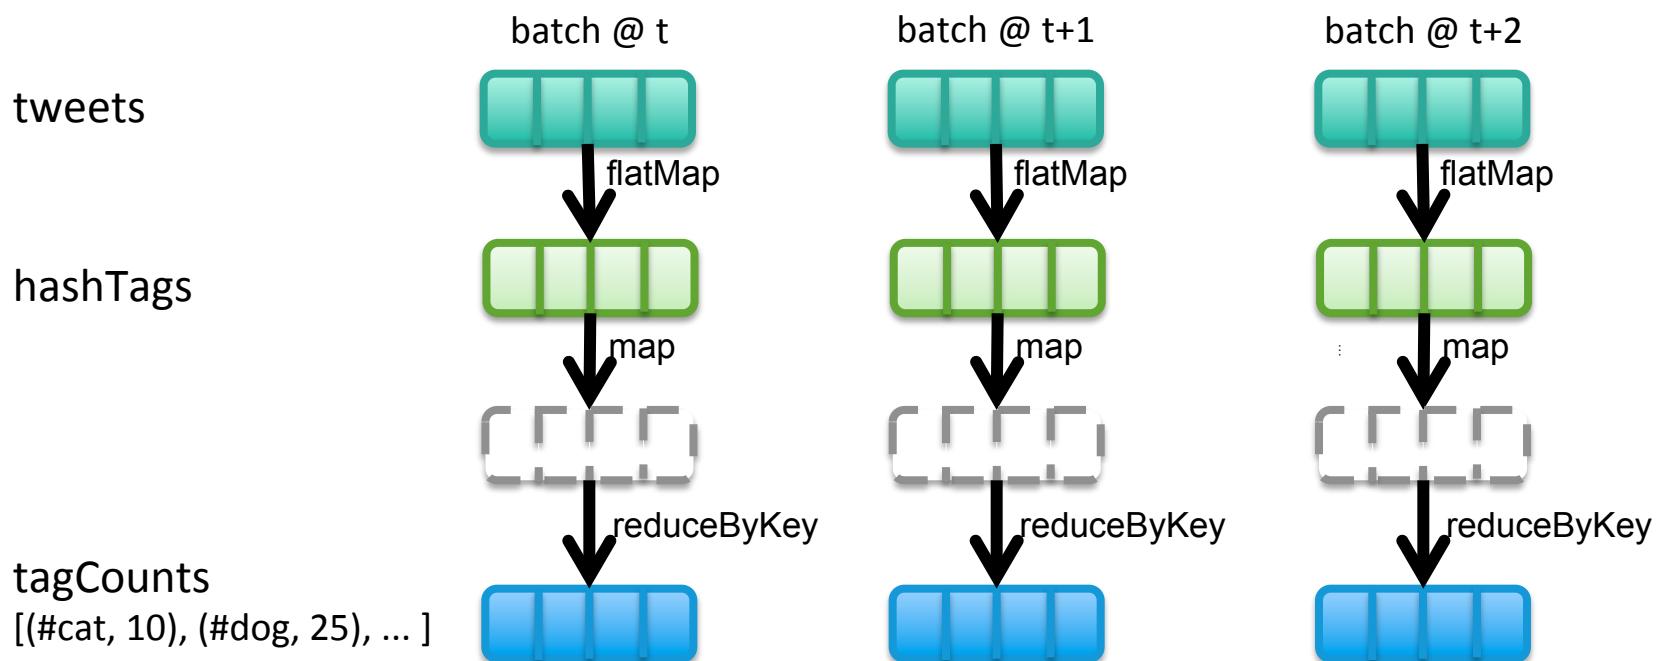
Vision - one stack to rule them all

- Explore data interactively using Spark Shell / PySpark to identify problems
- Use same code in Spark stand-alone programs to identify problems in production logs
- Use similar code in Spark Streaming to identify problems in live log streams

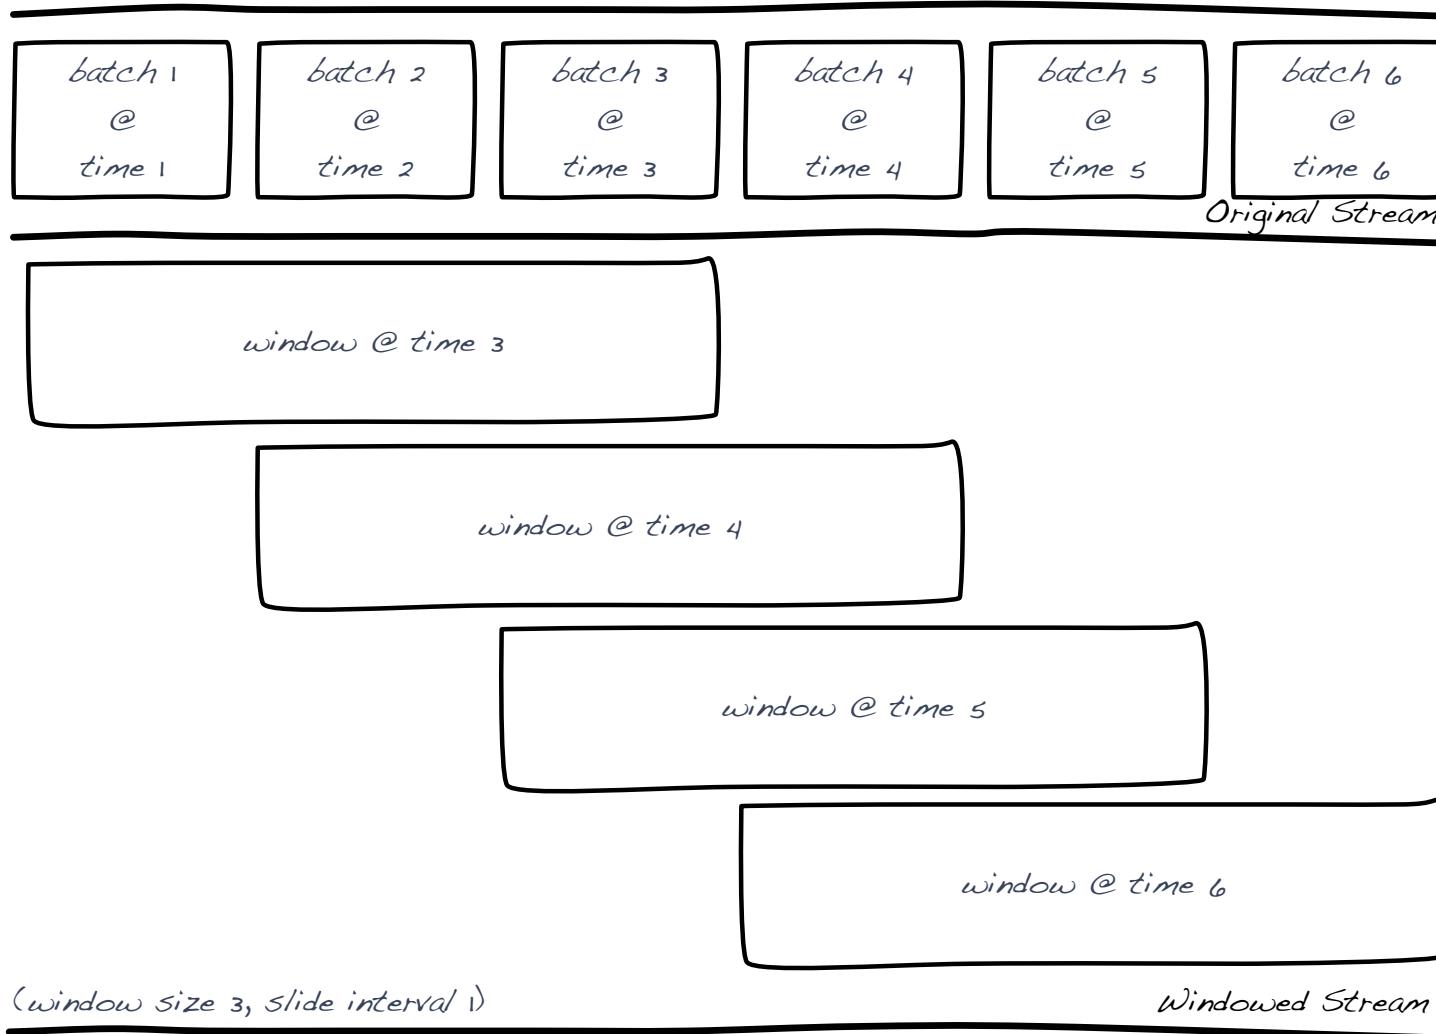
```
$ ./spark-shell
scala> val file = sc.hadoopFile("smallLogs")
...
scala> val filtered = file.filter(_.contains("ERROR"))
...
scala> val mapped = file.map(...)
object ProcessProductionData {
  def main(args: Array[String]) {
    val sc = new SparkContext(...)
    val file = sc.hadoopFile("productionLogs")
    val filtered =
      file.filter(_.contains("ERROR"))
    val mapped = file.map(...)
    ...
  }
}
object ProcessLiveStream {
  def main(args: Array[String]) {
    val sc = new StreamingContext(...)
    val stream = sc.kafkaStream(...)
    val filtered =
      stream.filter(_.contains("ERROR"))
    val mapped = stream.map(...)
    ...
  }
}
```

Example 2: Count the HashTags

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)
val hashTags = tweets.flatMap (status => getTags(status))
val tagCounts = hashTags.countByValue()
```

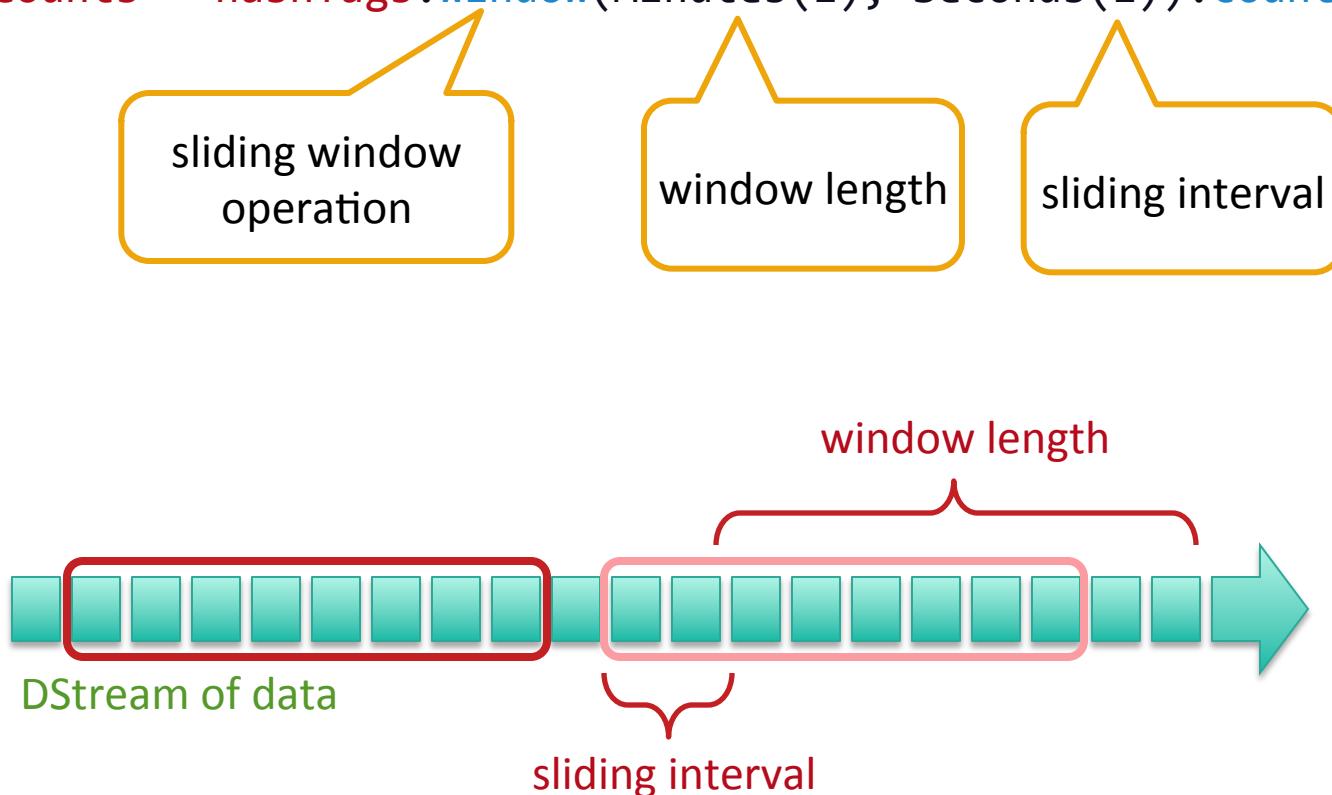


Window-based Operations on DStreams



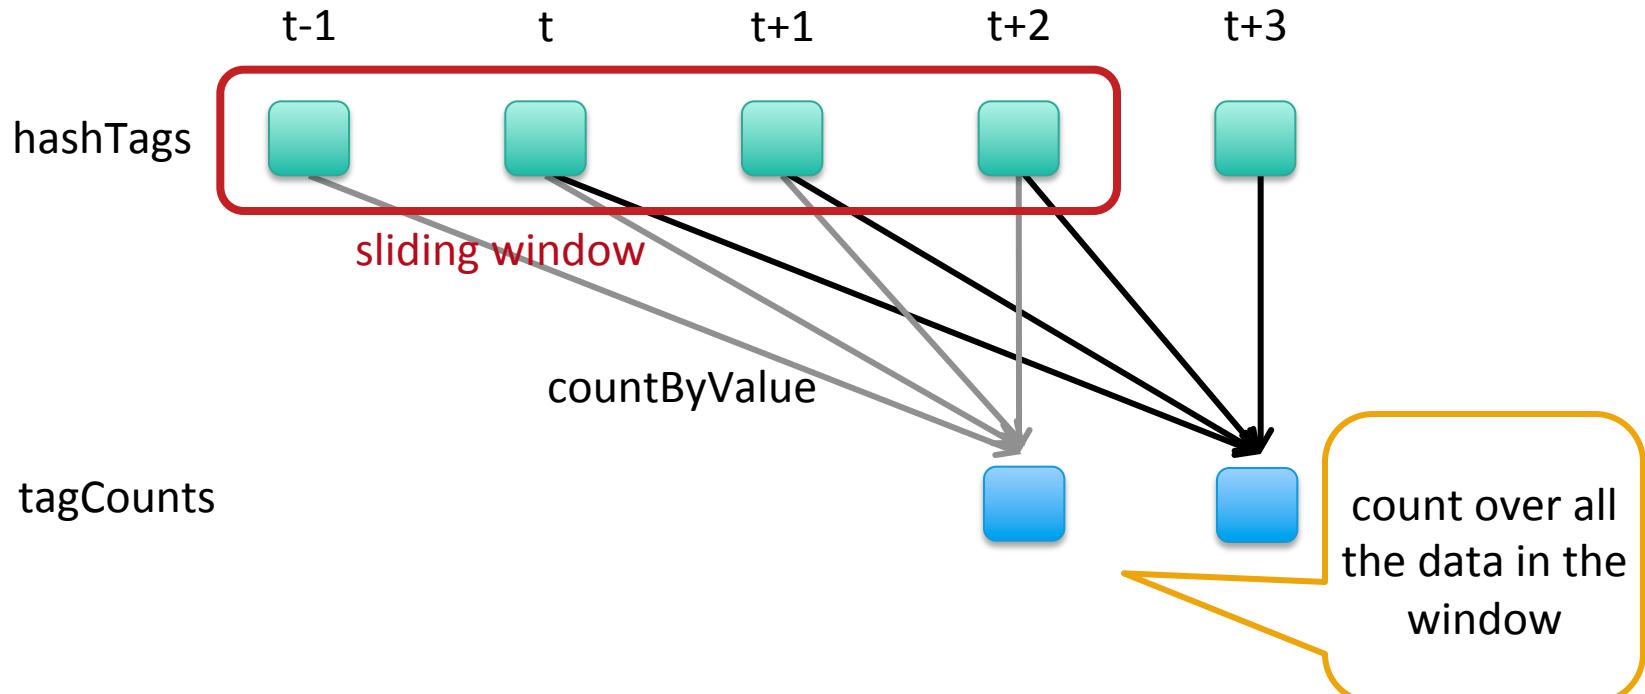
Example 3: Count the HashTags over last 1 min

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)
val hashTags = tweets.flatMap (status => getTags(status))
val tagCounts = hashTags.window(Minutes(1), Seconds(1)).countByValue()
```



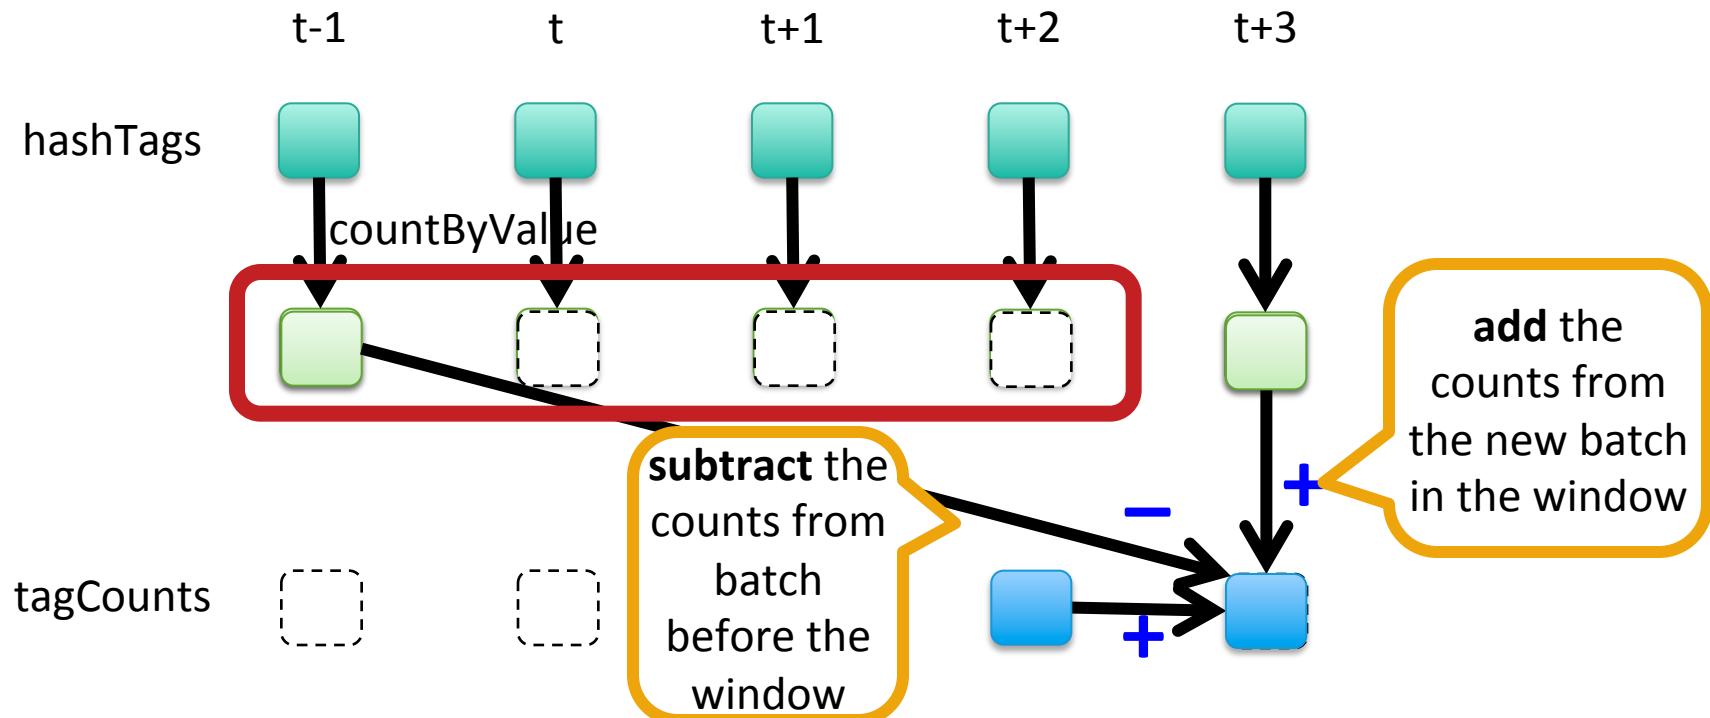
Example 3: Count the HashTags over last 1 min

```
val tagCounts = hashTags.window(Minutes(1), Seconds(1)).countByValue()
```



Example 3: Smart Window-based countByValue

```
val tagCounts = hashtags.countByValueAndWindow(Minutes(10),  
Seconds(1))
```



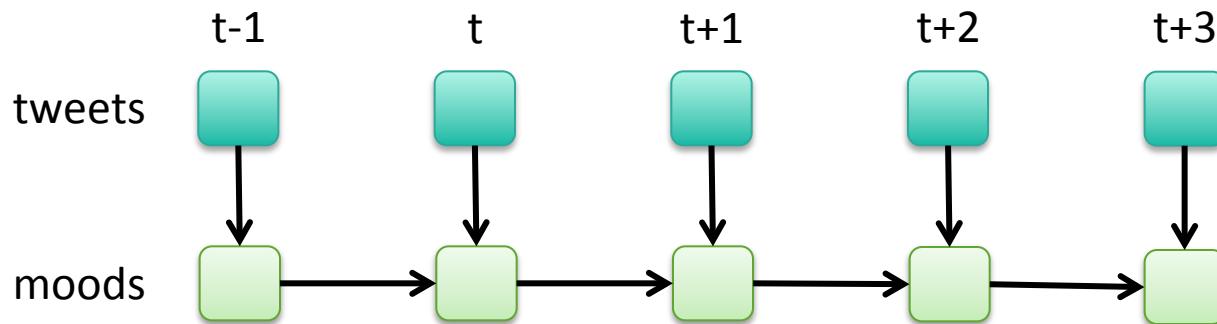
Smart window-based *reduce*

- Technique to incrementally compute count generalizes to many reduce operations
 - Need a function to “inverse reduce” (“subtract” for counting)
- Could have implemented counting as:
`hashTags.reduceByKeyAndWindow(_ + _, _ - _, Minutes(1), ...)`

Arbitrary Stateful Computations

- Maintain arbitrary state, track sessions and specify function to generate new state based on previous state and new data:
 - e.g. Maintain per-user mood as state, and update it with his/her tweets

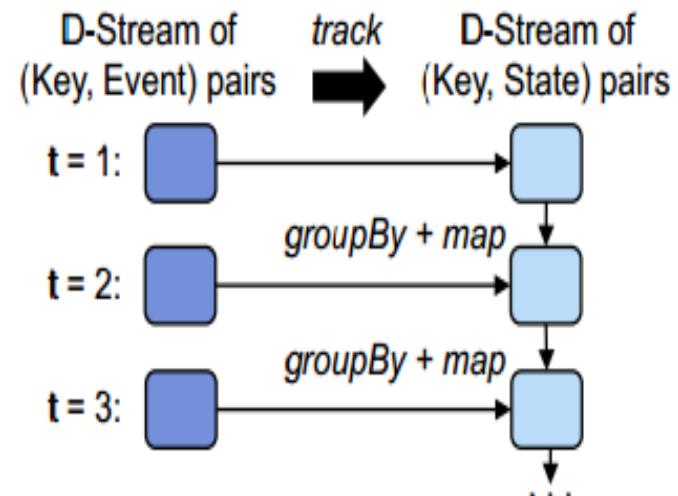
```
def updateMood(newTweets, lastMood) => newMood  
val moods = tweets.updateStateByKey(tweet => updateMood( _ ))
```



State Tracking

- A series of events → state changing

```
sessions = events.track(  
  (key, ev) => 1, // initialize function  
  (key, st, ev) => // update function  
    ev == Exit ? null : 1,  
    "30s") // timeout  
counts = sessions.count() // a stream  
of ints
```



Another Example: Word Count with Kafka

```
val context = new StreamingContext(conf, Seconds(1))
```

entry point of streaming functionality

```
val lines = KafkaUtils.createStream(context, ...)
```

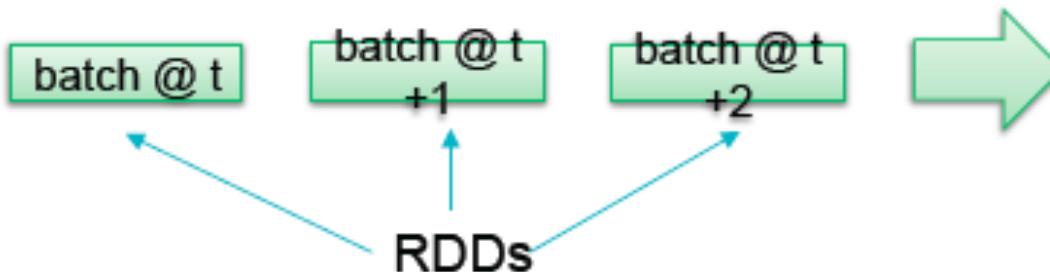
create DStream from Kafka data

Discretized Stream (DStream)

basic abstraction of Spark Streaming
series of RDDs representing a stream of

data

lines DStream



Another Example: Word Count with Kafka

```
val context = new StreamingContext(conf, Seconds(1))
```

```
val lines = KafkaUtils.createStream(context, ...)
```

```
val words = lines.flatMap(_.split(" "))
```

split lines into words

lines DStream

RDD @ t

RDD @ t+1

RDD @ t+2



flatMap

flatMap

flatMap

words DStream

RDD @ t

RDD @ t+1

RDD @ t+2



Another Example: Word Count with Kafka

```
val context = new StreamingContext(conf, Seconds(1))
```

```
val lines = KafkaUtils.createStream(context, ...)
```

```
val words = lines.flatMap(_.split(" "))
```

```
val wordCounts = words.map(x => (x, 1))
```

```
.reduceByKey(_ + _)
```

count the words

```
wordCounts.print()
```

print some counts on
screen

```
context.start()
```

start receiving and
transforming the data

Another Example: Word Count with Kafka

```
val context = new StreamingContext(conf, Seconds(1))

val lines = KafkaUtils.createStream(context, ...)

val words = lines.flatMap(_.split(" "))

val wordCounts = words.map(x => (x, 1))
    .reduceByKey(_ + _)

wordCounts.foreachRDD(rdd => /* do something */ )

context.start()
```

push data out to
storage systems

Many Transformations

Window operations

```
words.map(x => (x, 1)).reduceByKeyAndWindow(_ + _, Minutes(1))
```

Arbitrary stateful processing

```
def stateUpdateFunc(newData, lastState) => updatedState  
  
val stateStream = keyValueDStream.updateStateByKey(stateUpdateFunc)
```

Combine Batch and Stream Processing

- Inter-mix RDD and DStream operations
 - e.g., Join incoming tweets with a spam HDFS file to filter out bad tweets

```
tweets.transform(tweetsRDD => {  
    tweetsRDD.join(spamHDFSFile).filter(...)  
})
```

- Query streaming data using SQL, e.g.

```
select * from table_from_streaming_data
```

Another Example

```
val ssc = new StreamingContext(sc, Seconds(5))
val sqlContext = new SQLContext(sc)
val tweets = TwitterUtils.createStream(ssc, auth)
val transformed = tweets.filter(isEnglish).window(Minutes(1))

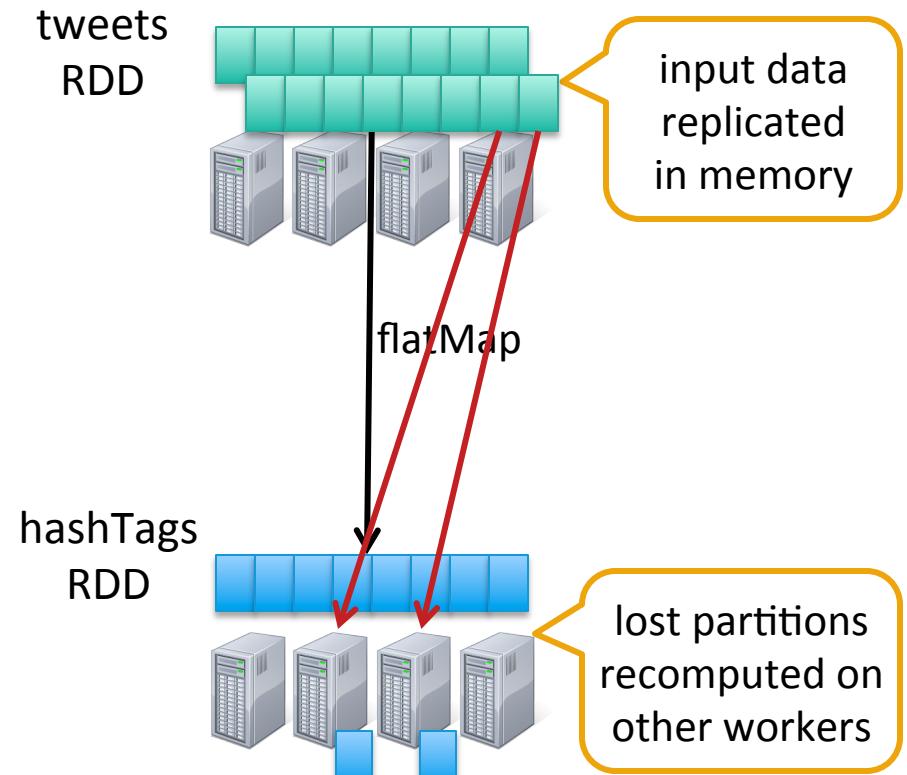
transformed.foreachRDD { rdd =>
    // Tweet is a case class containing necessary
    rdd.map(Tweet.apply(_)).registerAsTable("tweets")
}
```

```
SELECT text FROM tweets WHERE similarity(tweet) > 0.01
```

```
SELECT getClosestCountry(lat, long) FROM tweets
```

Fault Tolerance

- RDDs remember the operations that created them
- Batches of input data are replicated in memory for fault-tolerance
- Data lost due to worker failure, can be recomputed from replicated input data
- Therefore, all transformed data is fault-tolerant
- Exactly once semantics
 - No double counting



Input Sources

- Out of the box support for:
 - Kafka, Flume, Akka Actors, Raw TCP sockets, HDFS, etc
- Developers can write additional custom *receiver(s)*
 - Just define what to and when receiver is started and stopped
- Can also generate one's own sequence of RDDs and push them in as a “Stream”

Zero (Input) Data Loss during Streaming

For Non-replayable Sources, i.e. sources that do not support replay from any position (e.g. Flume, etc):

- Solved using Write Ahead Log (WAL) (since Spark 1.3)

For Replayable Sources, i.e. sources that allow data to be replayed from any position (e.g. Kafka, Kinesis, etc):

- Solved with more reliable Kafka and Kinesis Integrations (Spark 1.3-1.5)

Write Ahead Log (WAL) [since Spark 1.3]

- All received data synchronously written to HDFS and replayed when necessary after failure
- WAL can be enabled by setting Spark configuration flag:
`spark.streaming.receiver.writeAheadLog.enabled` to TRUE
- Can give end-to-end at least once guarantee for sources that can support acks, but do not support replays

Reliable Kinesis [since Spark 1.5]

- Save record sequence numbers instead of data to WAL
- Replay from Kinesis using sequence numbers
- Higher throughput than using WAL
- Can give at least once guarantee

Reliable Kafka [Spark 1.3, graduated in 1.5]

- New API: **Direct** Kafka stream:
 - Does not use receivers, does not use ZooKeeper to save offsets
 - Offset management (saving, replaying) by Spark Streaming
- Can provide up to 10x higher throughput than earlier receiver
 - <https://spark-summit.org/2015/events/towards-benchmarking-modern-distributed-streaming-systems/>
- Can give exactly-once guarantee (excluding o/p to storage)
- Can run Spark batch jobs directly on Kafka
 - # of RDD partitions = # of Kafka partitions, easy to reason about
 - <https://databricks.com/blog/2015/03/30/improvements-to-kafka-integration-of-spark-streaming.html>

Discretized Stream (DStream)

A sequence of RDDs representing
a stream of data

What does it take to define a DStream?

DStream Interface

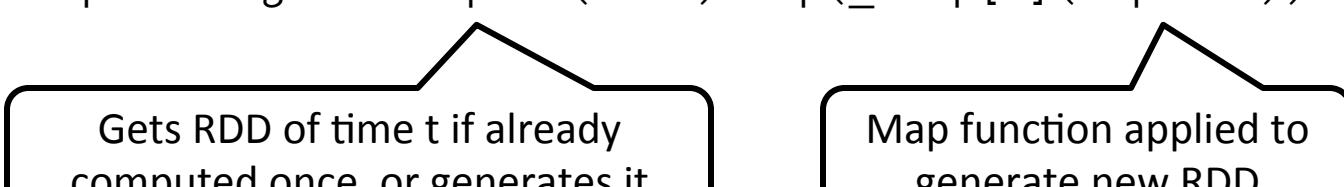
The DStream interface primarily defines how to generate an RDD in each batch interval

- List of *dependent* (parent) DStreams
- *Slide Interval*, the interval at which it will compute RDDs
- Function to *compute* RDD at a time t

Example: Mapped DStream

- *Dependencies:* Single parent DStream
- *Slide Interval:* Same as the parent DStream
- *Compute function for time t:* Create new RDD by applying map function on parent DStream's RDD of time t

```
override def compute(time: Time): Option[RDD[U]] = {  
    parent.getOrCompute(time).map(_.map[U](mapFunc))  
}
```

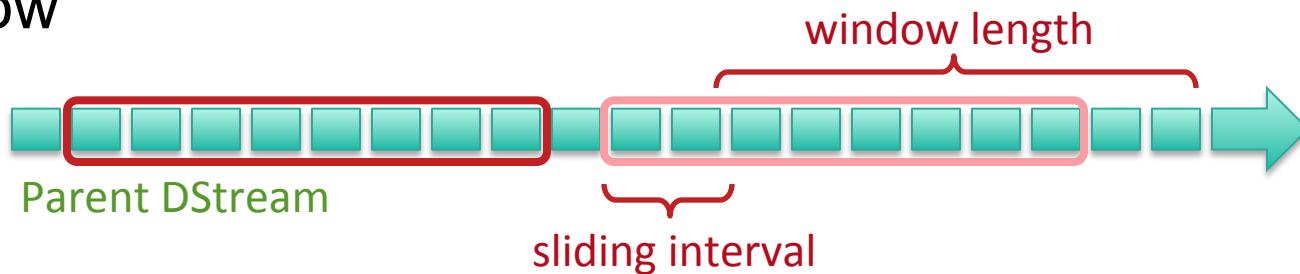


Gets RDD of time t if already computed once, or generates it

Map function applied to generate new RDD

Example: Windowed DStream

Window operation gather together data over a sliding window



Dependencies: Single parent DStream

Slide Interval: Window sliding interval

Compute function for time t: Apply union over all the RDDs of parent DStream between times t and $(t - \text{window length})$

Example: Network Input DStream

Base class of all input DStreams that receive data from the network

- *Dependencies:* None
- *Slide Interval:* Batch duration in streaming context
- *Compute function for time t:* Create a BlockRDD with all the blocks of data received in the last batch interval
- Associated with a **Network Receiver** object

Network Receiver

Responsible for receiving data and pushing it into Spark's data management layer (Block Manager)

Base class for all receivers - Kafka, Flume, etc.

Simple Interface:

- What to do *on starting* the receiver
 - Helper object *blockGenerator* to push data into Spark
- What to do *on stopping* the receiver

Example: Socket Receiver

- *On start:*

- Connect to remote TCP server

- While socket is connected,

- Receiving bytes and deserialize

- Deserialize them into Java objects

- Add the objects to *blockGenerator*

- *On stop:*

- Disconnect socket

Other functions in DStream interface

- *parentRememberDuration* – defines how long should
 - Window-based DStreams have
parentRememberDuration = window length
- *mustCheckpoint* – if set to true, the system will automatically enable periodic checkpointing
 - Set to true for stateful DStreams

DStream Graph

Spark Streaming program

```
t = ssc.twitterStream("...")  
    .map(...)  
t.foreach(...)
```



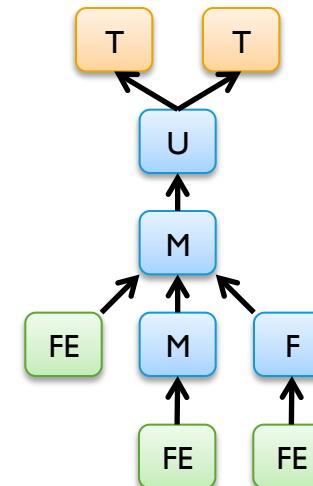
Dummy DStream signifying
an output operation

DStream Graph



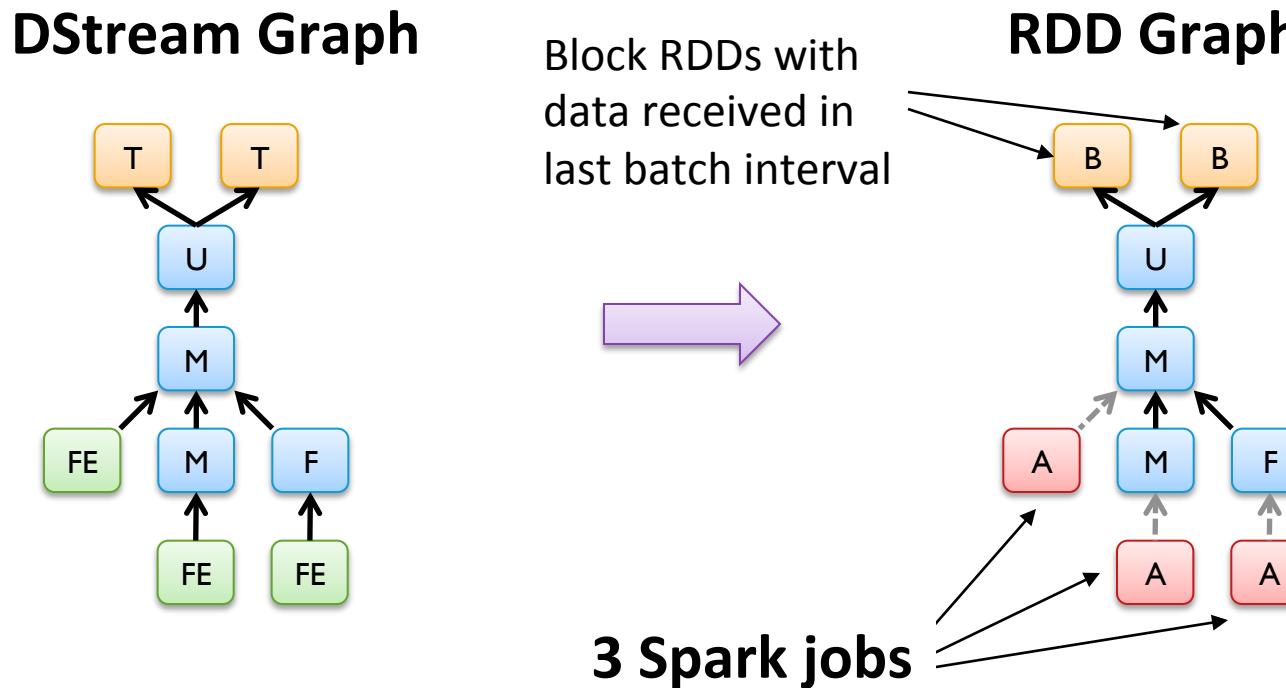
```
t1 = ssc.twitterStream("...")  
t2 = ssc.twitterStream("...")
```

```
t = t1.union(t2).map(...)  
  
t.saveAsHadoopFiles(...)  
t.map(...).foreach(...)  
t.filter(...).foreach(...)
```



DStream Graph → RDD Graphs → Spark jobs

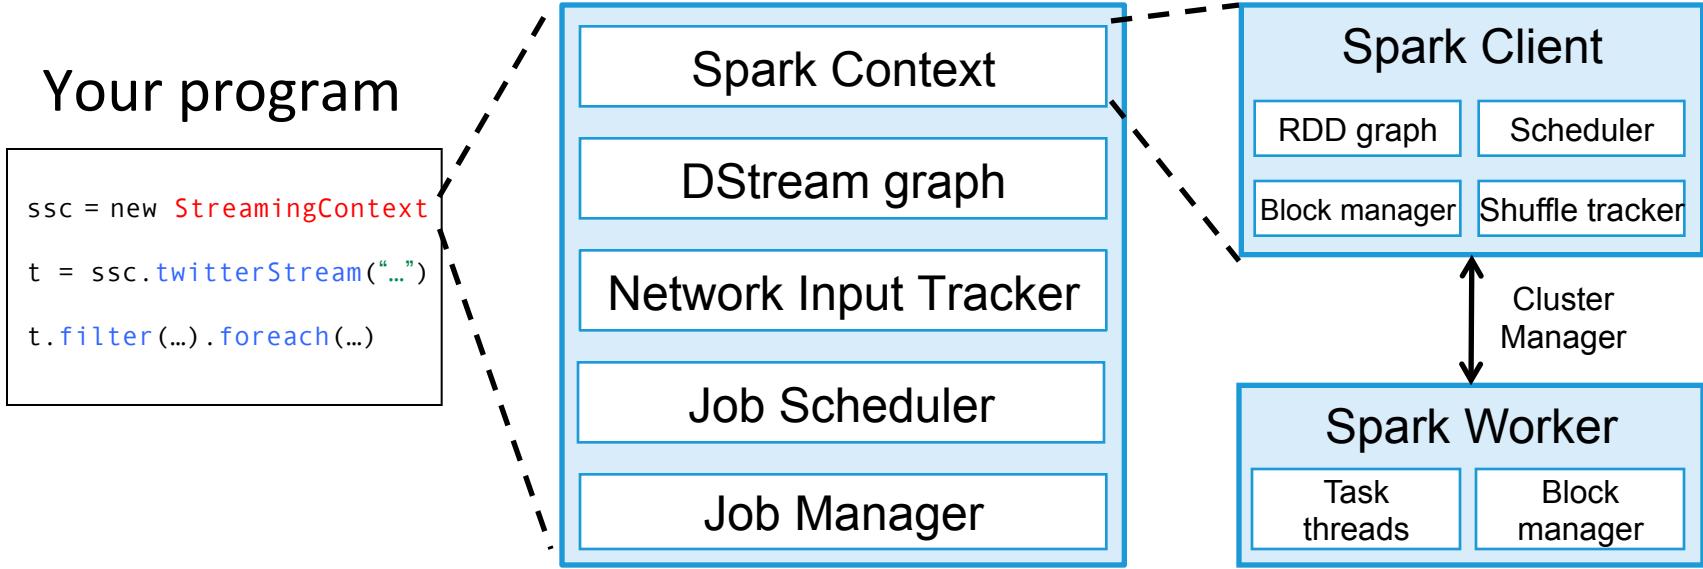
- Every interval, RDD graph is computed from DStream graph
- For each output operation, a Spark action is created
- For each action, a Spark job is created to compute it



Agenda

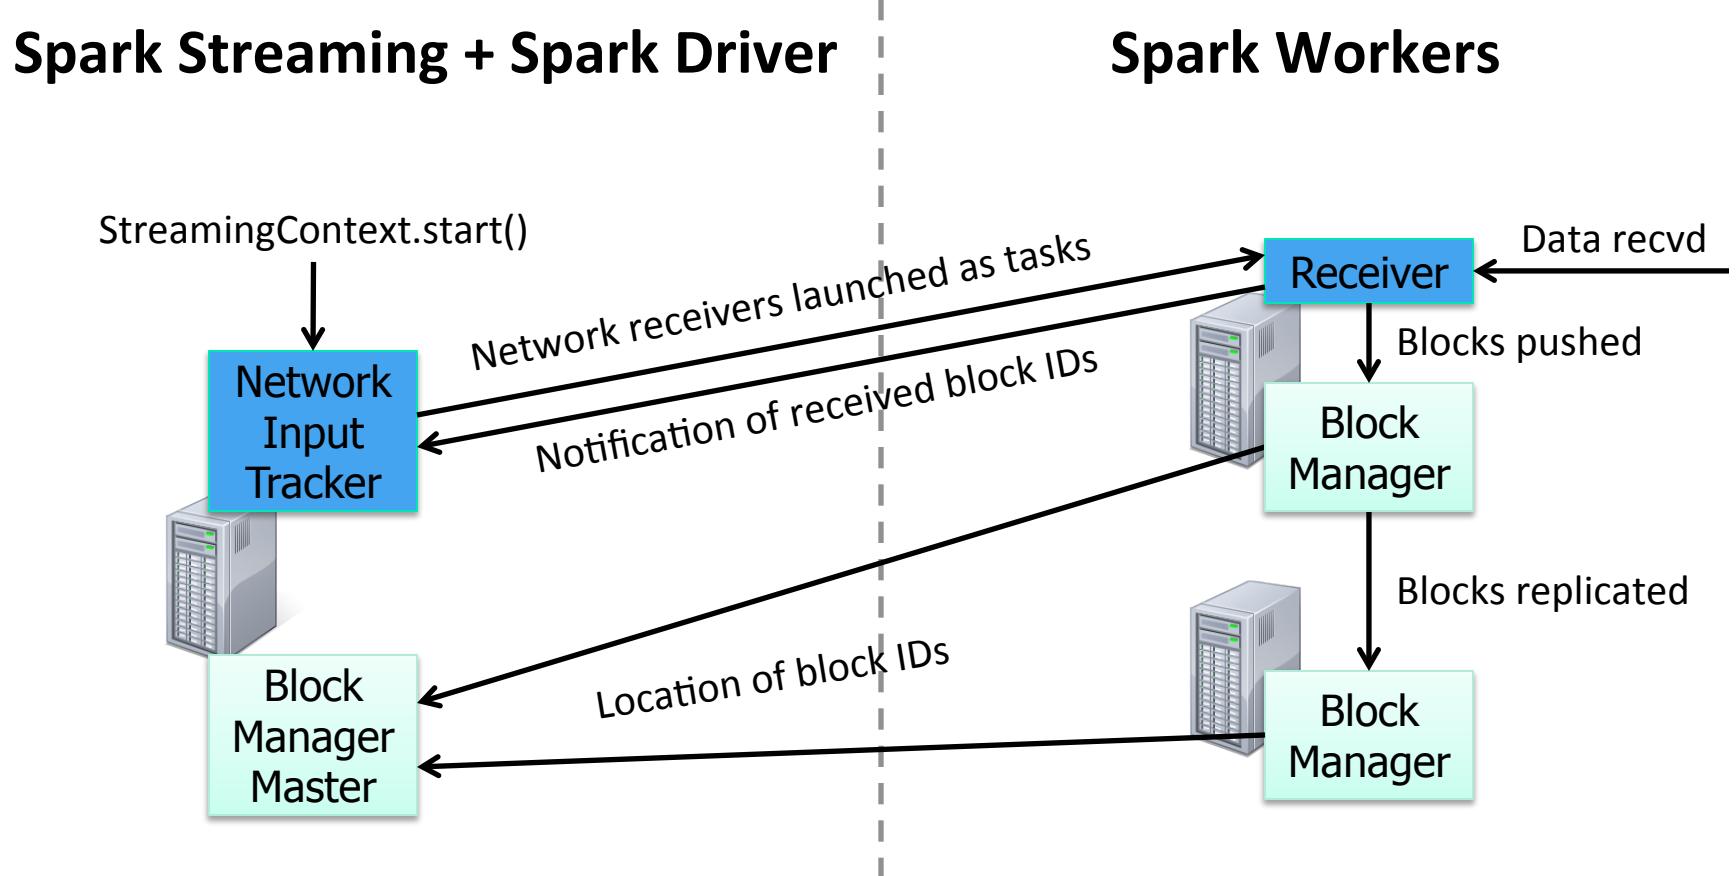
- Overview
- DStream Abstraction
- System Model
- Persistence / Caching
- RDD Checkpointing
- Performance Tuning

Components



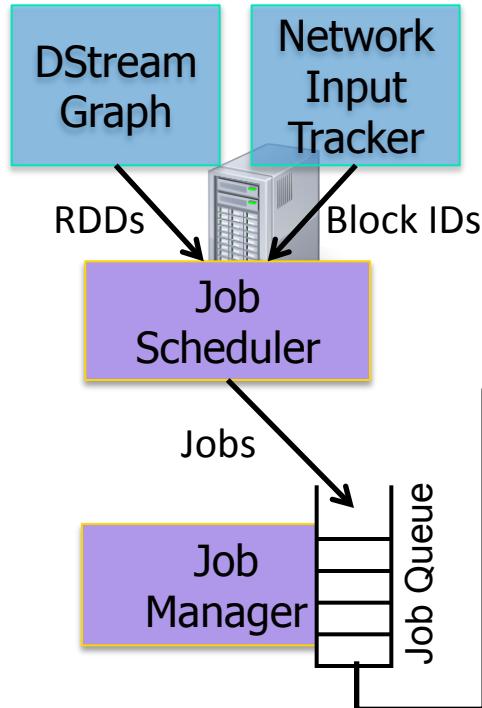
- **Network Input Tracker** – Keeps track of the data received by each network receiver and maps them to the corresponding input DStreams
- **Job Scheduler** – Periodically queries the DStream graph to generate Spark jobs from received data, and hands them to Job Manager for execution
- **Job Manager** – Maintains a job queue and executes the jobs in Spark

Execution Model – Receiving Data

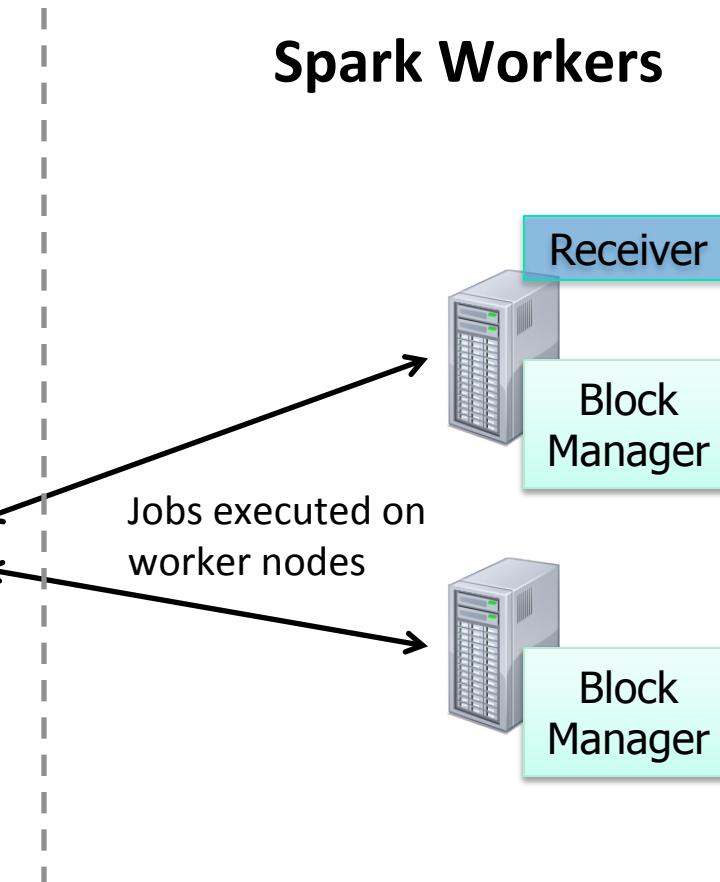


Execution Model – Job Scheduling

Spark Streaming + Spark Driver



Spark Workers



Job Scheduling

- Each output operation used generates a job
 - More jobs → more time taken to process batches → higher batch duration
- Job Manager decides how many concurrent Spark jobs to run
 - Default is 1, can be set using Java property `spark.streaming.concurrentJobs`
 - If you have multiple output operations, you can try increasing this property to reduce batch processing times and so reduce batch duration

Agenda

- Overview
- DStream Abstraction
- System Model
- Persistence / Caching
- RDD Checkpointing
- Performance Tuning

DStream Persistence

- If a DStream is set to persist at a storage level, then all RDDs generated by it set to the same storage level
- When to persist?
 - If there are multiple transformations / actions on a DStream
 - If RDDs in a DStream is going to be used multiple times
- Window-based DStreams are automatically persisted in memory

DStream Persistence

- Default storage level of DStreams is `StorageLevel.MEMORY_ONLY_SER` (i.e. in memory as serialized bytes)
 - Except for input DStreams which have `StorageLevel.MEMORY_AND_DISK_SER_2`
 - Note the difference from RDD's default level (no serialization)
 - Serialization reduces random pauses due to GC providing more consistent job processing times

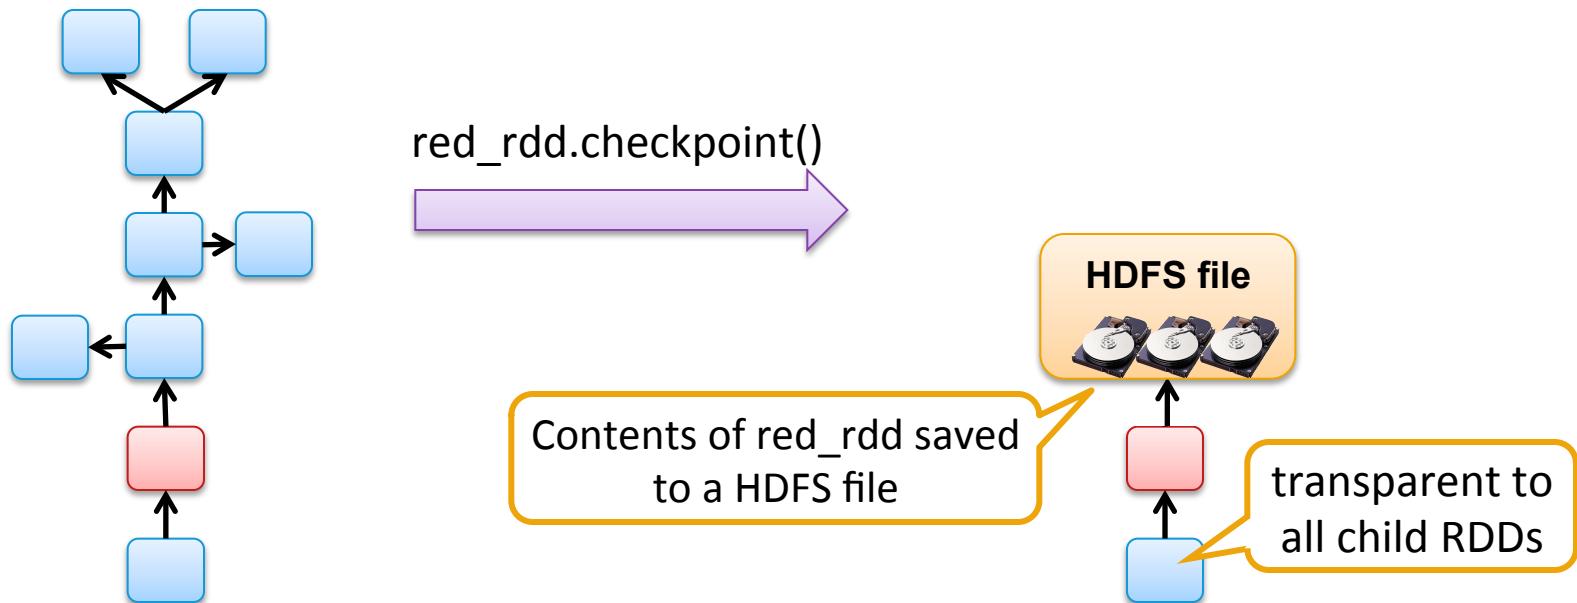
Agenda

- Overview
- DStream Abstraction
- System Model
- Persistence / Caching
- RDD Checkpointing
- Performance Tuning

What is RDD checkpointing?

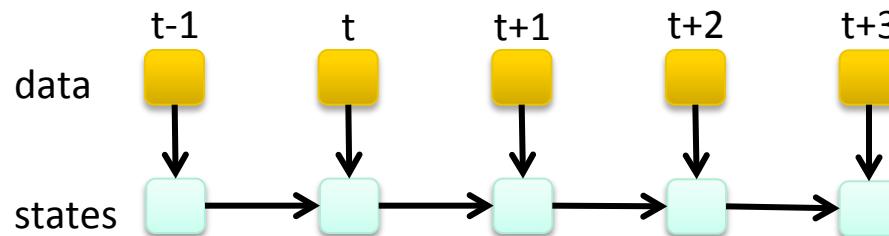
Saving RDD to HDFS to prevent RDD graph from growing too large

- Done internally in Spark transparent to the user program
- Done lazily, saved to HDFS the first time it is computed



Why is RDD checkpointing necessary?

Stateful DStream operators can have infinite lineages

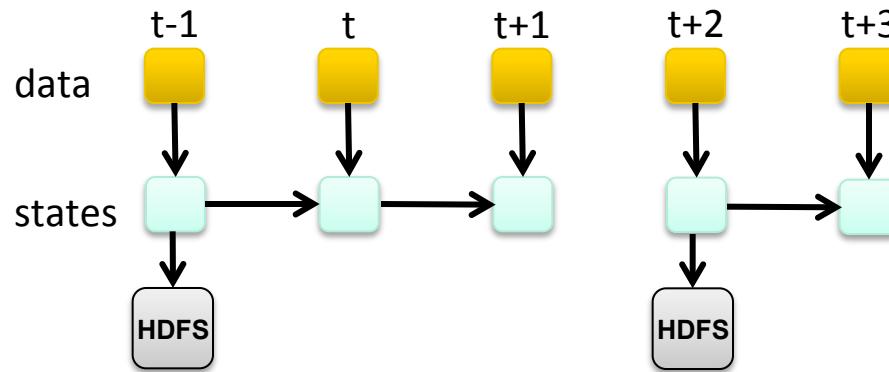


Large lineages lead to ...

- Large closure of the RDD object → large task sizes → high task launch times
- High recovery times under failure

Why is RDD checkpointing necessary?

Stateful DStream operators can have infinite lineages



Periodic RDD checkpointing solves this

Useful for iterative Spark programs as well

RDD Checkpointing

- Periodicity of checkpoint determines a tradeoff
 - Checkpoint too frequent: HDFS writing will slow things down
 - Checkpoint too infrequent: Task launch times may increase
 - Default setting checkpoints at most once in 10 seconds
 - Try to checkpoint once in about 10 batches

Agenda

- Overview
- DStream Abstraction
- System Model
- Persistence / Caching
- RDD Checkpointing
- Performance Tuning

Performance Tuning

Step 1

Achieve a stable configuration that can sustain the streaming workload

Step 2

Optimize for lower latency

Step 1: Achieving Stable Configuration

How to identify whether a configuration is stable?

- Look for the following messages in the log

Total delay: 0.01500 s for job 12 of time 1371512674000 ...

- If the total delay is continuously increasing, then unstable as the system is unable to process data as fast as its receiving!

- If the total delay stays roughly constant and around 2x the configured batch duration, then stable

Step 1: Achieving Stable Configuration

How to figure out a good stable configuration?

- Start with a low data rate, small number of nodes, reasonably large batch duration (5 – 10 seconds)
- Increase the data rate, number of nodes, etc.
- Find the bottleneck in the job processing
 - Jobs are divided into stages
 - Find which stage is taking the most amount of time

Step 1: Achieving Stable Configuration

How to figure out a good stable configuration?

- If the first map stage on raw data is taking most time, then try ...
 - Enabling delayed scheduling by setting property `spark.locality.wait`
 - Splitting your data source into multiple sub streams
 - Repartitioning the raw data into many partitions as first step
- If any of the subsequent stages are taking a lot of time, try ...
 - Try increasing the level of parallelism (i.e., increase number of reducers)
 - Add more processors to the system

Step 2: Optimize for Lower Latency

- Reduce batch size and find a stable configuration again
 - Increase levels of parallelism, etc.
- Optimize serialization overheads
 - Consider using Kryo serialization instead of the default Java serialization for both data and tasks
 - For data, set property

```
spark.serializer=spark.KryoSerializer
```
 - For tasks, set

```
spark.closure.serializer=spark.KryoSerializer
```
- Use Spark stand-alone mode rather than Mesos

Step 2: Optimize for Lower Latency

- Using concurrent mark sweep GC -xx:
+UseConcMarkSweepGC is recommended
- Reduces throughput a little, but also reduces large GC pauses and may allow lower batch sizes by making processing time more consistent
- Try disabling serialization in DStream/RDD persistence levels
- Increases memory consumption and randomness of GC related pauses, but may reduce latency by further reducing serialization overheads
- For a full list of guidelines for performance tuning
 - [Spark Tuning Guide](#)
 - [Spark Streaming Tuning Guide](#)

System Stability

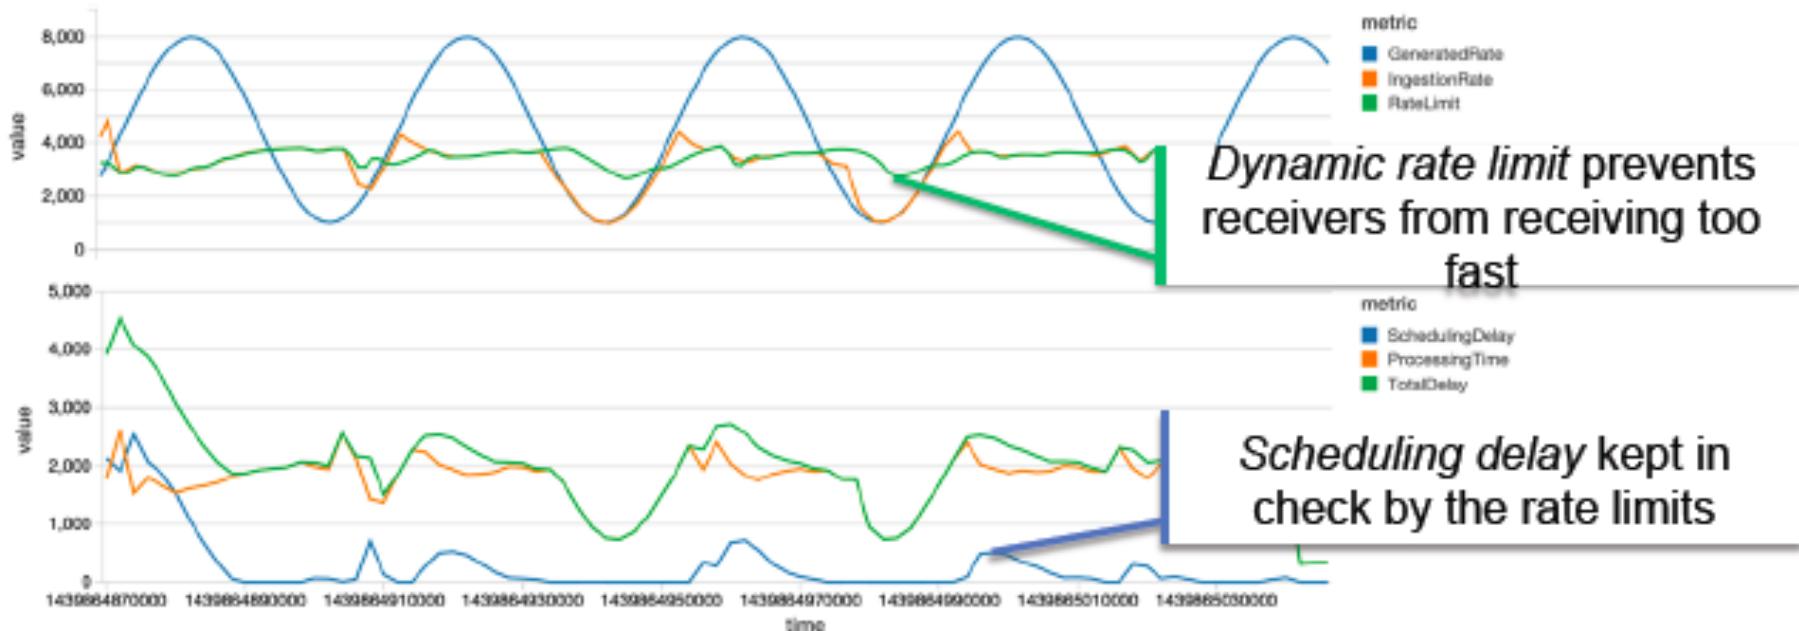
- Streaming applications may have to deal with variations in data rates and processing rates
- For stability, any streaming application must receive data only as fast as it can process
- Static rate limits on receivers [Spark 1.1]
 - But hard to figure out the right rate

Backpressure [Spark 1.5]

- System **automatically** and **dynamically** adapts rate limits to ensure stability under any processing conditions
- If sinks slow down, then the system automatically pushes back on the source to slow down receiving
- System uses batch processing times and scheduling delays experienced to set rate limits
- Well known PID controller theory (used in industrial control systems) is used to calculate appropriate rate limit
 - Contributed by Typesafe
- Enabled by setting Spark configuration flag
 - `spark.streaming.backpressure.enabled` to TRUE

Backpressure [Spark 1.5]

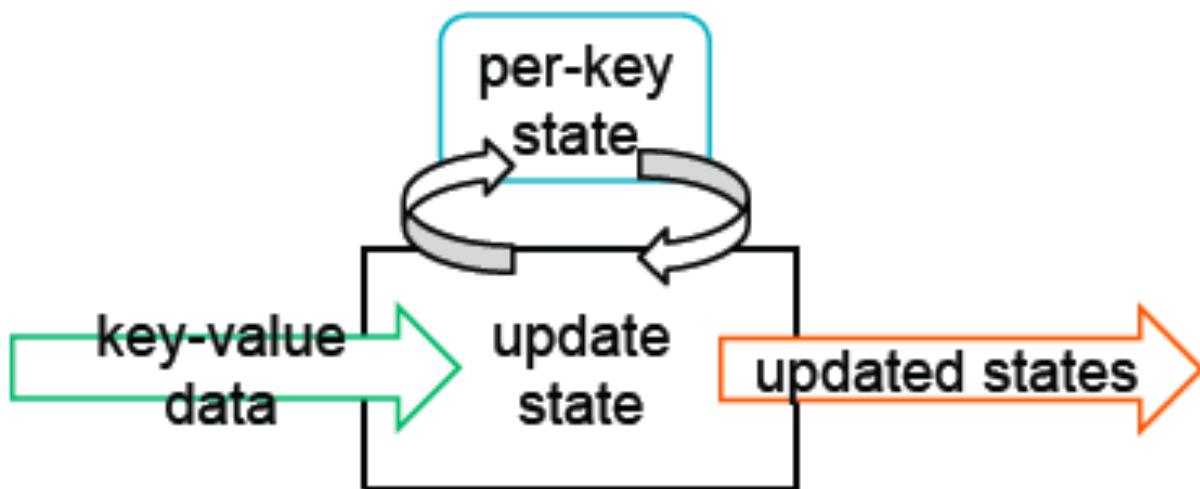
- System automatically and dynamically adapts rate limits



Improved State Management

Earlier stateful stream processing done with
`updateStateByKey`

```
def stateUpdateFunc(newData, lastState) => updatedState  
  
val stateDStream = keyValueDStream.updateStateByKey(stateUpdateFunc)
```



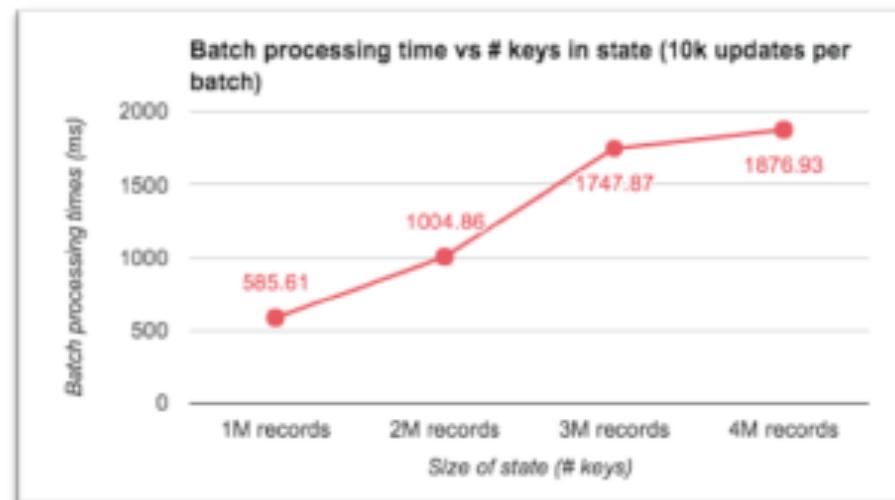
Improved State Management

Feedback from community about `updateStateByKey`

Need to keep much larger state

Processing times of batches increase with the amount state, limits performance

Need to expire keys that have received no data for a while



Improved State Management

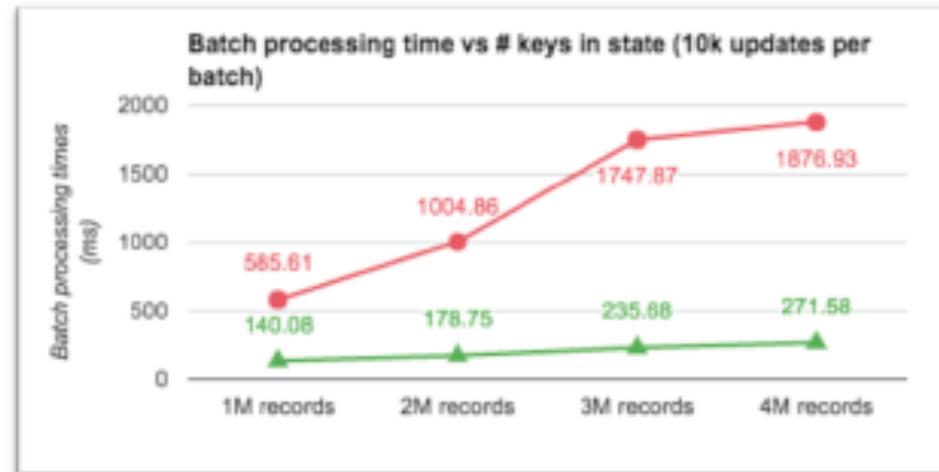
New API with timeouts: `trackStateByKey`

```
def updateFunc(values, state) => emittedData
  // call state.update(newState) to update state

keyValueDStream.trackStateByKey(  
  StateSpec.function(updateFunc).timeout(Minutes(10)))
```

Can provide order of magnitude higher performance than `updateStateByKey`

<https://issues.apache.org/jira/browse/SPARK-2629>



Visualizations

Streaming Statistics

Running batches of 1 second for 13 minutes 14 seconds since 2015/06/06 10:57:31 (794 completed batches, 1653 records)



Stats over last 1000 batches

For stability
Scheduling delay should be approx 0
Processing Time approx < batch interval

Visualizations

Details of individual batches

Active Batches (1)

Batch Time	Input Size	Scheduling Delay	Processing Time	Status
2015/06/08 11:10:46				ISSIN

Completed Batches (last 794)

Batch Time
2015/06/08 11:10:45
2015/06/08 11:10:44
2015/06/08 11:10:43
2015/06/08 11:10:42
2015/06/08 11:10:41
2015/06/08 11:10:40

Details of batch at 2015/06/29 15:34:00

Batch Duration: 30 s
Input data size: 0 records
Scheduling delay: 0 ms
Processing time: 72 ms
Total delay: 72 ms
Input Metadata:

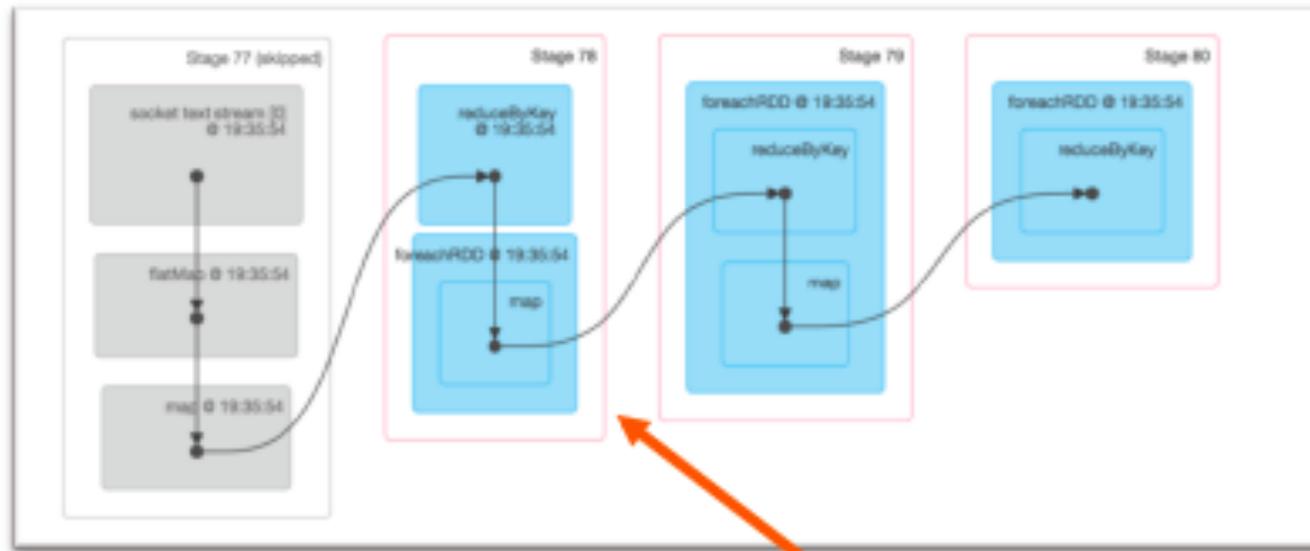
Input	Metadata
File stream [1]	file:/Users/zsx/streaming/a.txt, file:/Users/zsx/streaming/b.txt, file:/Users/zsx/streaming/c.txt, file:/Users/zsx/
Kafka direct stream [0]	OffsetRange(topic: 'test', partition: 0, range: [1 -> 1]), OffsetRange(topic: 'test2', partition: 0, range: [1 -> 1])

Kafka offsets processed in each batch,
Can help in debugging bad data

List of Spark jobs in each batch

Output Op Id	Description	Duration	Job Id	Duration	Stages: Succeeded/Total	Tags
0	print at DirectKafkaWordCount.scala:69	59 ms	6	48 ms	2/2	
			7	7 ms	1/1 (1 skipped)	
			8	4 ms	1/1 (1 skipped)	

Visualizations

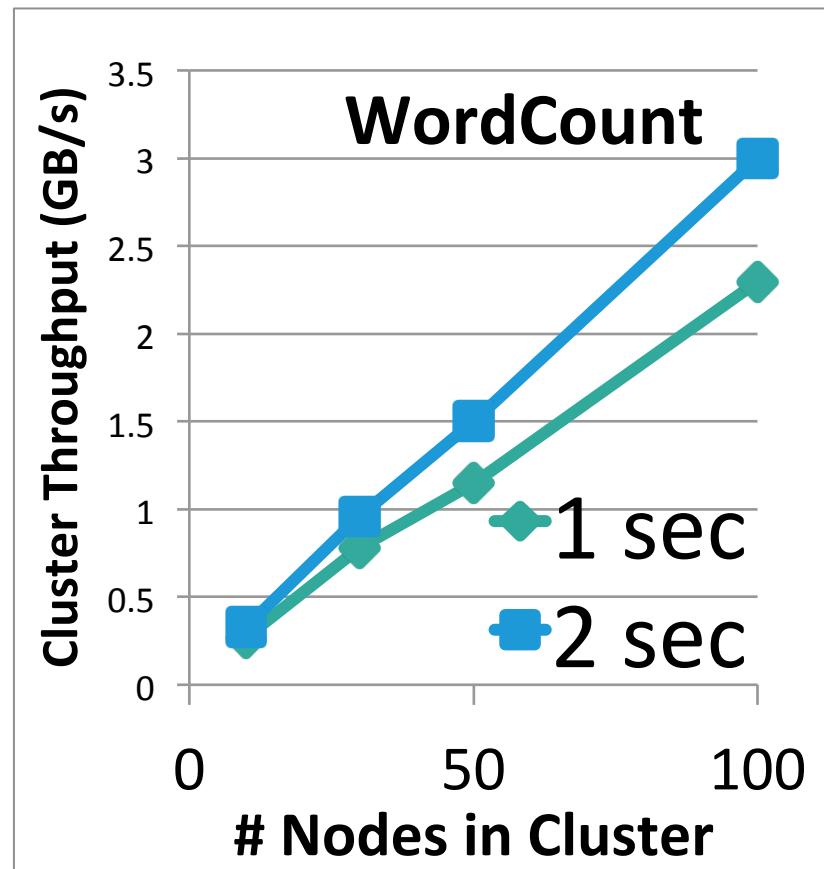
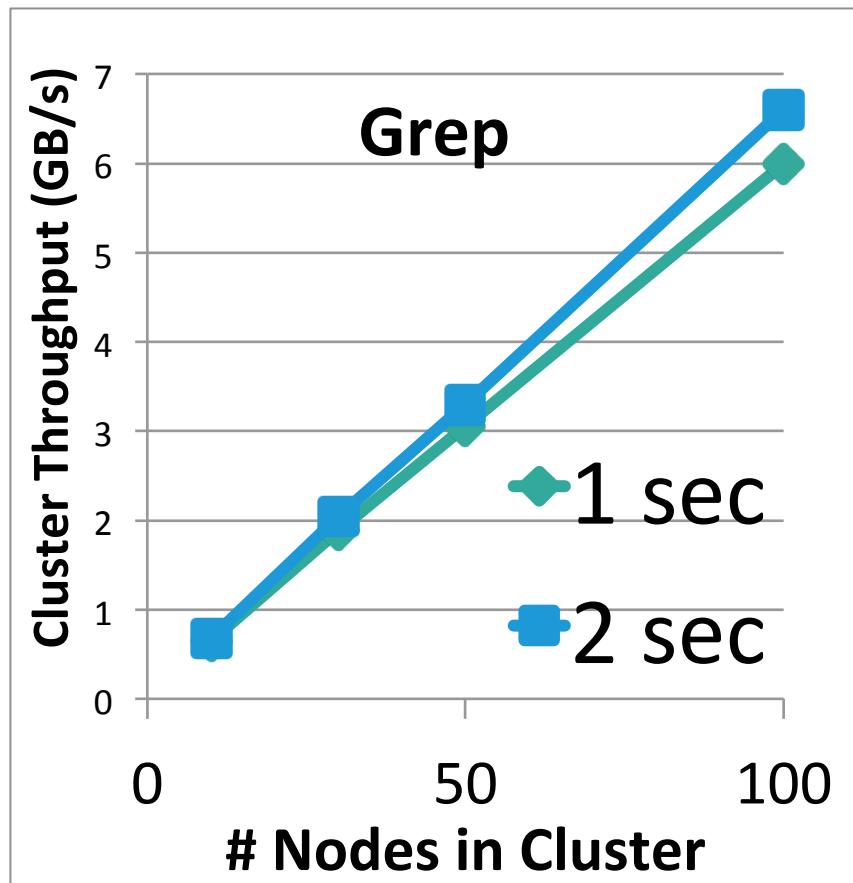


Full DAG of RDDs and stages generated by Spark Streaming

Performance

Can process **6 GB/sec (60M records/sec)** of data on 100 nodes at **sub-second latency**

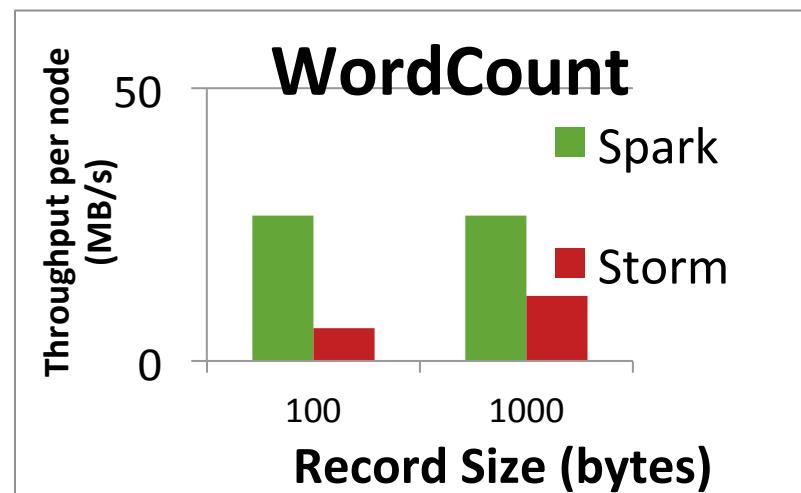
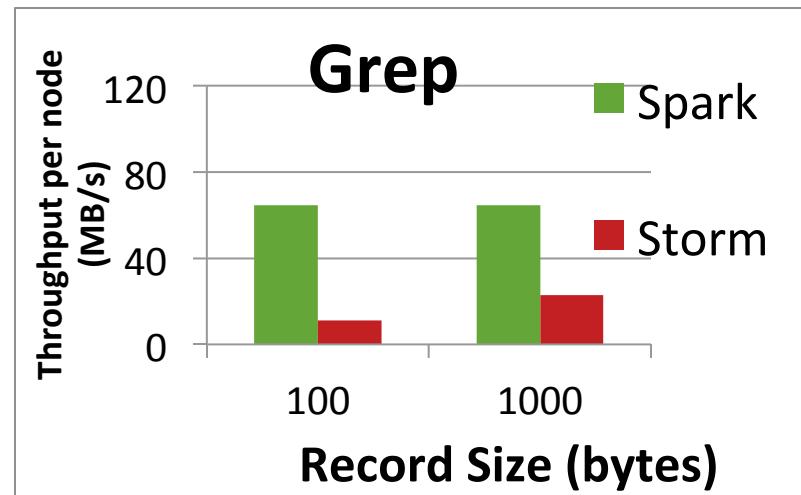
- Tested with 100 streams of data on 100 EC2 instances with 4 cores each



Comparison with Storm and S4

Higher throughput
than Storm

- Spark Streaming: **670k** records/second/node
- Storm: **115k** records/second/node
- Apache S4: 7.5k records/second/node



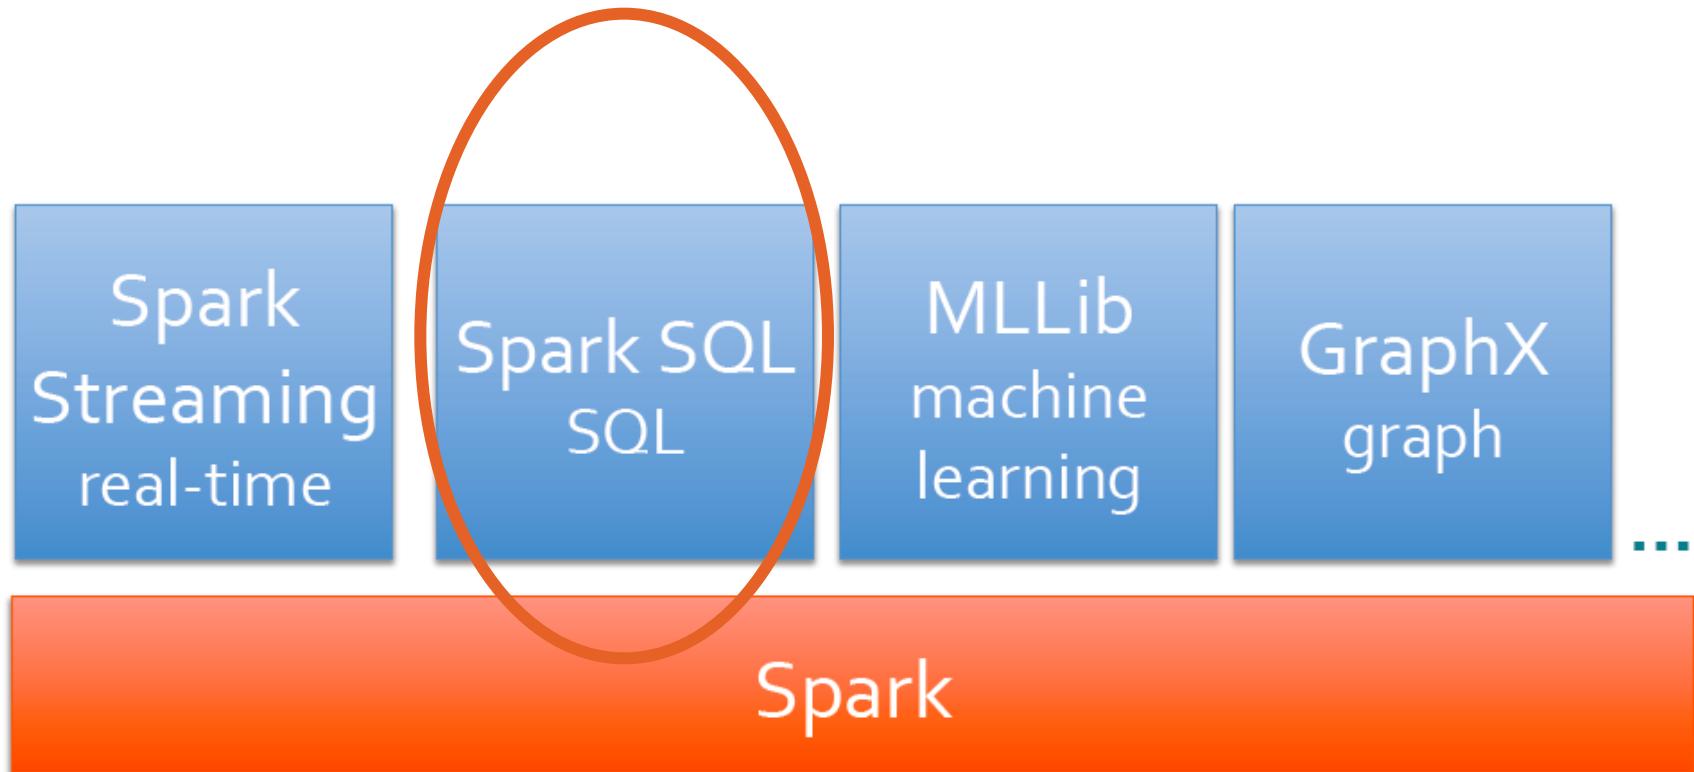
Small code base

- 5000 LOC for Scala API (+ 1500 LC for Java API)
 - Most DStream code mirrors the RDD code

Future Extensions on Spark Streaming

- API and Libraries Support of Streaming DataFrames
 - Logical-to-physical plan optimizations
 - Tungsten-based binary optimizations
 - Support for event-time based windowing
 - Support for Out-of-Order Data
- Add Native Infrastructure support for Dynamic Allocation for Streaming
 - Dynamically scale the cluster resources based on processing load
 - Need to work with Backpressure to scale up/down while maintaining stability
- Programmable monitoring by exposing more info via Streaming Listener
- Performance Enhancement: higher throughput and lower latency, specially for Stateful Ops, e.g. trackStateByKey

Generality of RDDs in Spark



Before SQL support was available from Spark



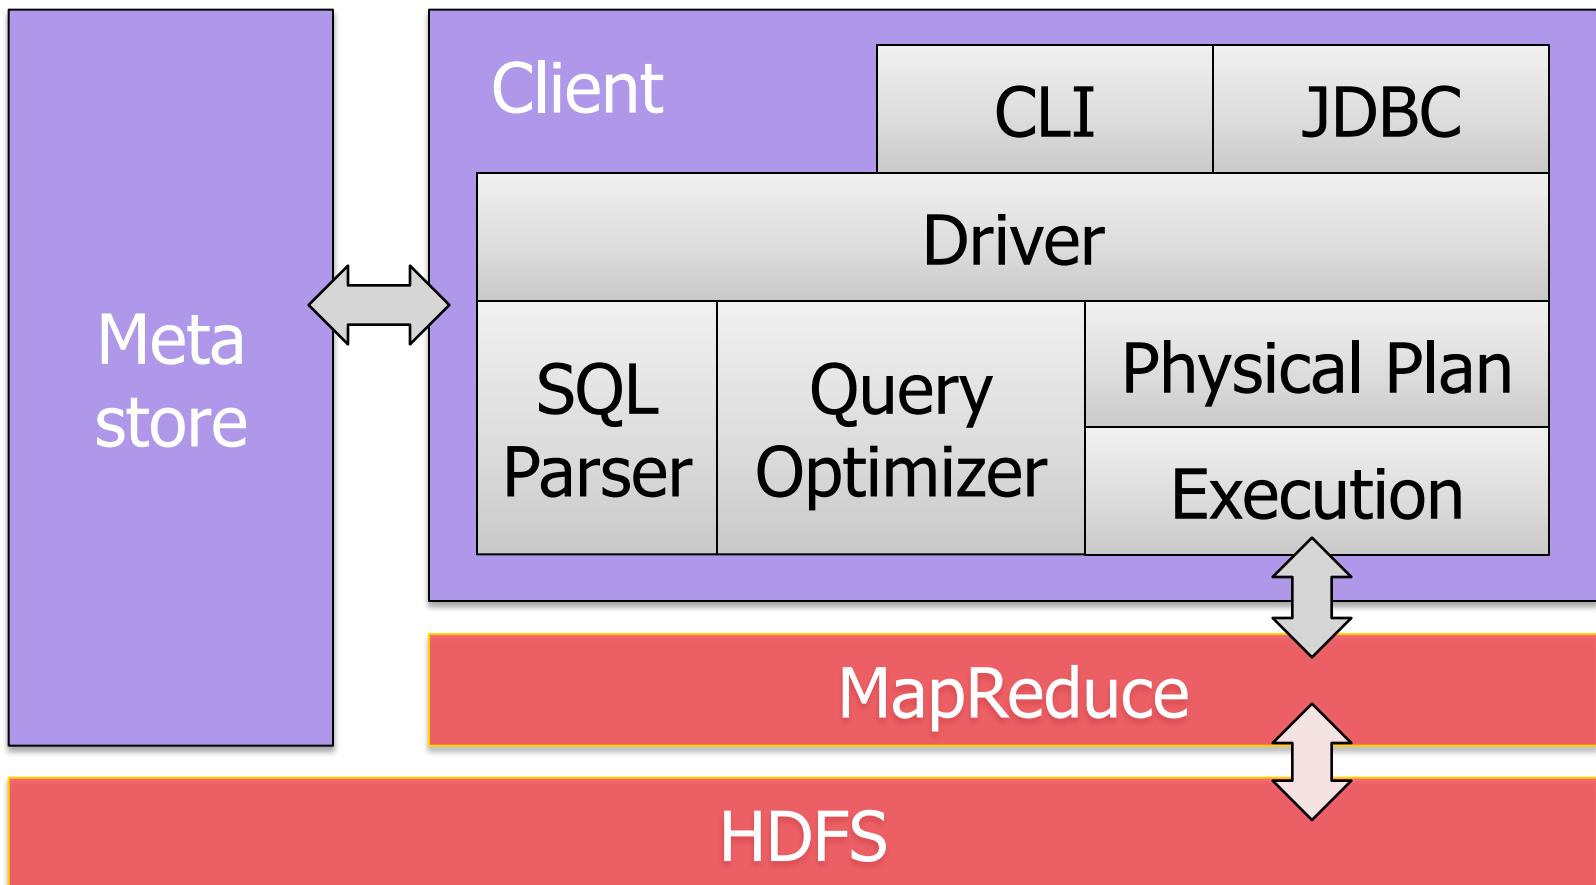
- The Spark Core Engine does not understand the structure of the data in RDDs or the semantics of user functions → limited optimization.
- However, most data is structured, e.g. JSON, CSV, Avro, Parquet, Hive, etc
 - => Programming/ Operations via the RDD API inevitably ends up with a lot of tuples (_1, _2, ...)
- Functional Transformations, e.g. Map/Reduce are still not as Intuitive as SQL for a lot of Experienced System/Data Analysts.

SQL support in Spark - Take 1: The Shark Story

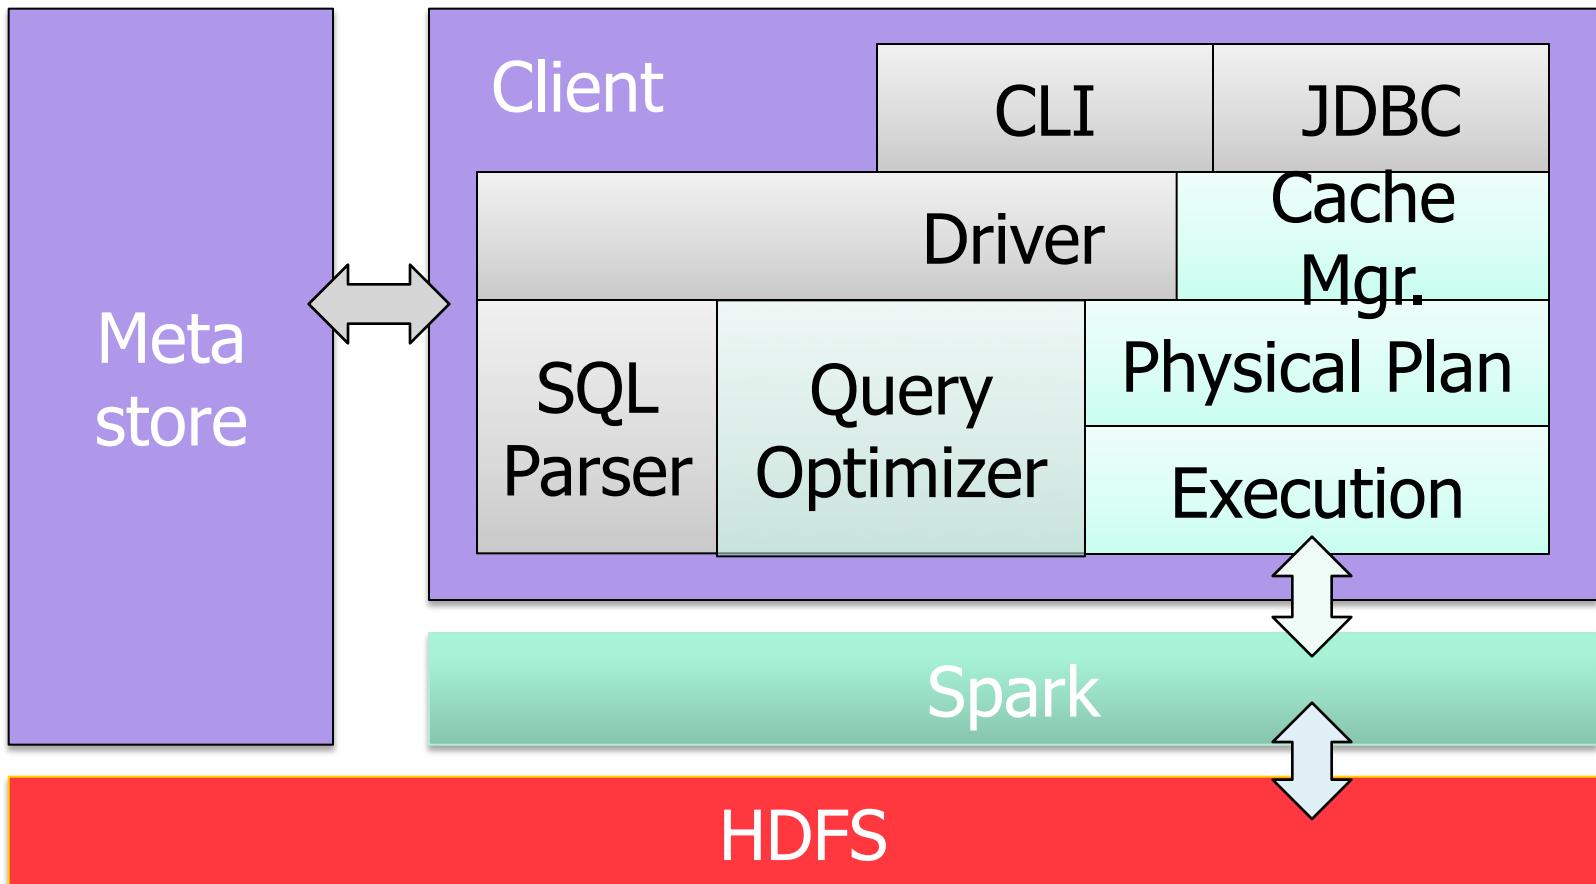
- Hive is great, but Hadoop's execution engine makes even the smallest queries take minutes
- Scala is good for programmers, but many data users only know SQL
- **Initial Approach: Make Hive to run on Spark**

SHARK = Hive on Spark

Original Hive Architecture



Shark Architecture



[Engle et al, SIGMOD 2012]

Efficient In-Memory Storage

- Simply caching Hive records as Java objects is inefficient due to high per-object overhead
- Instead, Shark employs column-oriented storage using **arrays of primitive types**

Row Storage

1	john	
2	mike	
3	sally	

Column Storage

1	2	3
john	mike	sally

Efficient In-Memory Storage

- Simply caching Hive records as Java objects is inefficient due to high per-object overhead
- Instead, Shark employs column-oriented storage using **arrays of primitive types**

Row Storage

Column Storage

Benefit: similarly compact size to serialized data,
but >5x faster to access



But Shark was short-lived (2011-2014)

Limitations of **SHARK**

- Can only be used to query external data in Hive catalog → limited data sources
- Can only be invoked via SQL string from Spark
→ error prone
- Hive optimizer tailored for MapReduce
→ difficult to extend
- As a result, BDAS Project decided to switch to Spark SQL and stopped development of Shark in 2014
 - The Apache Hive community still runs a Hive-over-Spark effort, as well as the Stinger/ Stinger.Next efforts to make Hive/HiveQL to be SQL compatible and low-latency

Take 2: Spark SQL Overview

- Part of the core distribution since Spark 1.0 (April 2014)
 - Optionally alongside or replacing existing Hive deployments
- Run SQL/ HiveQL queries including UDFs, UDAFs and SerDes, e.g.



```
SELECT COUNT(*)  
FROM hiveTable  
WHERE hive_udf(data)
```

- Connect existing Business Intelligence (BI) tools to Spark through JDBC



- Bindings in Python, Scala and Java



The Approach of Spark SQL

- Introduce a Tightly Integrated way to work with a new abstraction of Structured Data called **SchemaRDD**, which is a Distributed Collection of Rows (i.e. a Table) with Named Columns
 - SchemaRDD was renamed to **DataFrame** in Spark 1.3
- Support the Transformation of RDDs using SQL: In particular, **DataFrames** (aka SchemaRDDs) is an abstraction which supports:
 - Selecting, Filtering, Aggregating and Plotting Structured data (cf. R or Python-based Pandas)
- Evaluated lazily → unmaterialized *logical* plan
- Data source integration Support for: Hive, Parquet, JSON and ...

Relationship between Spark SQL and Shark

- Shark modified the Hive backend to run over Spark but had two challenges:
 - Limited integration with Spark programs
 - Hive Optimizer not designed for Spark
- Spark SQL reuses some parts of Shark by Borrowing:
 - Hive Data Loading
 - In-memory Column-store
- while Adding:
 - RDD-aware Optimizer
 - Richer Language Interfaces

What is an RDD ?

- Dependencies
- Partitions (with optional locality information)
- Compute Function: `Partition=>Iterator[T]`

Opaque Computation

What is an RDD ?

- Dependencies
- Partitions (with optional locality information)
- Compute Function: Partition=>Iterator[T]

[T]
Opaque Data

Opaque Data

Why Structure ?

- What do we mean by “Structure” [verb] :
 - Construct or Arrange according to a plan ; Give a pattern or organization to.
- By definition, structure will LIMIT what can be expressed.
- In practice, it is still possible to accommodate a vast majority of computations

BUT

- By Limiting the space of what can be expressed
ENABLES Optimization

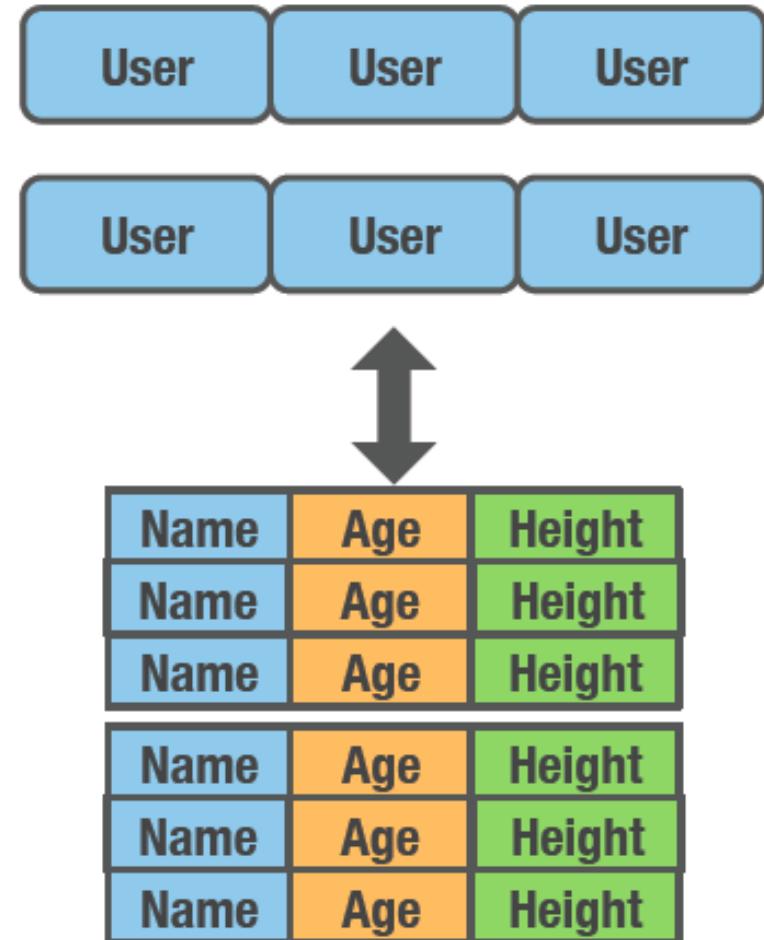
Adding Schema to RDDs

Spark + RDDs

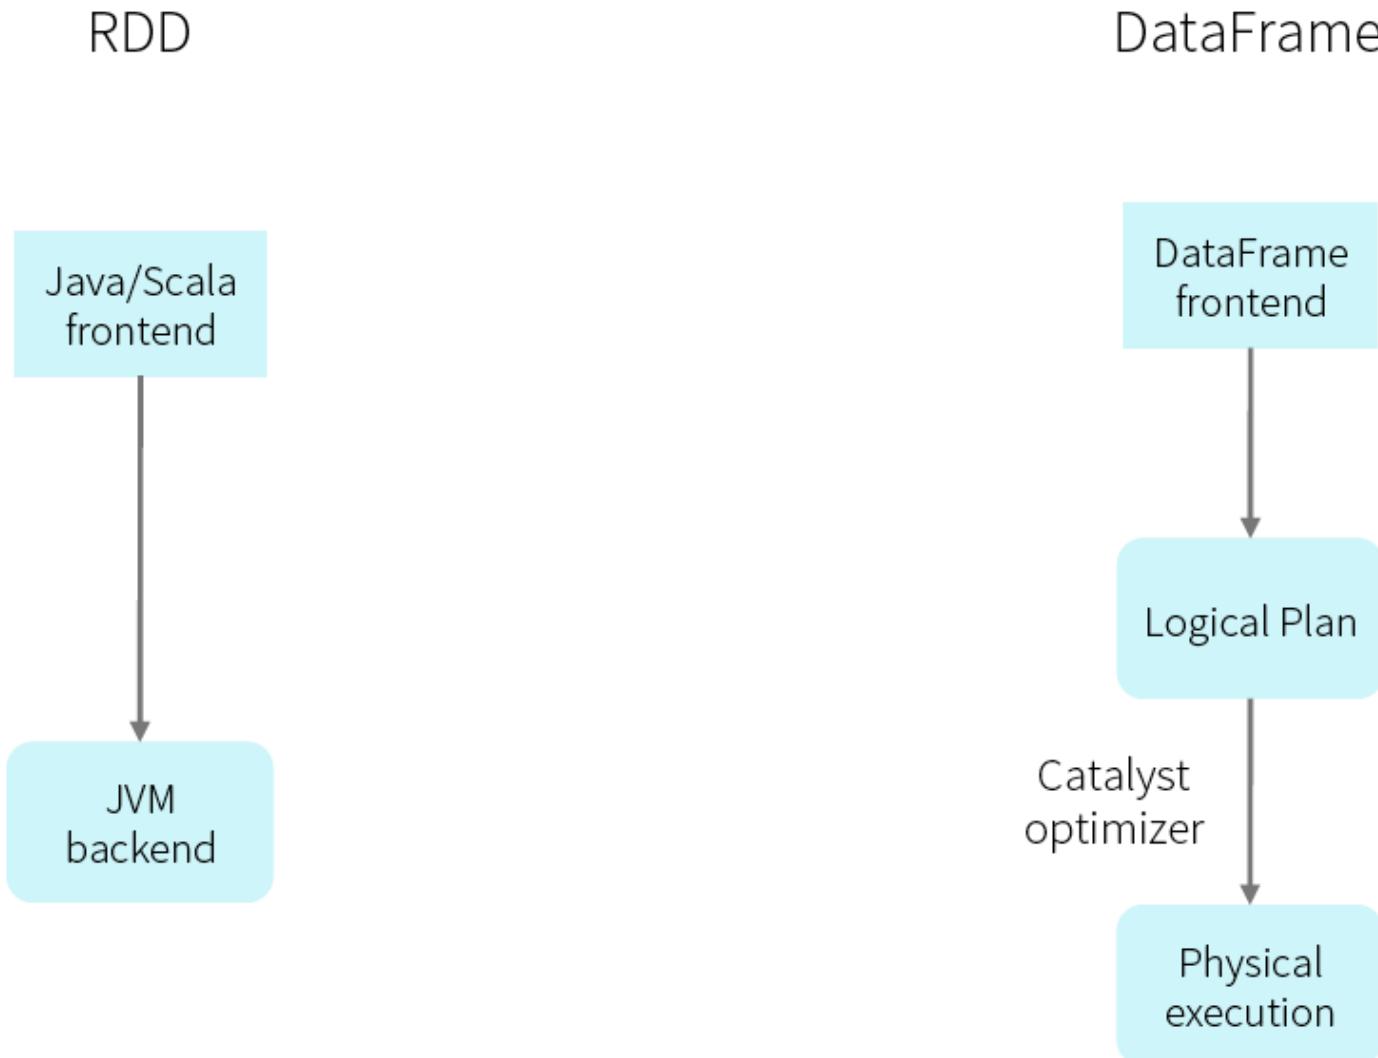
- **Functional** transformations
on Partitioned Collections
of *Opaque Objects*

SQL + DataFrames (aka SchemaRDDs)

- **Declarative** transformations
on Partitioned Collections
of *Tuples*



Comparing the Approaches of RDD vs. DataFrame



Data Model for DataFrame

- Nested data model
- Supports both primitive SQL types (boolean, integer, double, decimal, string, date, timestamp) and complex types (structs, arrays, maps, and unions); also user defined types.
- First class support for complex data types

DataFrame Operations

- Relational operations (select, where, join, groupBy) via a DSL
- Operators take *expression* objects
- Operators build up an abstract syntax tree (AST), which is then optimized by *Catalyst*.

```
employees
    .join(dept, employees("deptId") === dept("id"))
    .where(employees("gender") === "female")
    .groupBy(dept("id"), dept("name"))
    .agg(count("name"))
```

- Alternatively, register as temp SQL table and perform traditional SQL query strings

```
users.where(users("age") < 21)
        .registerTempTable("young")
ctx.sql("SELECT count(*), avg(age) FROM young")
```

Programming Interface for Spark SQL

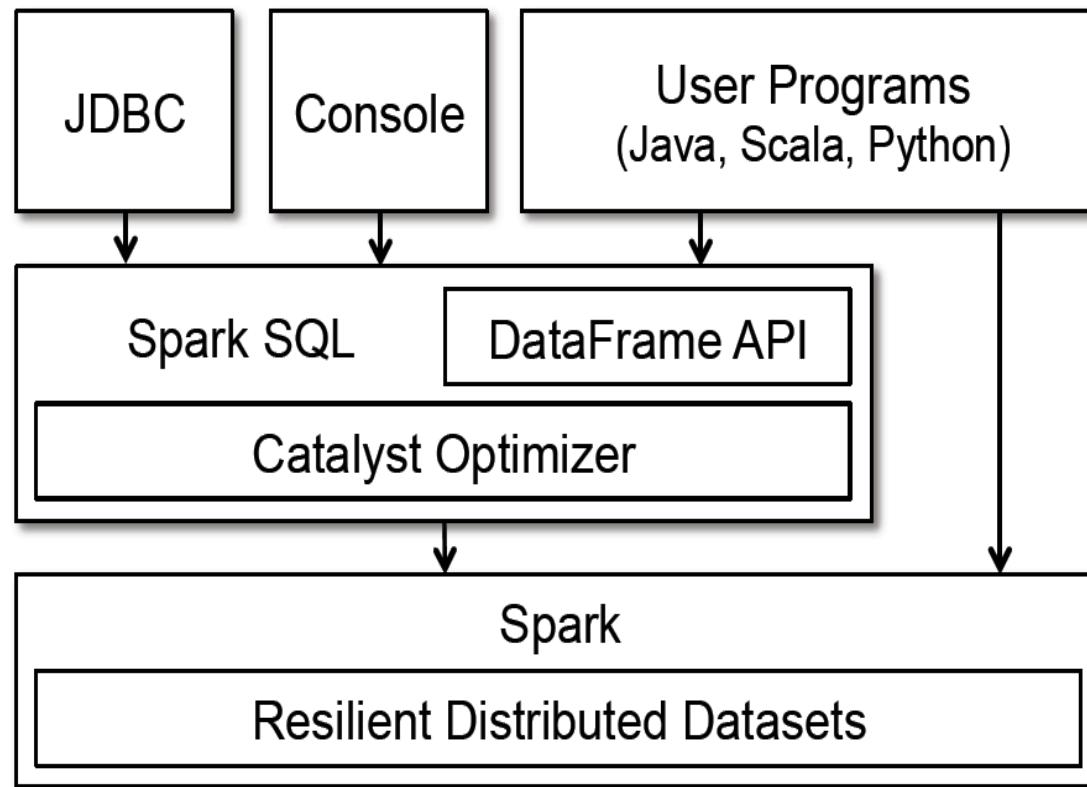
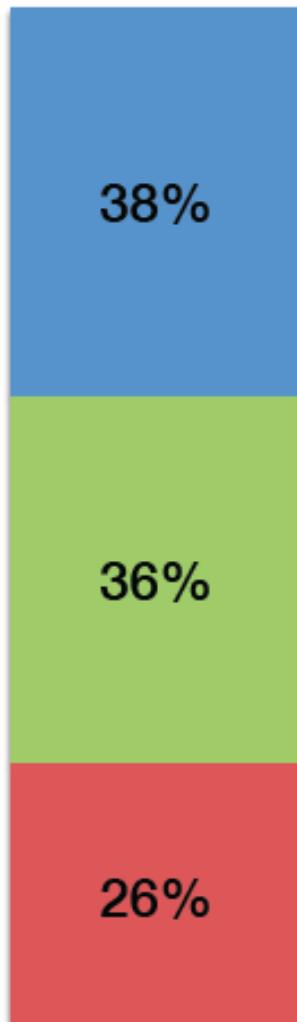


Figure 1: Interfaces to Spark SQL, and interaction with Spark.

Spark SQL Components



■ Catalyst Optimizer

- Relational algebra + expressions
- Query optimization

■ Spark SQL Core

- Execution of queries as RDDs
- Reading in Parquet, JSON ...

■ Hive Support

- HQL, MetaStore, SerDes, UDFs

Getting Started: Spark SQL

- SQLContext / HiveContext
 - Entry point for all SQL functionality
 - Wraps/Extend existing Spark Context

```
from pyspark.sql import SQLContext  
sqlCtx = SQLContext(sc)
```

OR

```
ctx = new HiveContext()  
users = ctx.table("users")  
young = users.where(users("age") < 21)  
println(young.count())
```

Sample Input Data

- A text file filled with people's names and ages:

Michael, 30

Andy, 31

...

RDDs as Relations (Scala)

```
val sqlContext = new org.apache.spark.sql.SQLContext(sc)
import sqlContext._

// Define the schema using a case class.
case class Person(name: String, age: Int)
// Create an RDD of Person objects and register it as a table.
val people =
  sc.textFile("examples/src/main/resources/people.txt")
    .map(_.split(","))
    .map(p => Person(p(0), p(1).trim.toInt))

people.registerAsTable("people")
```

RDDs as Relations (Python)

```
# Load a text file and convert each line to a dictionary.  
lines = sc.textFile("examples/.../people.txt")  
  
parts = lines.map(lambda l: l.split(","))  
people = parts.map(lambda p: Row(name=p[0],age=int(p[1])))  
  
# Infer the schema, and register the SchemaRDD as a table  
peopleTable = sqlctx.inferSchema(people)  
peopleTable.registerAsTable("people")
```

RDDs as Relations (Java)

```
public class Person implements Serializable {  
    private String _name;  
    private int _age;  
    public String getName() { return _name; }  
    public void setName(String name) { _name = name; }  
    public int getAge() { return _age; }  
    public void setAge(int age) { _age = age; }  
}  
  
JavaSQLContext ctx = new org.apache.spark.sql.api.java.JavaSQLContext(sc)  
JavaRDD<Person> people = ctx.textFile("examples/src/main/resources/  
people.txt").map(  
    new Function<String, Person>() {  
        public Person call(String line) throws Exception {  
            String[] parts = line.split(",");  
            Person person = new Person();  
            person.setName(parts[0]);  
            person.setAge(Integer.parseInt(parts[1].trim()));  
            return person;  
        }  
    });  
  
JavaSchemaRDD schemaPeople = sqlCtx.applySchema(people, Person.class);
```

Querying using Spark SQL (Python)

```
# SQL can be run over SchemaRDDs that have been registered  
# as a table.  
teenagers = sqlCtx.sql("""  
    SELECT name FROM people WHERE age >= 13 AND age <= 19""")  
  
# The results of SQL queries are RDDs and support all the normal  
# RDD operations.  
teenNames = teenagers.map(lambda p: "Name: " + p.name)
```

Support of Existing Tools, and New Data Sources

- SparkSQL includes a server that exposes its data using JDBC/ODBC
 - Query data from HDFS/S3
 - Including formats like Hive/Parquet/JSON
 - Support for caching data IN-MEMORY

Caching Tables In-Memory

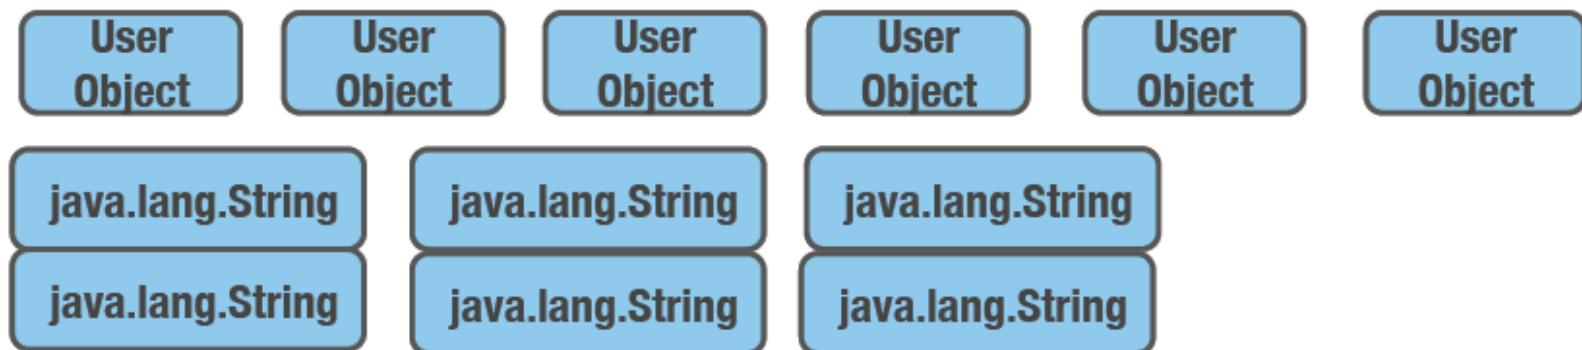
- SparkSQL can cache tables using an in-memory columnar format:
 - Scan only required columns
 - Fewer allocated objects (less Garbage Collection)
 - Automatically selects best compression

e.g.

```
cacheTable("people") or dataframe.cache()
```

Caching Comparison

Spark MEMORY_ONLY Caching



SchemaRDD Columnar Caching

ByteBuffer

Name	Name
Name	Name
Name	Name

ByteBuffer

Age	Age
Age	Age
Age	Age

ByteBuffer

Height	Height
Height	Height
Height	Height

Language Integrated UDFs

```
registerFunction("countMatches",
  lambda (pattern, text):
    re.subn(pattern, '', text)[1])

sql("SELECT countMatches('a', text)...")
```

Reading Data stored in Hive

```
from pyspark.sql import HiveContext
hiveCtx = HiveContext(sc)

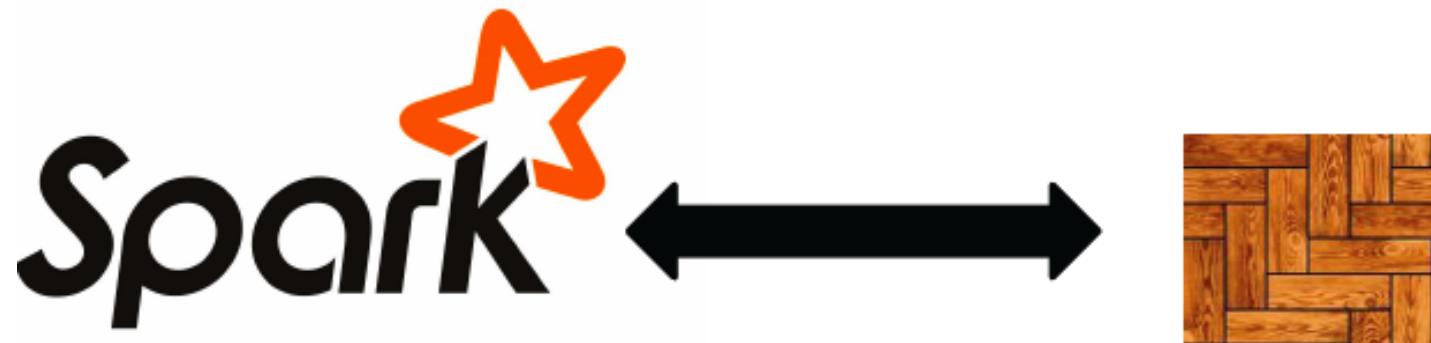
hiveCtx.hql("""
    CREATE TABLE IF NOT EXISTS src (key INT, value STRING)""")

hiveCtx.hql("""
    LOAD DATA LOCAL INPATH 'examples/.../kv1.txt' INTO TABLE src""")

# Queries can be expressed in HiveQL.
results = hiveCtx.hql("FROM src SELECT key, value").collect()
```

Parquet Compatibility

- Native support for reading data in Parquet
 - Columnar storage avoids reading unneeded data
 - RDDs can be written to Parquet files, preserving the schema
 - Convert other slower formats into Parquet for repeated querying



Using Parquet

```
# SchemaRDDs can be saved as Parquet files, maintaining the
# schema information.
peopleTable.saveAsParquetFile("people.parquet")

# Read in the Parquet file created above. Parquet files are
# self-describing so the schema is preserved. The result of
# Loading a parquet file is also a SchemaRDD.
parquetFile = sqlCtx.parquetFile("people.parquet")

# Parquet files can be registered as tables used in SQL.
parquetFile.registerAsTable("parquetFile")
teenagers = sqlCtx.sql("""
    SELECT name FROM parquetFile WHERE age >= 13 AND age <= 19""")
```

JSON Support

- Use jsonFile or jsonRDD to convert a collection of JSON objects into a DataFrame
- Infer and Union the schema of each record
- Maintain nested structures and arrays

JSON Example

```
# Create a SchemaRDD from the file(s) pointed to by path
people = sqlContext.jsonFile(path)

# Visualized inferred schema with printSchema().
people.printSchema()
# root
#   |-- age: integer
#   |-- name: string

# Register this SchemaRDD as a table.
people.registerTempTable("people")
```

Data Sources API

- Allow easy integration with new sources of structured data:

```
CREATE TEMPORARY TABLE episodes
USING com.databricks.spark.avro
OPTIONS (
    path "./episodes.avro"
)
```

<https://github.com/databricks/spark-avro>

Much More than SQL: DataFrames as A Unified Interface for the Processing of Structured Data



Much More than SQL: Simplifying Inputs and Outputs

Spark SQL's Data Source API can read and write
DataFrames using a variety of formats.

Built-In

{ JSON }



JDBC

Parquet



MySQL



PostgreSQL



Amazon S3

H2

External



dBase

APACHE
HBASE

elasticsearch.



amazon
web services
Amazon Redshift

and more...

Unified and Simplified Interface to Read/ Write Data in Many different Formats

```
df = sqlContext.read \  
  .format("json") \  
  .option("samplingRatio", "0.1") \  
  .load("/home/michael/data.json")
```

```
df.write \  
  .format("parquet") \  
  .mode("append") \  
  .partitionBy("year") \  
  .saveAsTable("fasterData")
```

read and write functions create new builders for doing I/O

Unified and Simplified Interface to Read/ Write Data in Many different Formats

```
df = sqlContext.read \
    .format("json") \
    .option("samplingRatio", "0.1") \
    .load("/home/michael/data.json")
```

}

```
df.write \
    .format("parquet") \
    .mode("append") \
    .partitionBy("year") \
    .saveAsTable("fasterData")
```

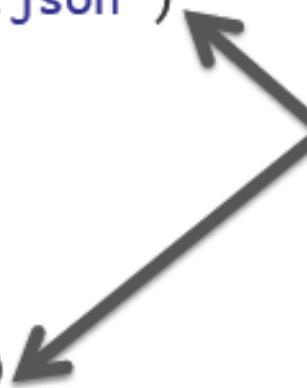
}

- Builder methods are used to specify:
- Format
 - Partitioning
 - Handling of existing data
 - and more

Unified and Simplified Interface to Read/ Write Data in Many different Formats

```
df = sqlContext.read \
    .format("json") \
    .option("samplingRatio", "0.1") \
    .load("/home/michael/data.json")
```

```
df.write \
    .format("parquet") \
    .mode("append") \
    .partitionBy("year") \
    .saveAsTable("fasterData")
```



load(...), save(...) or
saveAsTable(...)
functions create
new builders for
doing I/O

ETL using Custom Data Sources

```
sqlContext.read
  .format("com.databricks.spark.jira")
  .option("url", "https://issues.apache.org/jira/rest/api/latest/search")
  .option("user", "marmbrus")
  .option("password", "*****")
  .option("query", """
    |project = SPARK AND
    |component = SQL AND
    |(status = Open OR status = "In Progress" OR status = Reopened)""".stripMargin)
  .load()
  .repartition(1)
  .write
  .format("parquet")
  .saveAsTable("sparkSqlJira")
```

Write Less Codes with DataFrames

- Common operations can be expressed concisely as higher level operation calls to the DataFrame API:
 - Selecting required Columns
 - Joining Different Data Sources
 - Aggregation (Count, Sum, Average, etc)
 - Filtering

Write Less Codes: An Example of Computing Average



```
private IntWritable one =
  new IntWritable(1)
private IntWritable output =
  new IntWritable()
protected void map(
  LongWritable key,
  Text value,
  Context context) {
  String[] fields = value.split("\t")
  output.set(Integer.parseInt(fields[1]))
  context.write(one, output)
}

IntWritable one = new IntWritable(1)
DoubleWritable average = new DoubleWritable()

protected void reduce(
  IntWritable key,
  Iterable<IntWritable> values,
  Context context) {
  int sum = 0
  int count = 0
  for(IntWritable value : values) {
    sum += value.get()
    count++
  }
  average.set(sum / (double) count)
  context.write(key, average)
}

data = sc.textFile(...).split("\t")
data.map(lambda x: (x[0], [x[1], 1])) \
  .reduceByKey(lambda x, y: [x[0] + y[0], x[1] + y[1]]) \
  .map(lambda x: [x[0], x[1][0] / x[1][1]]) \
  .collect()
```

Write Less Code: Example of Computing Average

■ Using RDDs

```
■ data = sc.textFile(...).split("\t")
■ data.map(lambda x: (x[0], [int(x[1]), 1])) \
■     .reduceByKey(lambda x, y: [x[0] + y[0], x[1] + y[1]]) \
■     .map(lambda x: [x[0], x[1][0] / x[1][1]]) \
■     .collect()
```

Using SQL

```
SELECT name, avg(age)
FROM people
GROUP BY name
```

Using Pig

```
P = load '/people' as (name, name);
G = group P by name;
R = foreach G generate ... AVG(G.age);
```

Using DataFrames

```
sqlCtx.table("people") \
    .groupBy("name") \
    .agg("name", avg("age")) \
    .collect()
```

Read Less Data with DataFrames & SparkSQL

“The fastest way to process big data is to never read it.”

- SparkSQL can help the program to read less data automatically by performing BEYOND naïve scanning:
 - Using Columnar formats (e.g. Parquet) and prune irrelevant Columns and Blocks of data
 - Push filters to the source
 - Converting to more efficient formats, e.g. turning string comparisons into integer comparisons for dictionary encoded data
 - Using Partitioning (i.e., /year=2-14/month=02/..)
 - Skipping data using statistics (i.e. min, max)
 - Pushing predicates into storage systems (i.e. JDBC)

Intermix DataFrame Operations with Custom Codes (Python, Java, R, Scala)

```
zipToCity = udf(lambda zipCode: <custom logic here>)
```

```
def add_demographics(events):
    u = sqlCtx.table("users")
    events \
        .join(u, events.user_id == u.user_id) \
        .withColumn("city", zipToCity(df.zip))
```



Augments any DataFrame that contains user_id

Takes and returns a DataFrame

Integration with RDDs

- Internally, DataFrame execution is done with Spark RDDs
=> Easy Interoperation with outside sources and custom algorithms

External Input

```
def buildScan(  
    requiredColumns: Array[String],  
    filters: Array[Filter]): RDD[Row]
```

Custom Processing

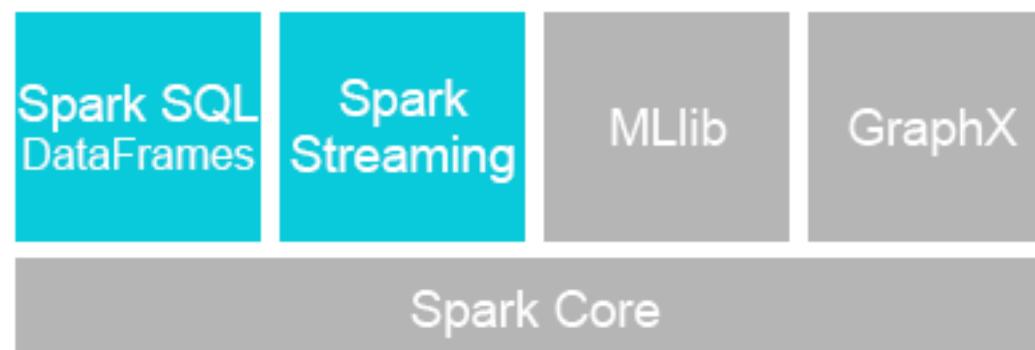
```
queryResult.rdd.mapPartitions { iter =>  
    ... Your code here ...  
}
```

Combine SQL with Streaming

- Interactively query streaming data with SQL and Dataframes

```
// Register each batch in stream as table
kafkaStream.foreachRDD { batchRDD =>
    batchRDD.toDF.registerTempTable("events")
}
```

```
// Interactively query table
sqlContext.sql("select * from events")
```



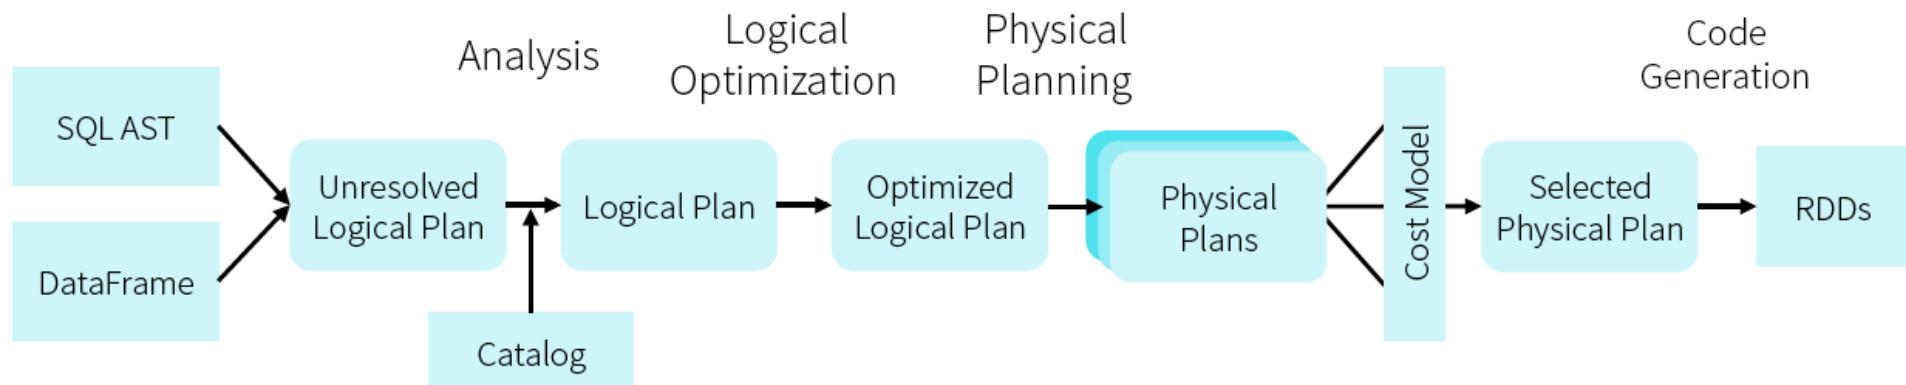
DataFrame & SparkSQL Demo

Demo:

- *Using Spark SQL to read, write, and transform data in a variety of formats:*

<http://people.apache.org/~marmbrus/talks/frame.demo.pdf>

Plan Optimization and Execution for the entire Pipelines

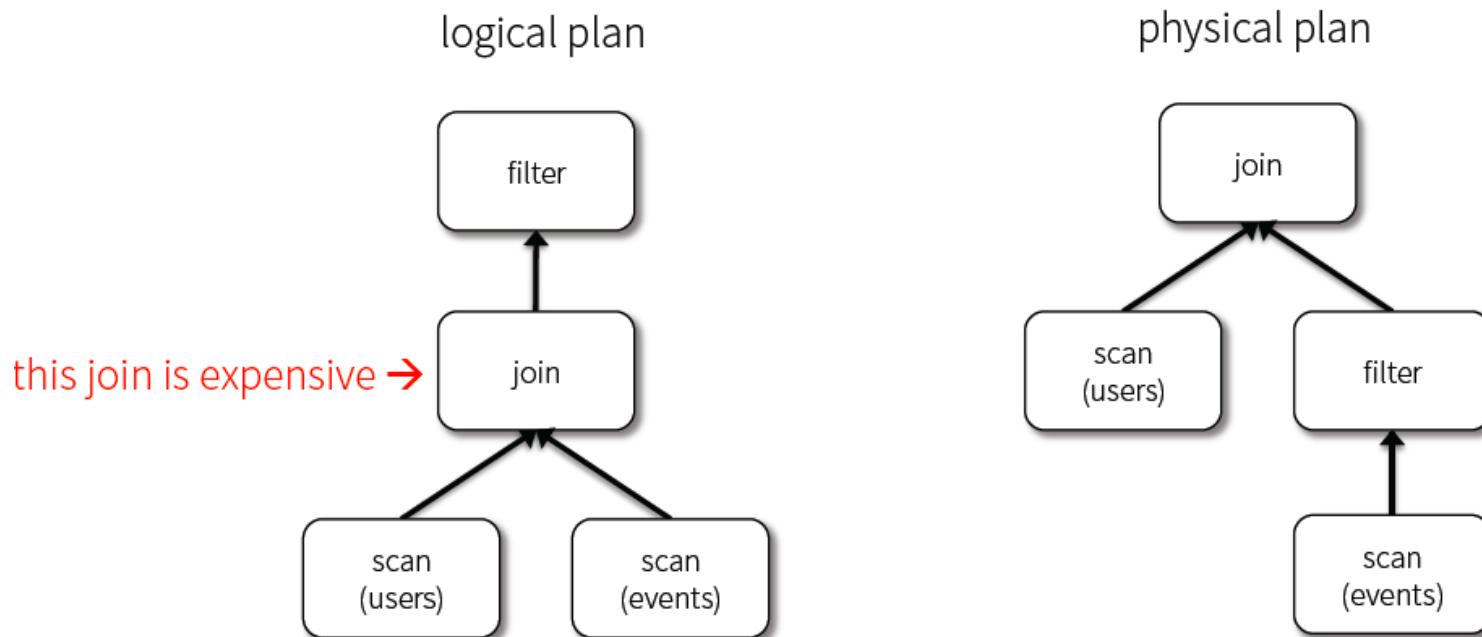


DataFrames and SQL share the same optimization/execution pipeline

- Optimization happens as late as possible
=> Spark SQL can optimize *even across different functions!*

Optimization Example

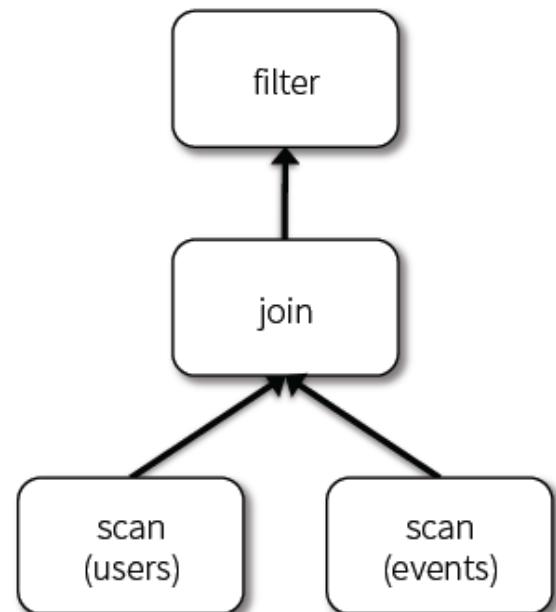
```
joined = users.join(events, users.id == events.uid)  
filtered = joined.filter(events.date >= "2015-01-01")
```



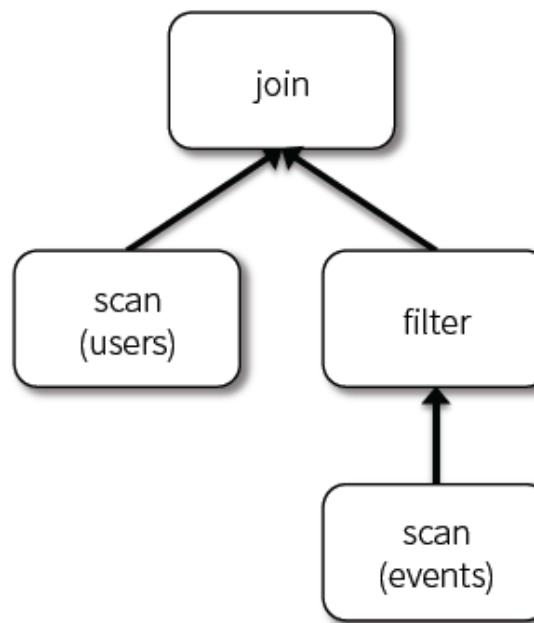
Optimization Example

```
joined = users.join(events, users.id == events.uid)  
filtered = joined.filter(events.date > "2015-01-01")
```

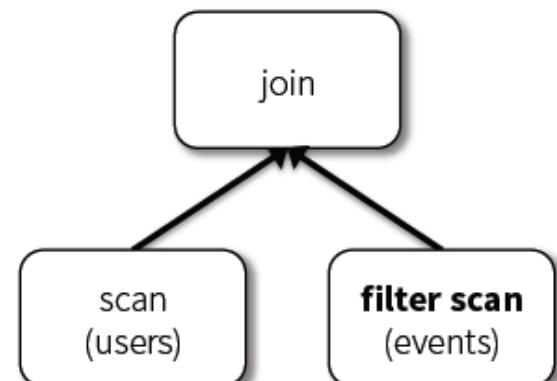
logical plan



optimized plan



optimized plan
with intelligent data sources

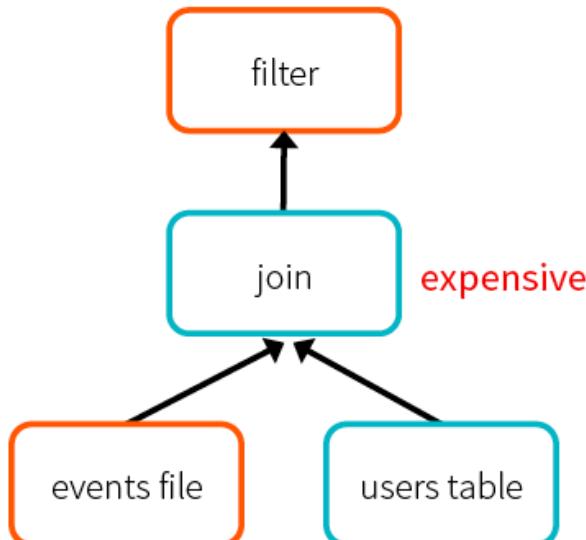


Optimization Example

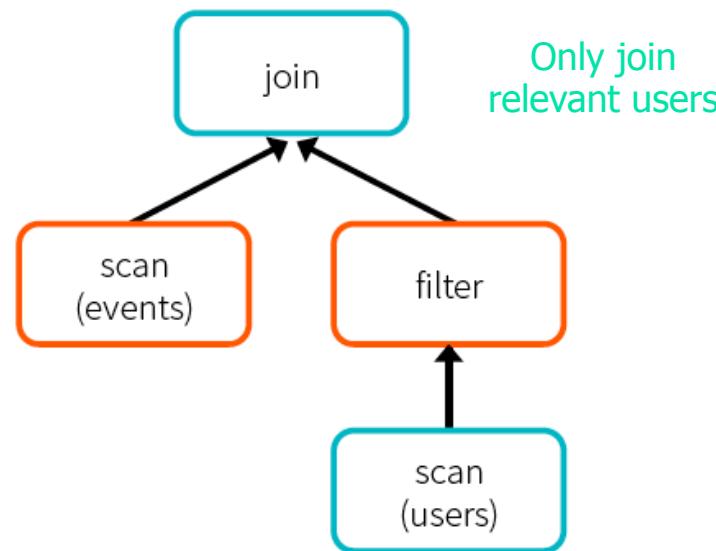
```
def add_demographics(events):
    u = sqlCtx.table("users")                      # Load Hive table
    events \
        .join(u, events.user_id == u.user_id) \      # Join on user_id
        .withColumn("city", zipToCity(df.zip))        # Run udf to add city column

events = add_demographics(sqlCtx.load("/data/events", "json"))
training_data = events.where(events.city == "New York").select(events.timestamp).collect()
```

Logical Plan



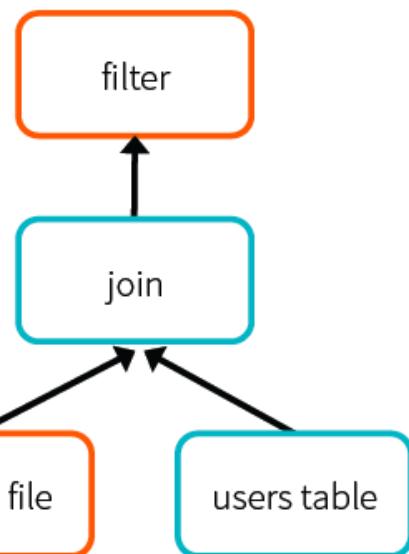
Physical Plan



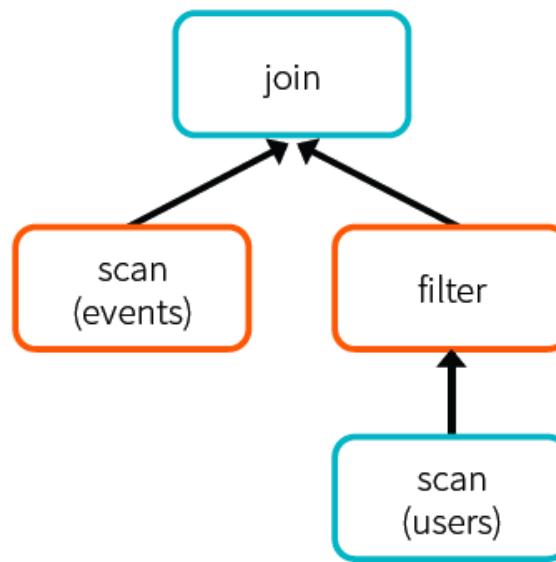
Optimization Example

```
def add_demographics(events):
    u = sqlCtx.table("users")                                # Load partitioned Hive table ←
    events \
        .join(u, events.user_id == u.user_id) \
        .withColumn("city", zipToCity(df.zip))               # Join on user_id
                                                               # Run udf to add city column
events = add_demographics(sqlCtx.load("/data/events", "parquet")) ←
training_data = events.where(events.city == "New York").select(events.timestamp).collect()
```

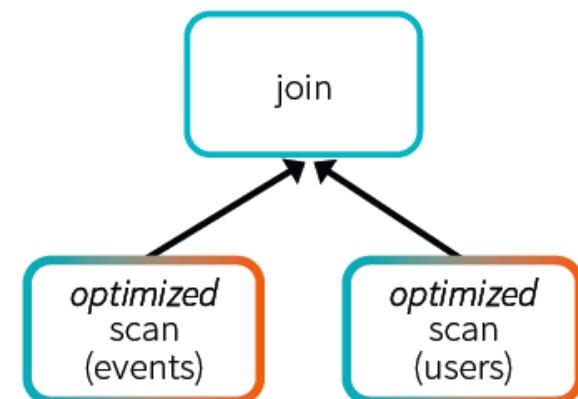
Logical Plan



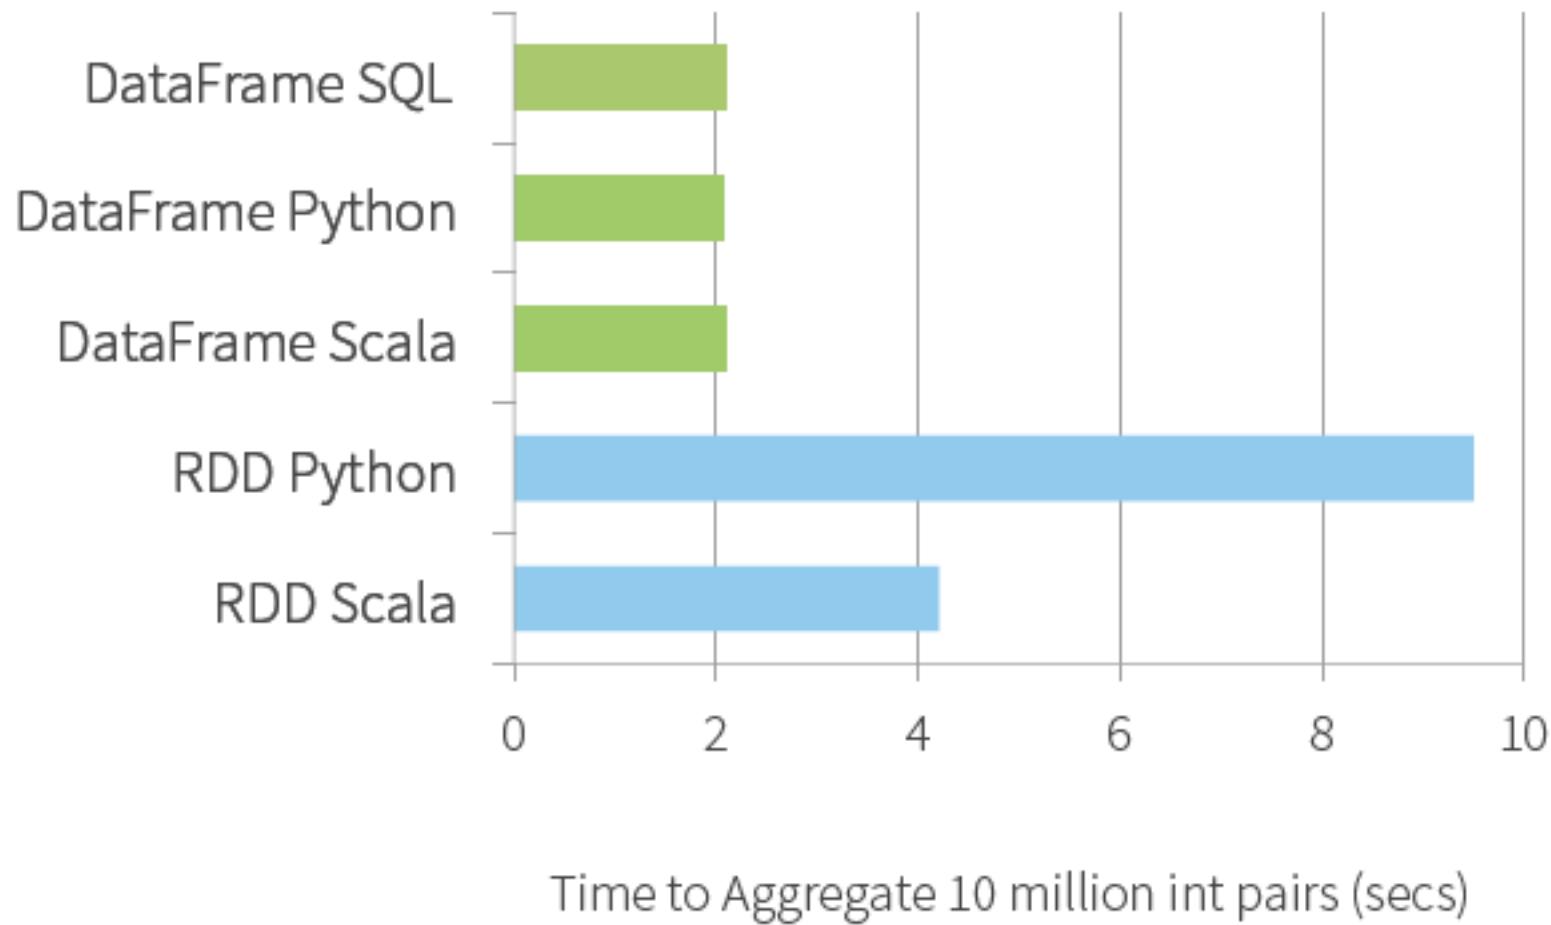
Physical Plan



Physical Plan
with Predicate Pushdown
and Column Pruning



Named Columns (vs. Opaque Objects in RDDs) Enable Performance Optimization



More Performance Comparison

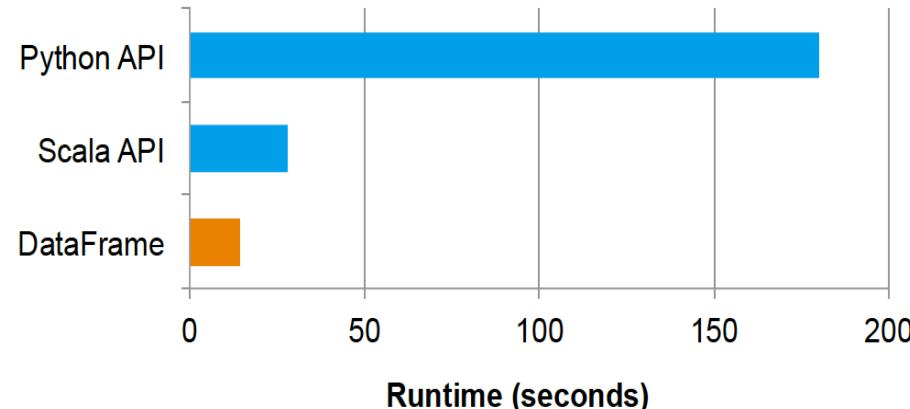


Figure 9: Performance of an aggregation written using the native Spark Python and Scala APIs versus the DataFrame API.

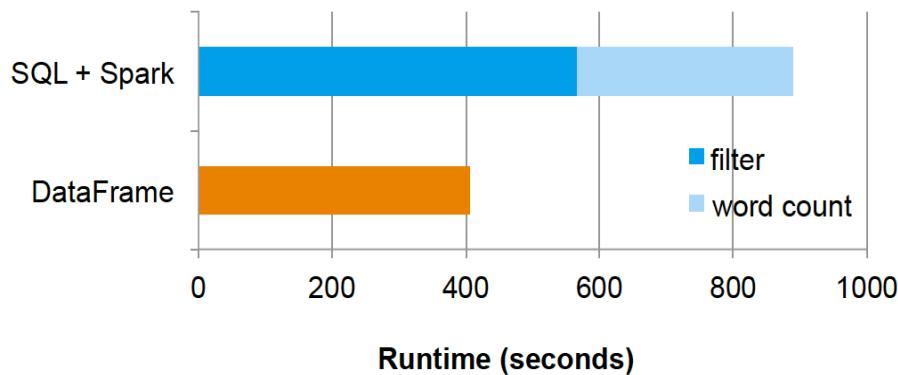
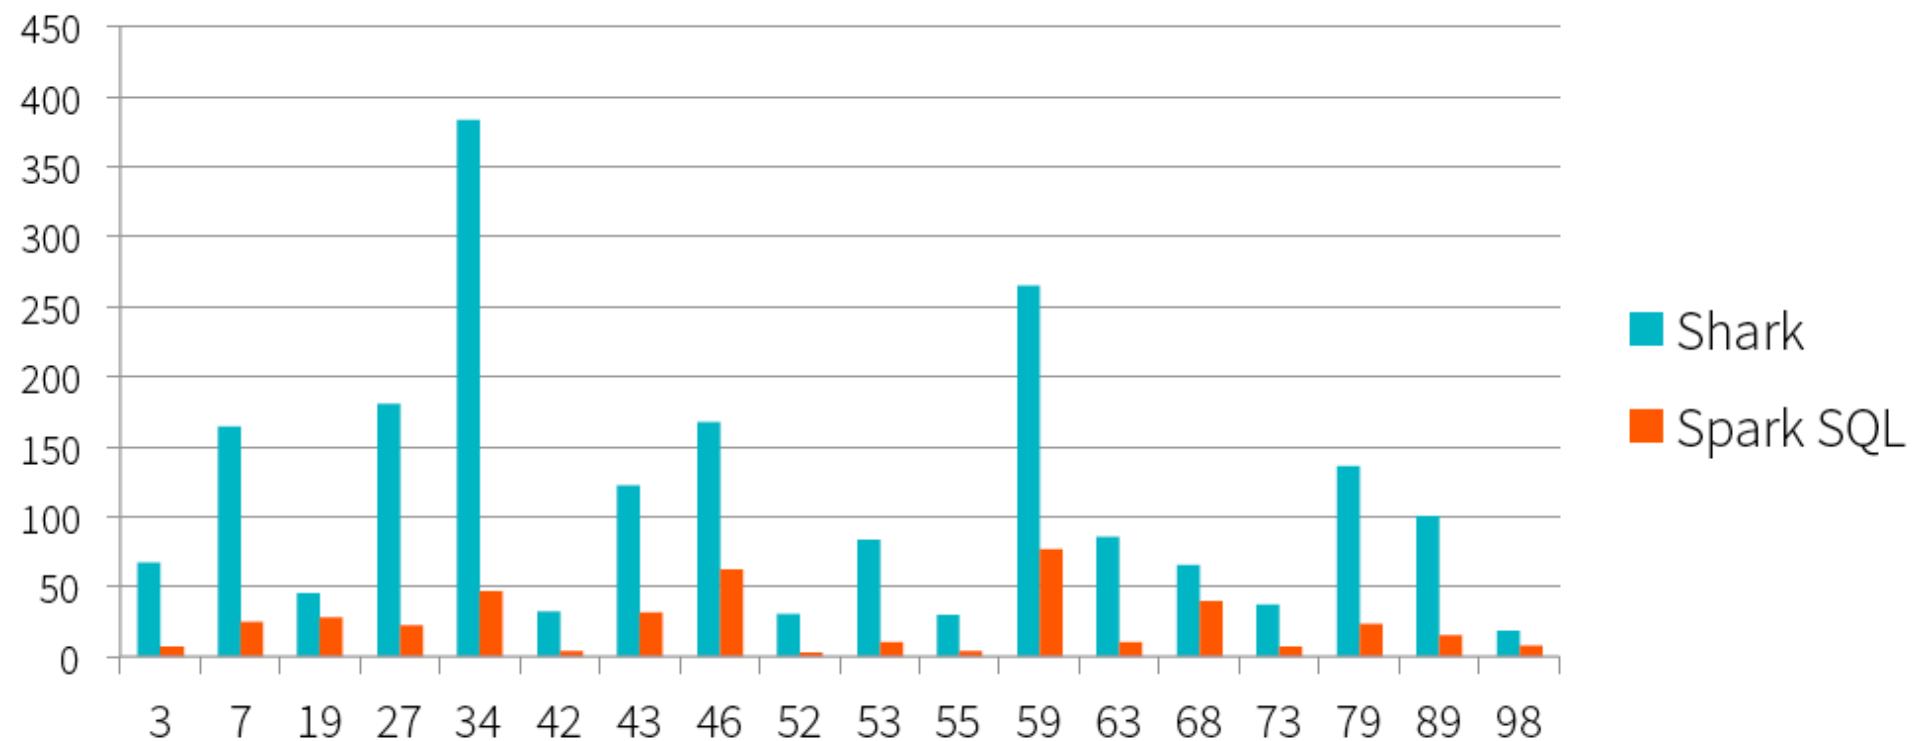


Figure 10: Performance of a two-stage pipeline written as a separate Spark SQL query and Spark job (above) and an integrated DataFrame job (below).

Performance Comparison: Spark SQL vs. Shark

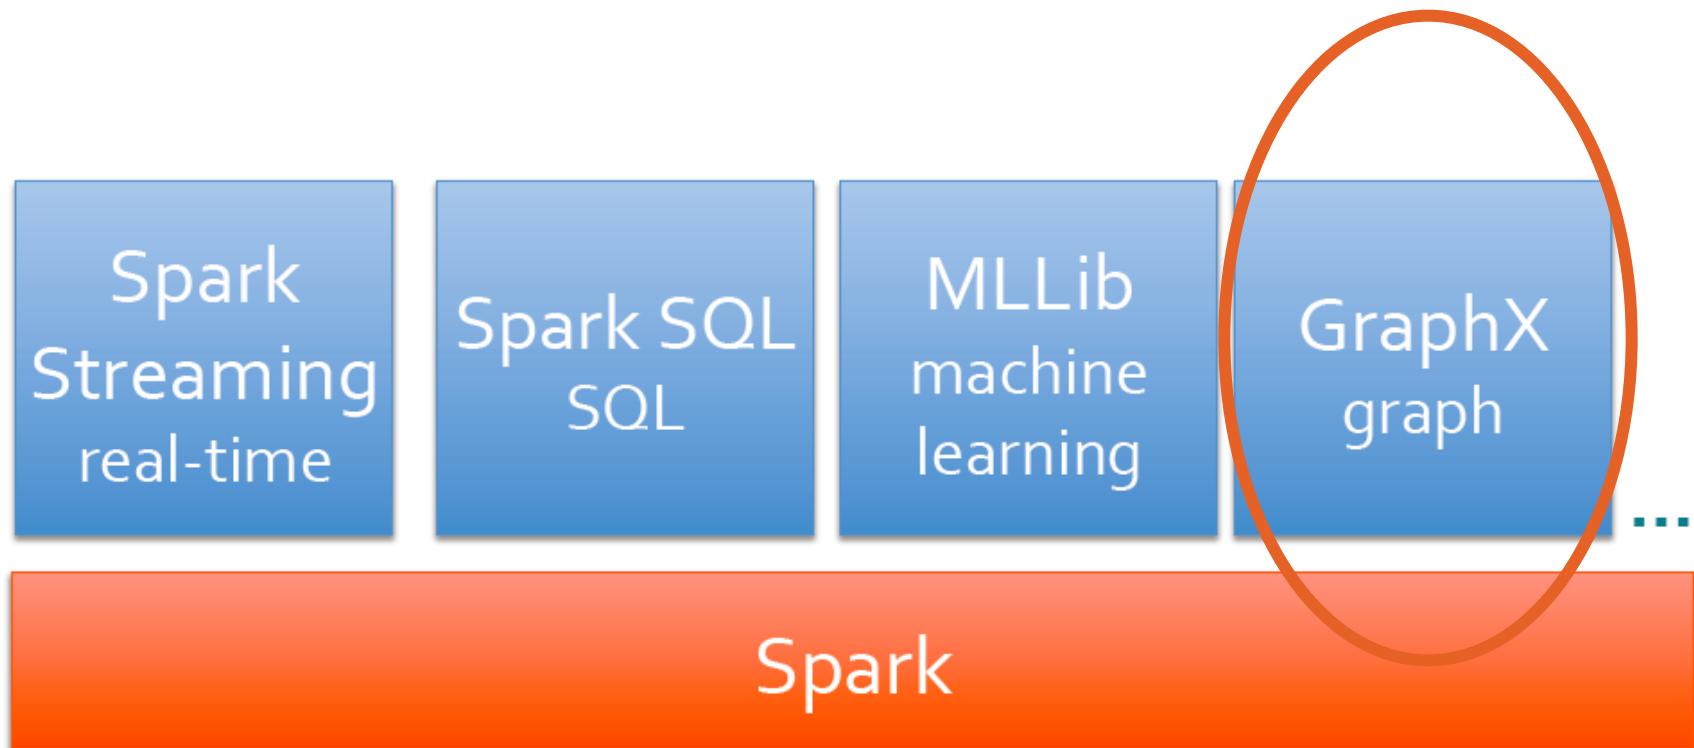
TPC-DS Performance



Lessons Learnt from Spark SQL

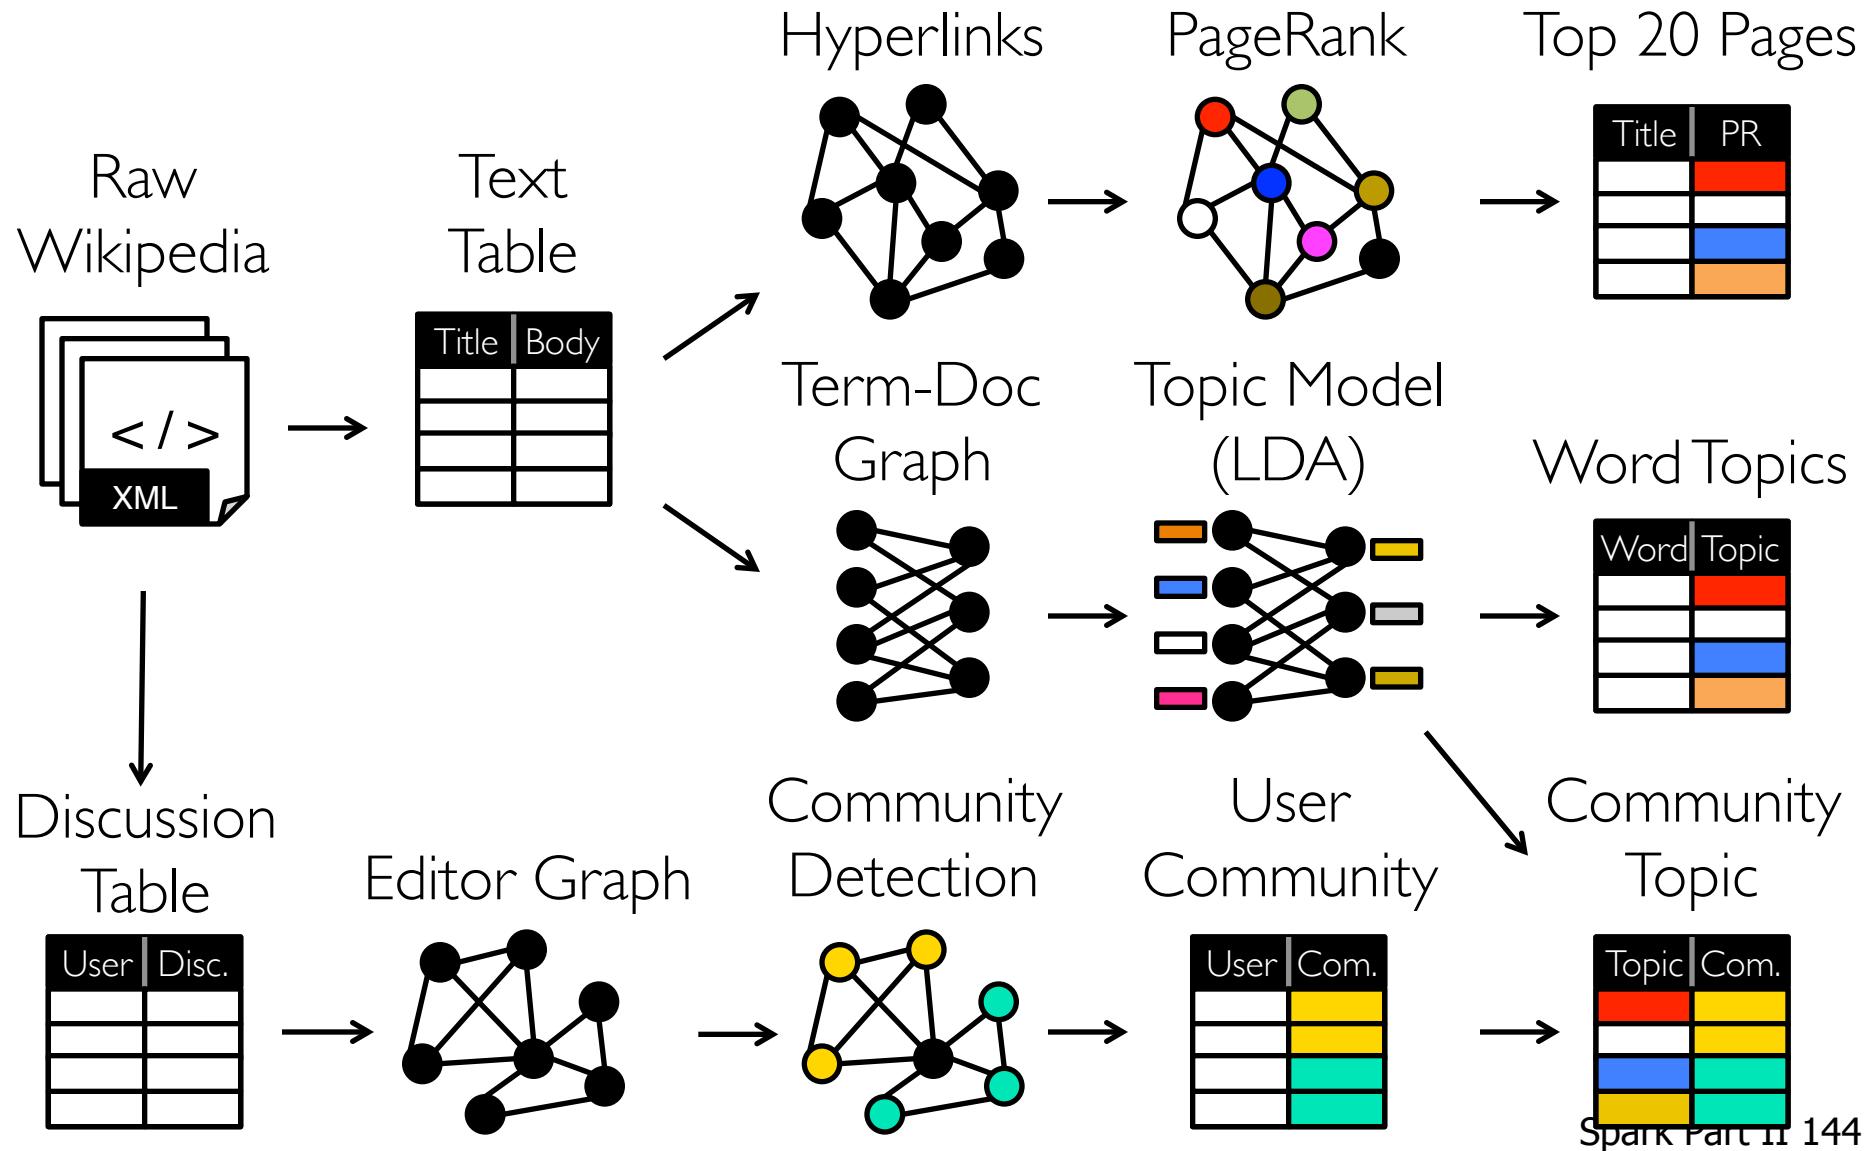
- SQL is wildly popular and important for real-world customers
- Schema is very useful
 - In most data pipelines, even the ones that start with unstructured data end up having some implicit structure
 - Key-value abstraction (under RDD) is too limited
 - Nevertheless, Support for Semi/Un-structured data is critical !
- Separation of Logical vs. Physical Plan is important for Performance Optimizations, e.g. join selection.

Generality of RDDs in Spark



A Short-version of GraphX: *Unifying Data-Parallel and Graph-Parallel Analytics*

Graphs are Central to Analytics



Tables and Graphs are composable views of the same physical data

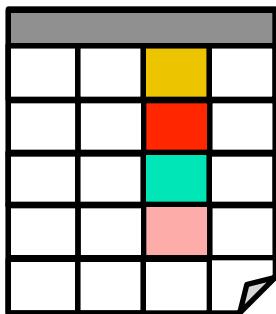
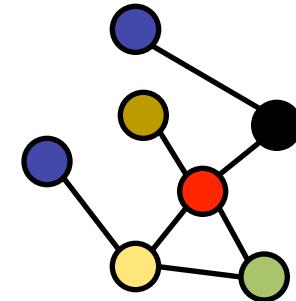
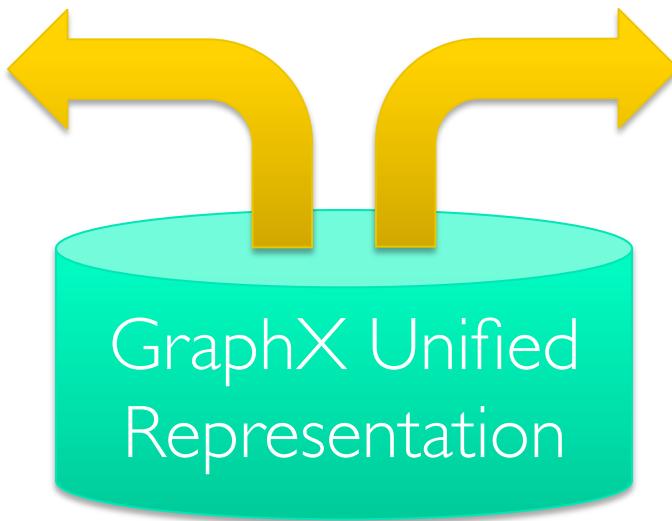


Table View



Graph View

Each view has its own operators that
exploit the semantics of the view
to achieve efficient execution

Separate Systems to Support Each View

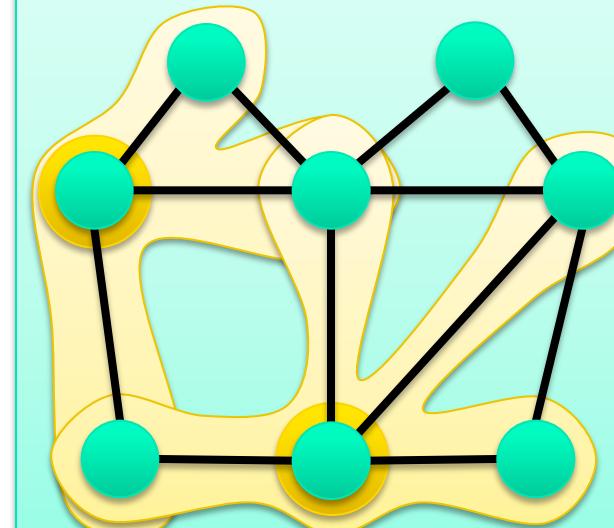
Table View



Graph View

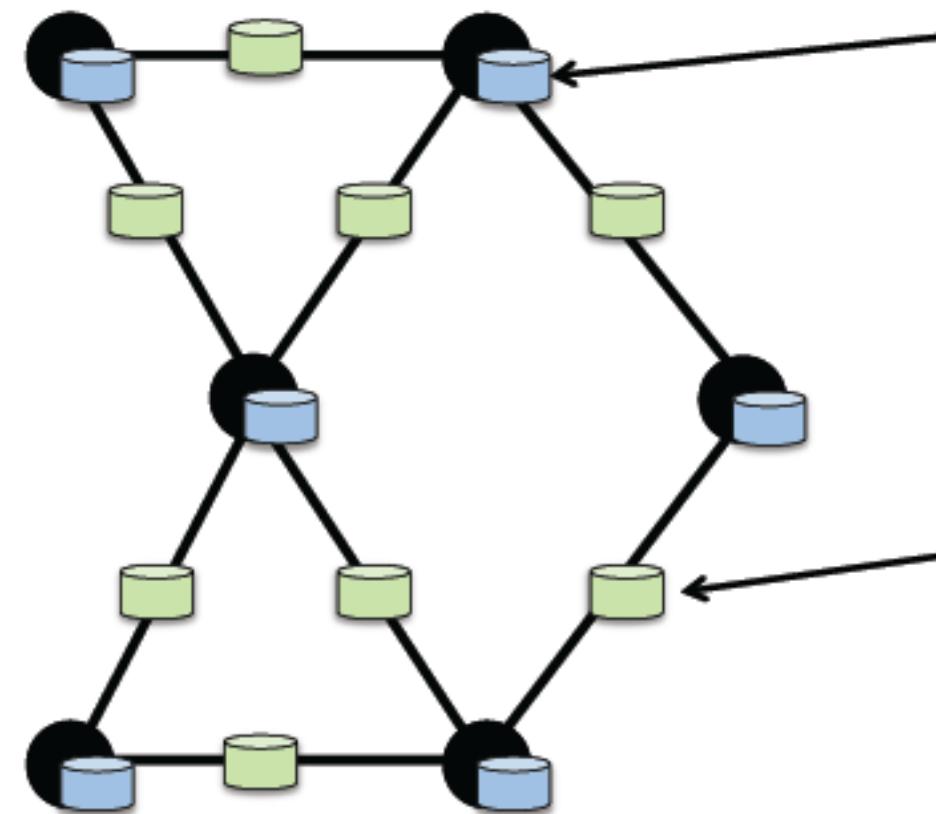


Dependency Graph



The GraphX API

Property Graphs



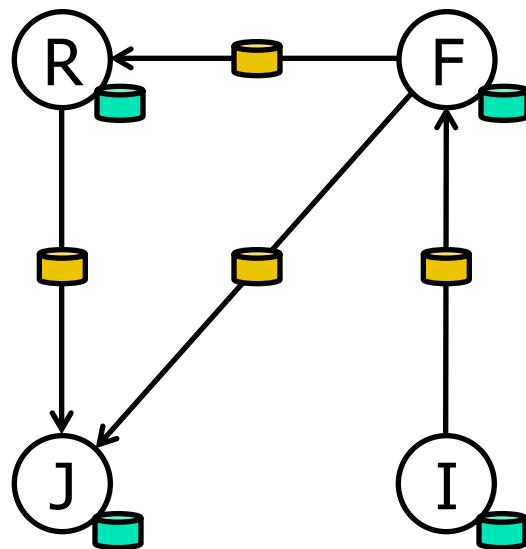
Vertex Property:

- User Profile
- Current PageRank Value

Edge Property:

- Weights
- Relationships
- Timestamps

Property Graph



View a Graph as a Table

Vertex Property Table

Id	Property (V)
Rxin	(Stu., Berk.)
Jegonzal	(PstDoc, Berk.)
Franklin	(Prof., Berk)
Istoica	(Prof., Berk)

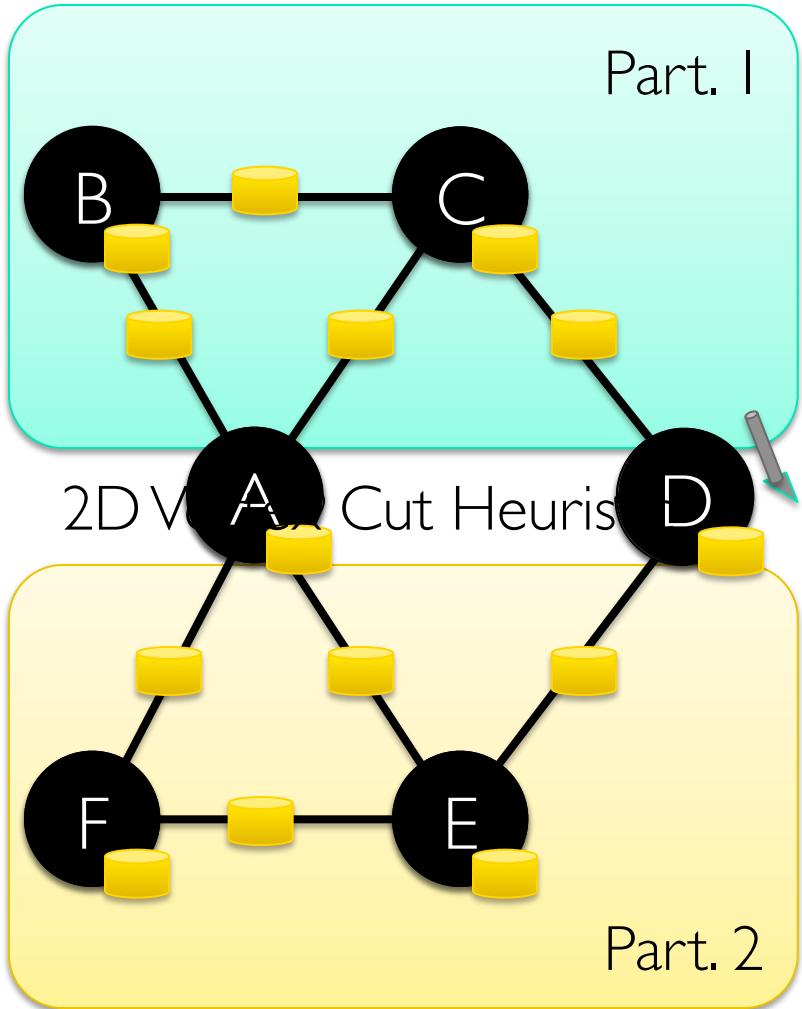
Edge Property Table

SrcId	DstId	Property (E)
rxin	jegonzal	Friend
franklin	rxin	Advisor
istoica	franklin	Coworker
franklin	jegonzal	PI

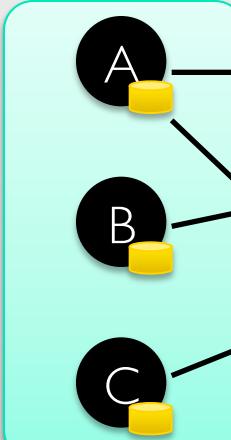
GraphX System Design

Distributed Graphs as Tables (RDDs)

Property Graph



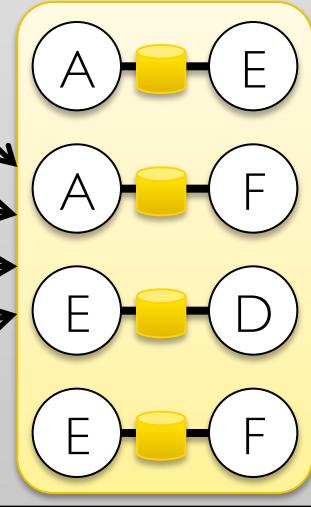
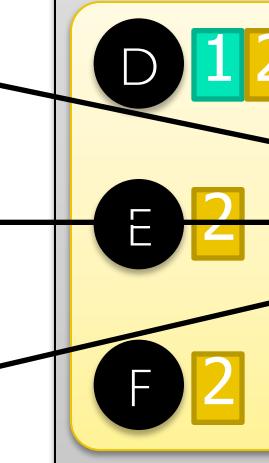
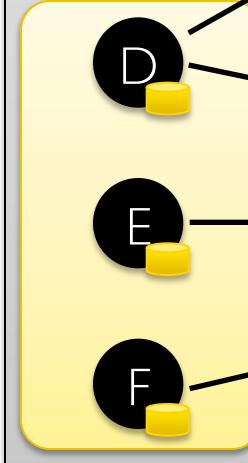
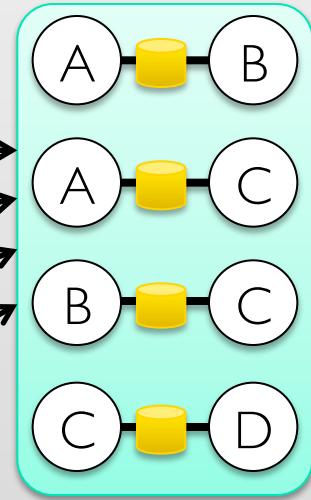
Vertex Table (RDD)



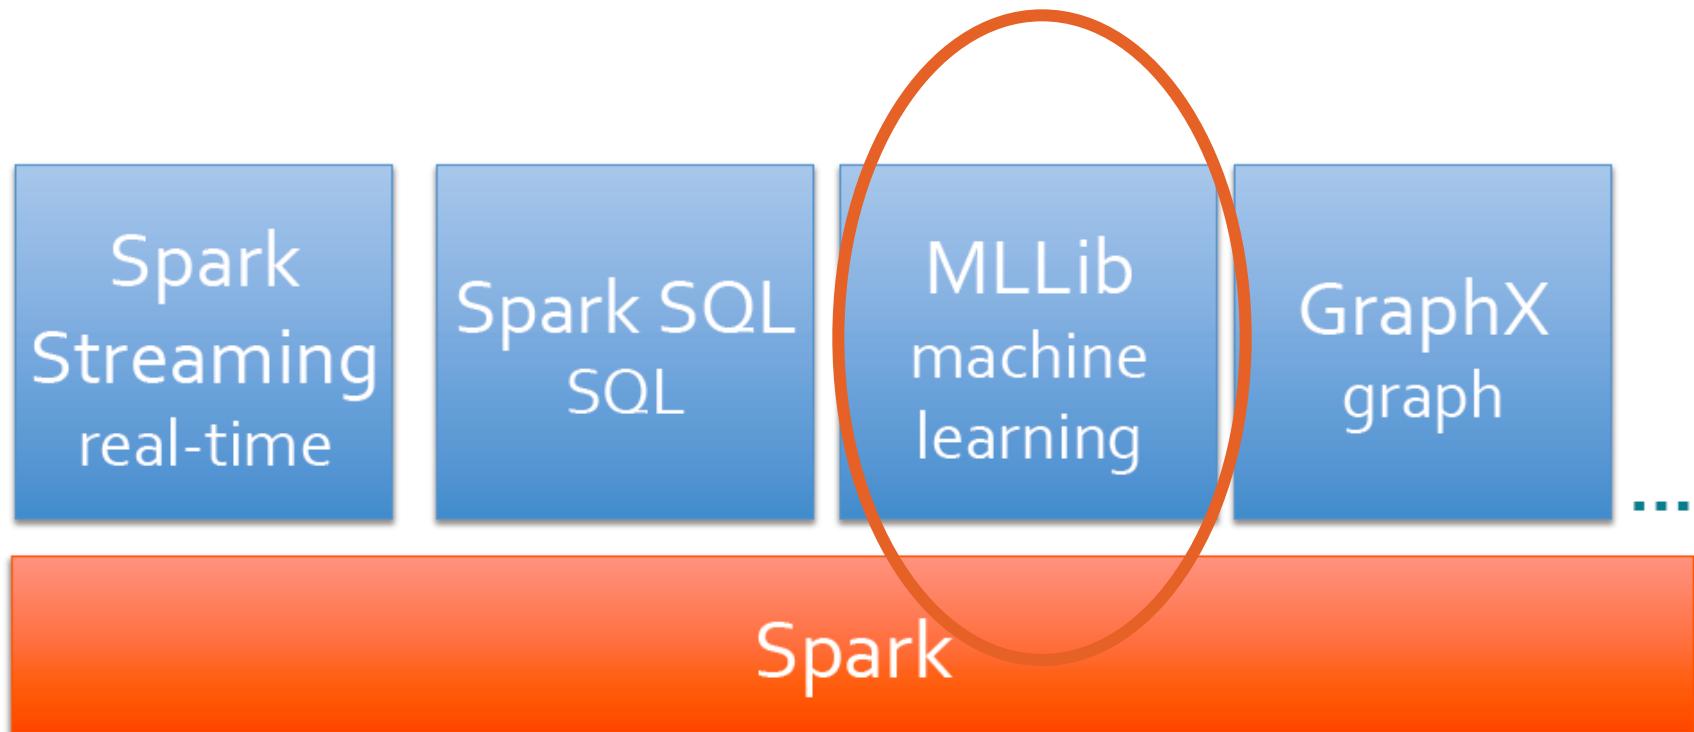
Routing Table (RDD)



Edge Table (RDD)

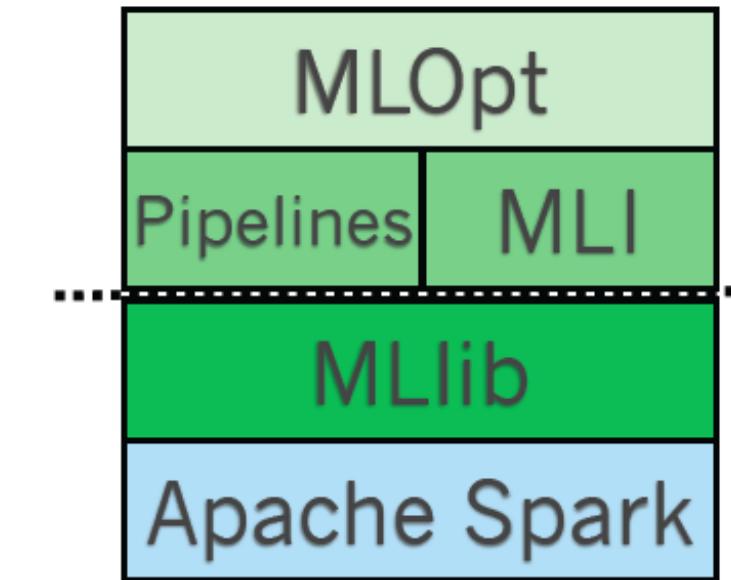


About Spark for MLlib



Machine Learning Support for Spark

- MLlib: Spark's core ML library
- The spark.ml package: High-level API built on the top of DataFrames for constructing Machine Learning Pipelines
 - Estimators, transformers, pipelines
- KeystoneML framework: Provide a richer set of operators than those in spark.ml
 - Featurizers for images, text, speech and example pipelines to reproduce state-of-the-art academic results in ML
 - Other related research projects:
- MLBase: MLOpt, MLI (UCB AMPLabs)
- Velox (UCB AMPLabs)
- Parameter Server (CMU)



Machine Learning Support for Spark

- **DataFrames**
 - Structured Data
 - Familiar API based on R & Python Pandas
 - Distributed, Optimized Implementation
- **Machine Learning Pipelines**
 - Simple construction and Tuning of Machine Learning workflows

Machine Learning Support for Spark (via MLlib or spark.ml)

Classification

- Logistic regression w/ elastic net
- Naive Bayes
- Streaming logistic regression
- Linear SVMs
- Decision trees
- Random forests
- Gradient-boosted trees
- Multilayer perceptron
- One-vs-rest

Regression

- Least squares w/ elastic net
- Isotonic regression
- Decision trees
- Random forests
- Gradient-boosted trees
- Streaming linear methods

Recommendation

- Alternating Least Squares

Frequent itemsets

- FP-growth
- Prefix span

Feature extraction & selection

- Binarizer
- Bucketizer
- Chi-Squared selection
- CountVectorizer
- Discrete cosine transform
- ElementwiseProduct
- Hashing term frequency
- Inverse document frequency
- MinMaxScaler
- Ngram
- Normalizer
- One-Hot Encoder
- PCA
- PolynomialExpansion
- RFormula
- SQLTransformer
- Standard scaler
- StopWordsRemover
- StringIndexer
- Tokenizer
- StringIndexer
- VectorAssembler
- VectorIndexer
- VectorSlicer
- Word2Vec

Clustering

- Gaussian mixture models
- K-Means
- Streaming K-Means
- Latent Dirichlet Allocation
- Power Iteration Clustering

Statistics

- Pearson correlation
- Spearman correlation
- Online summarization
- Chi-squared test
- Kernel density estimation

Linear algebra

- Local dense & sparse vectors & matrices
- Distributed matrices
 - Block-partitioned matrix
 - Row matrix
 - Indexed row matrix
 - Coordinate matrix
- Matrix decompositions

Model import/export

Pipelines

List based on Spark 1.5

Machine Learning Workflows are Complex

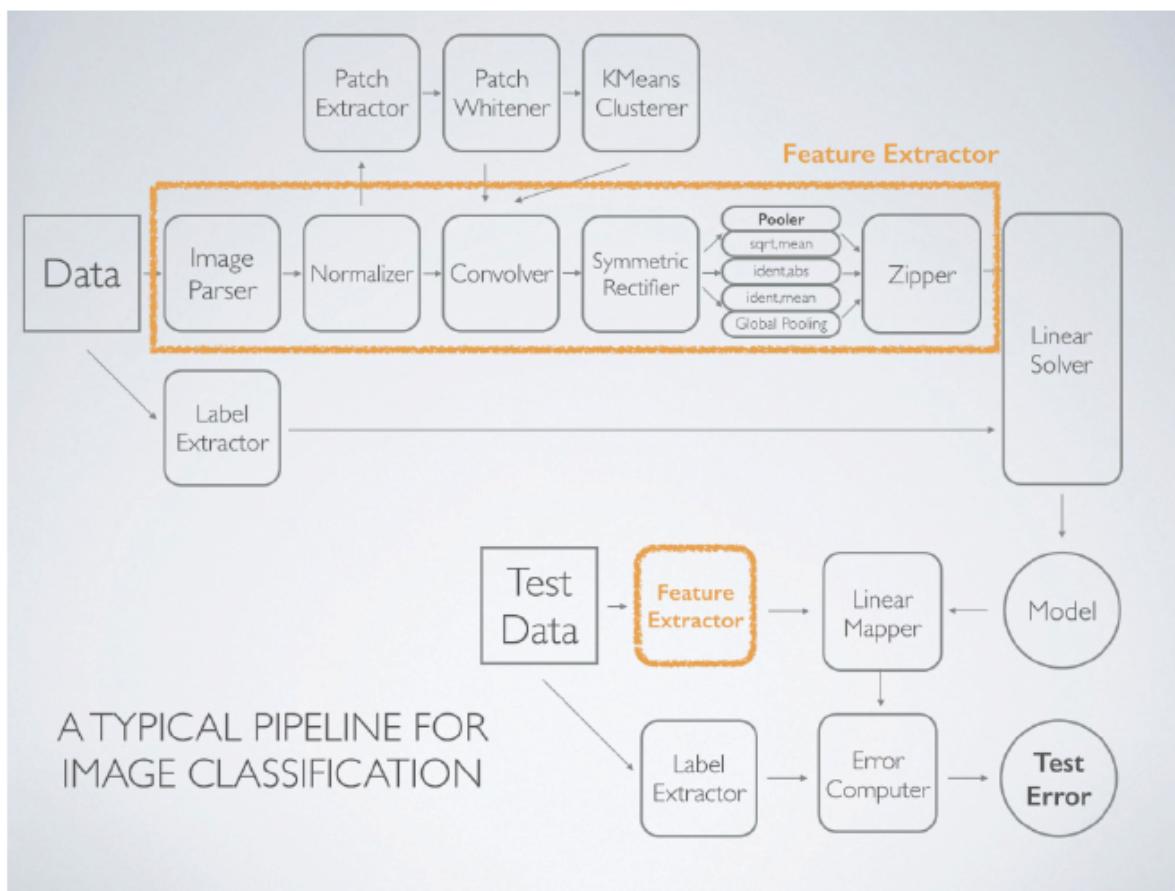


Image classification pipeline*

- Specify pipeline
- Inspect & debug
- Re-run on new data
- Tune parameters

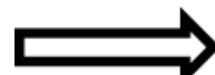
* Evan Sparks. "ML Pipelines."
amplab.cs.berkeley.edu/ml-pipelines

Example: Text Classification

Goal: Given a text document, predict its topic.

Features

Subject: Re: Lexan Polish?
Suggest McQuires #1 plastic
polish. It will help somewhat
but nothing will remove deep
scratches without making it
worse than it already is.
McQuires will do something...

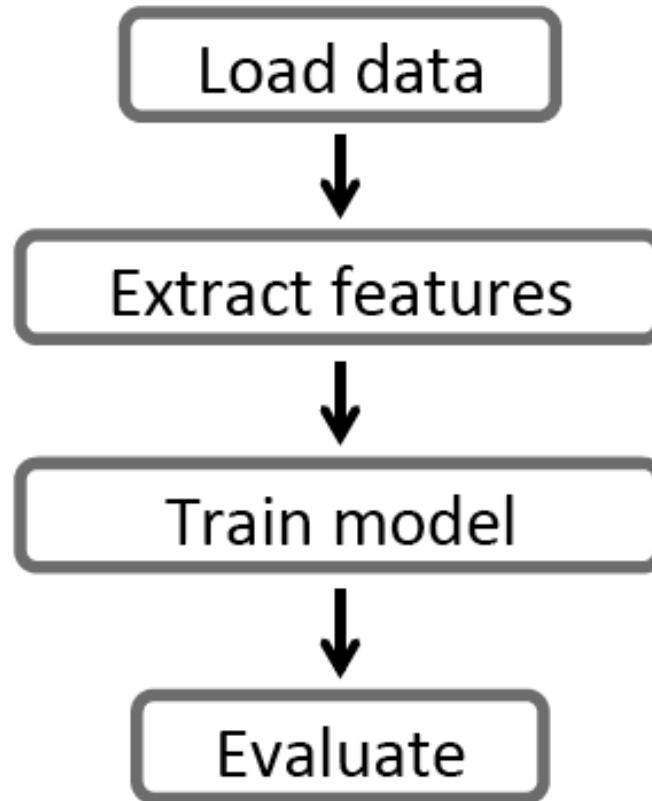


Label

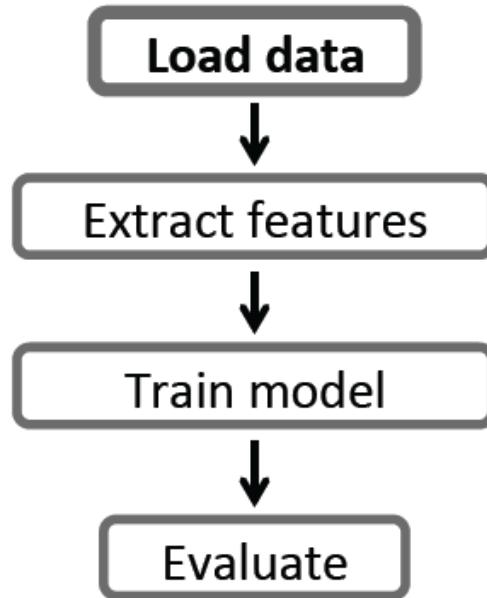
1: about science
0: not about science

Dataset: “20 Newsgroups”
From UCI KDD Archive

Machine Learning Workflow



Load Data



Data sources for DataFrames

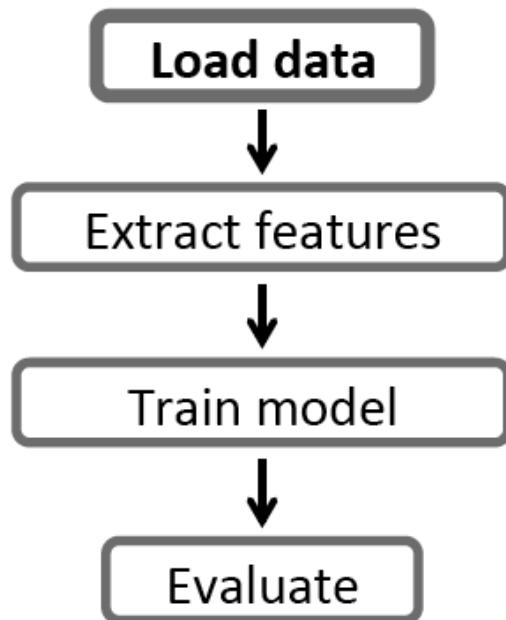
built-in



external



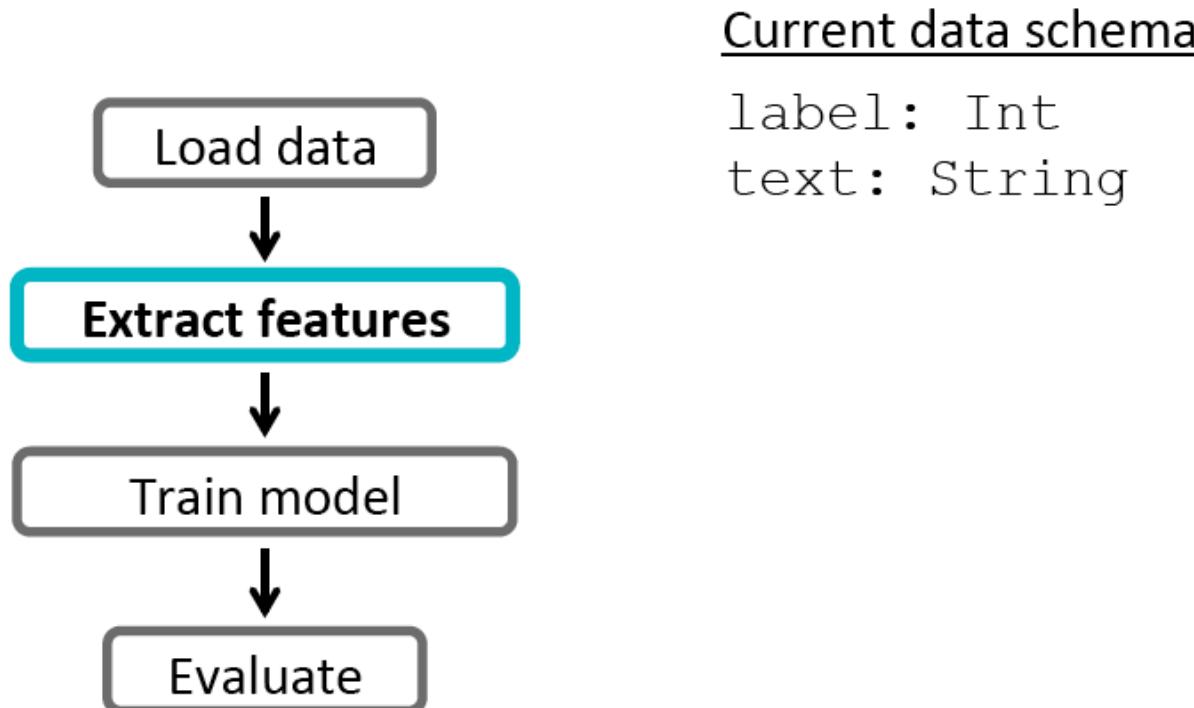
Load Data



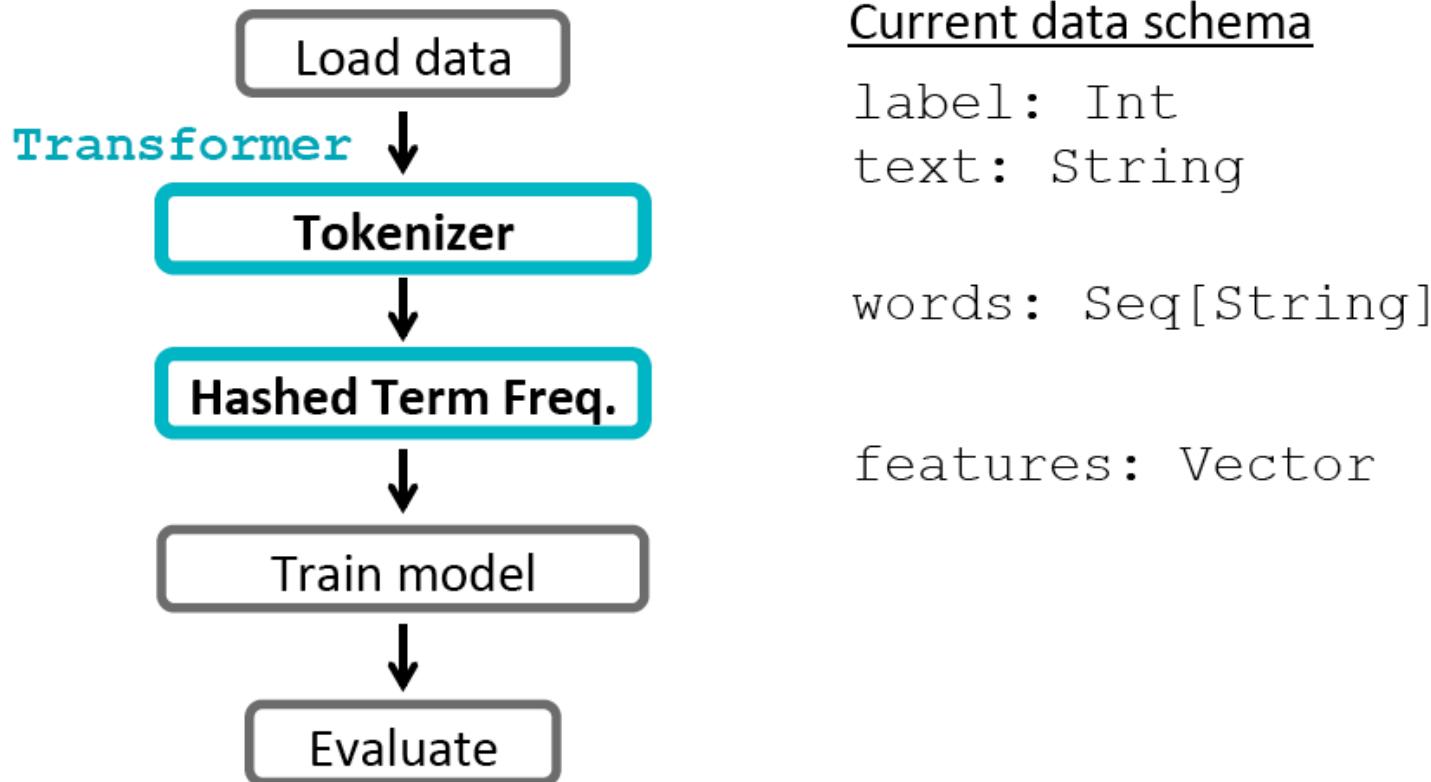
Current data schema

label: Int
text: String

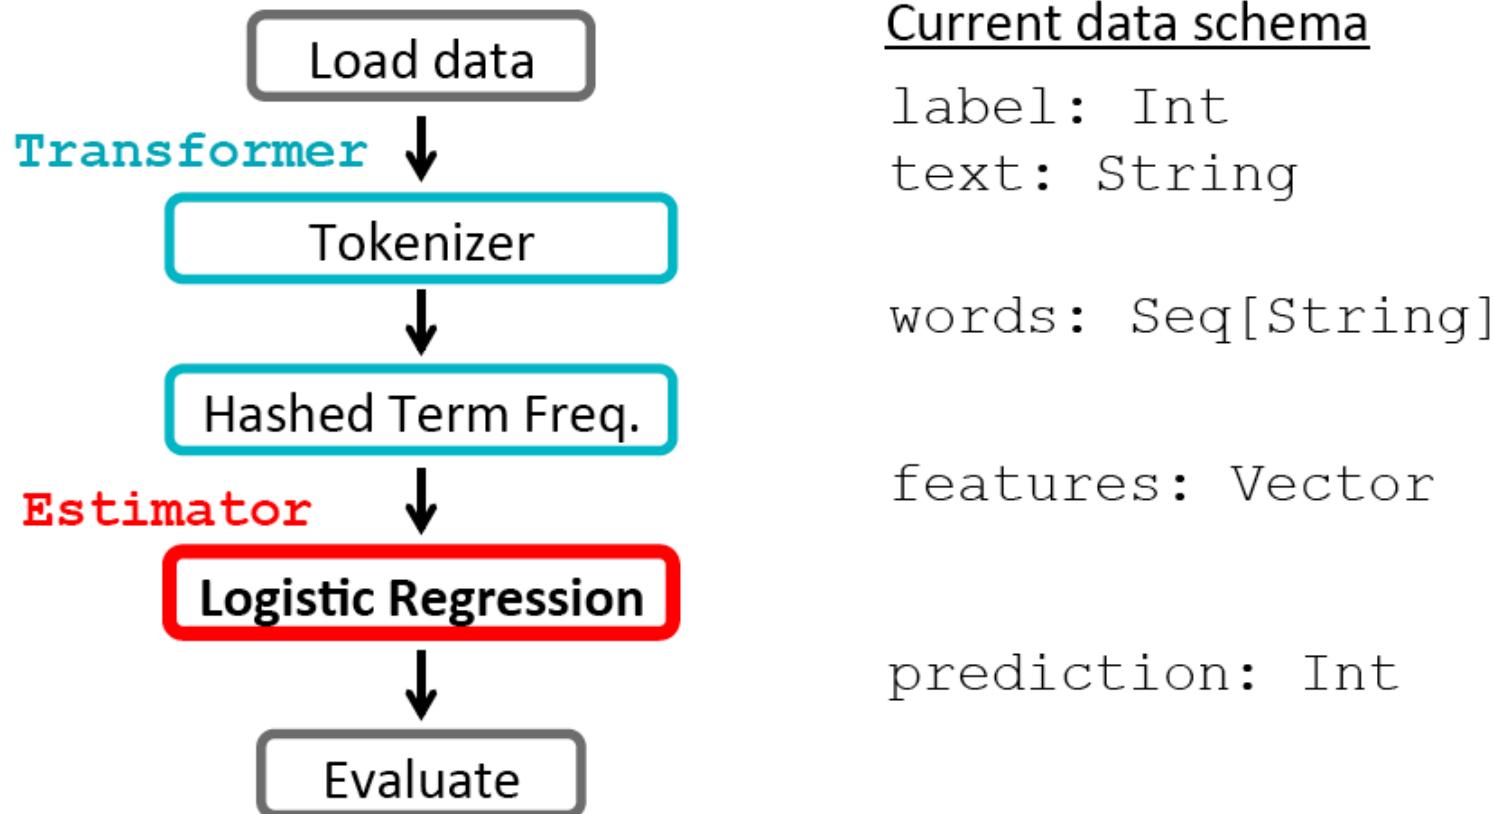
Extract Features



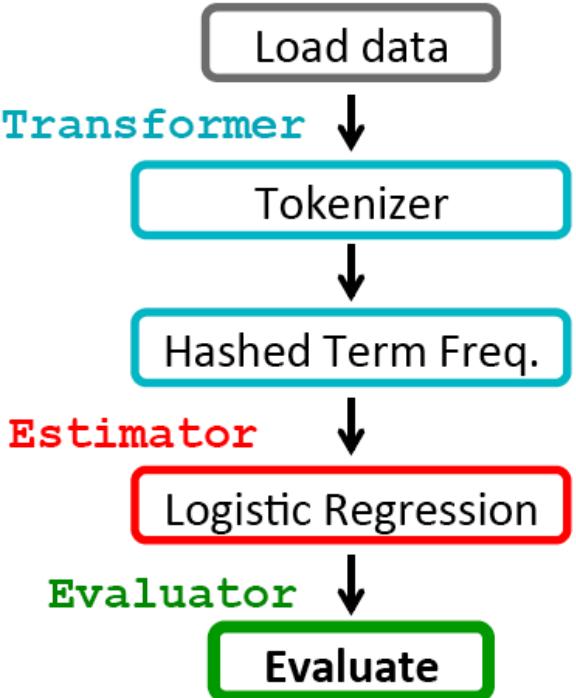
Extract Features



Train a Model



Evaluate the Model



Current data schema

label: Int

text: String

words: Seq[String]

features: Vector

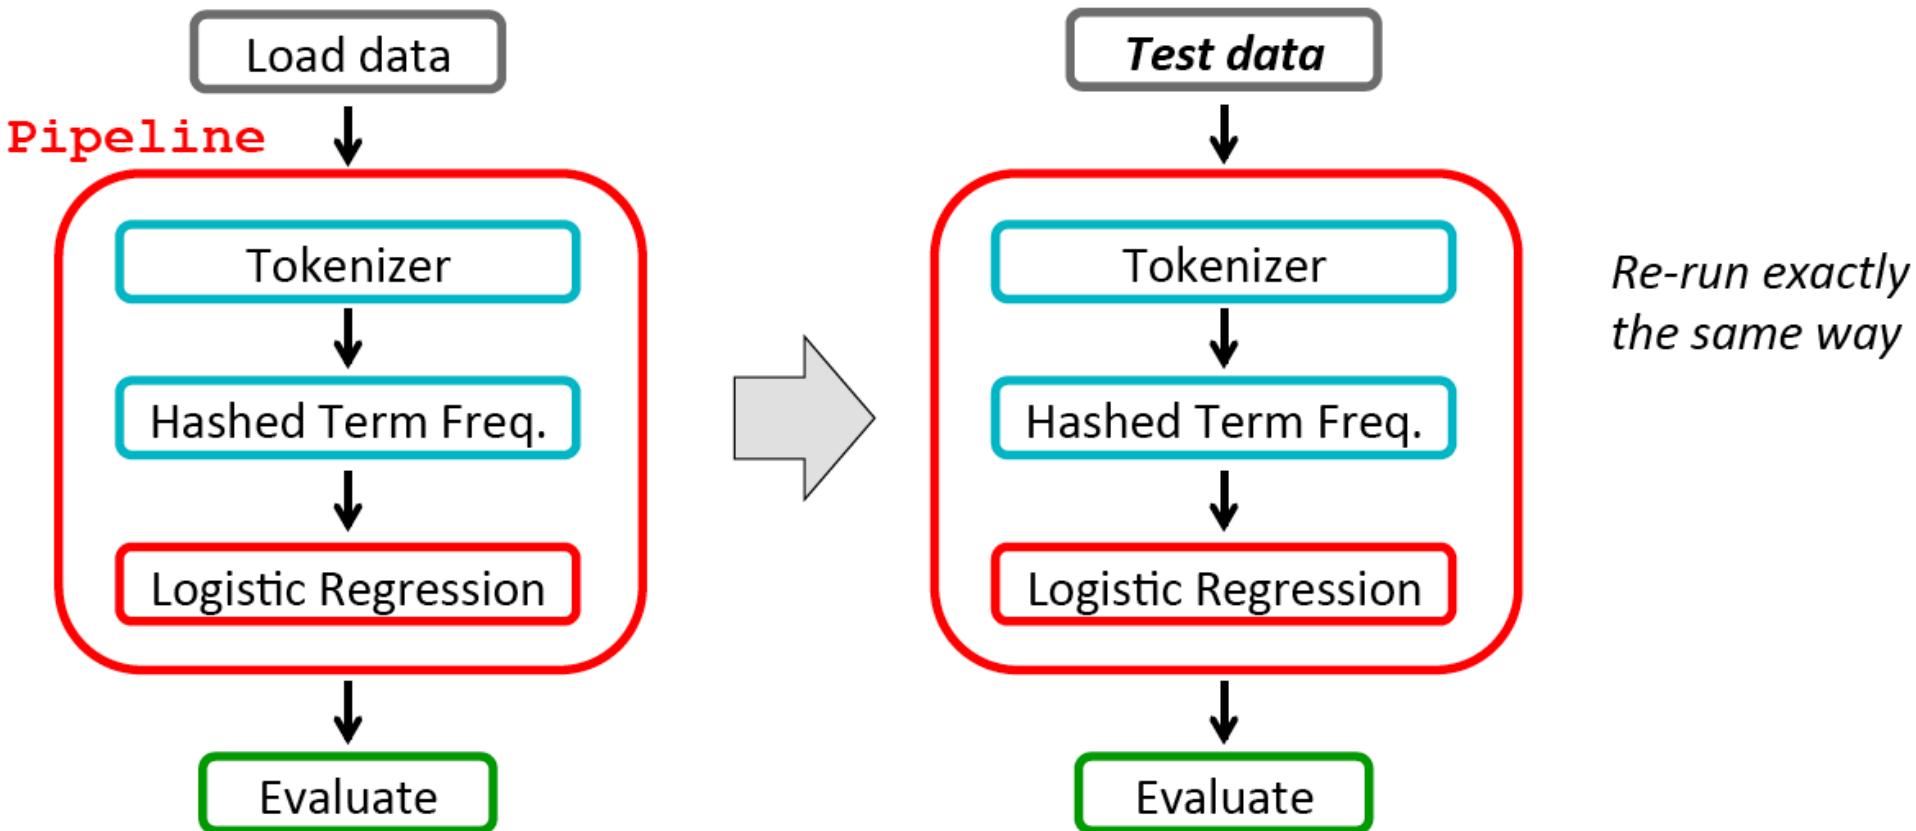
prediction: Int

By default, always append new columns

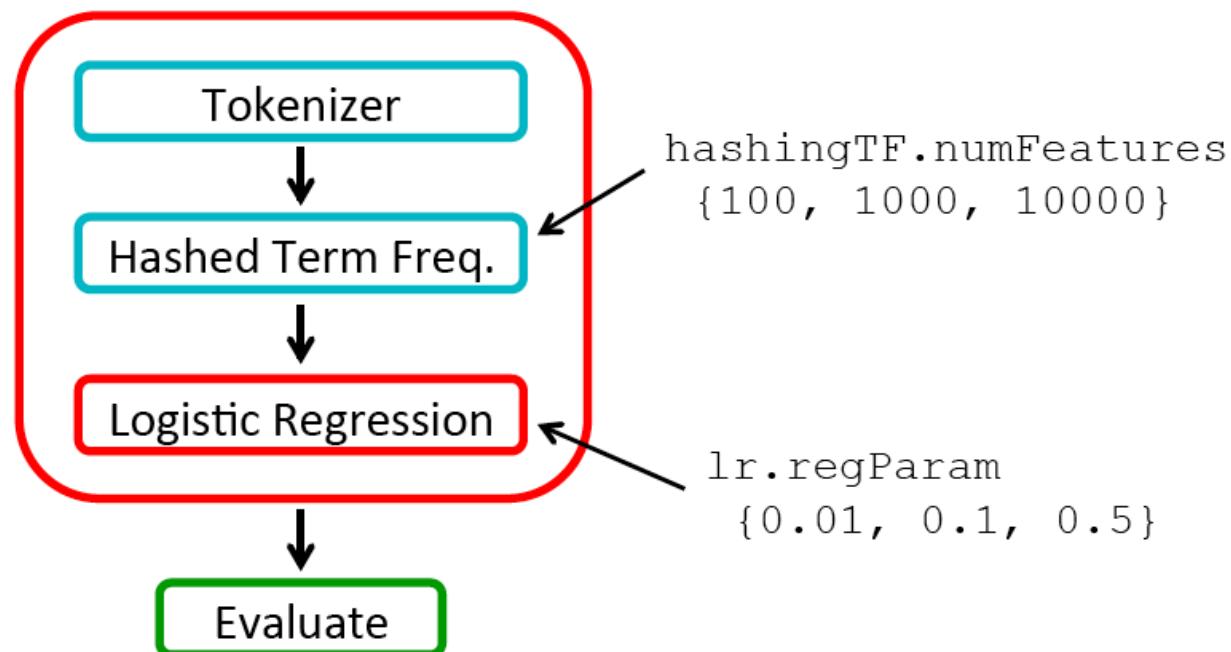
→ Can go back & inspect intermediate results

→ Made efficient by DataFrame optimizations

Machine Learning Pipelines



Parameter Tuning



CrossValidator

Given:

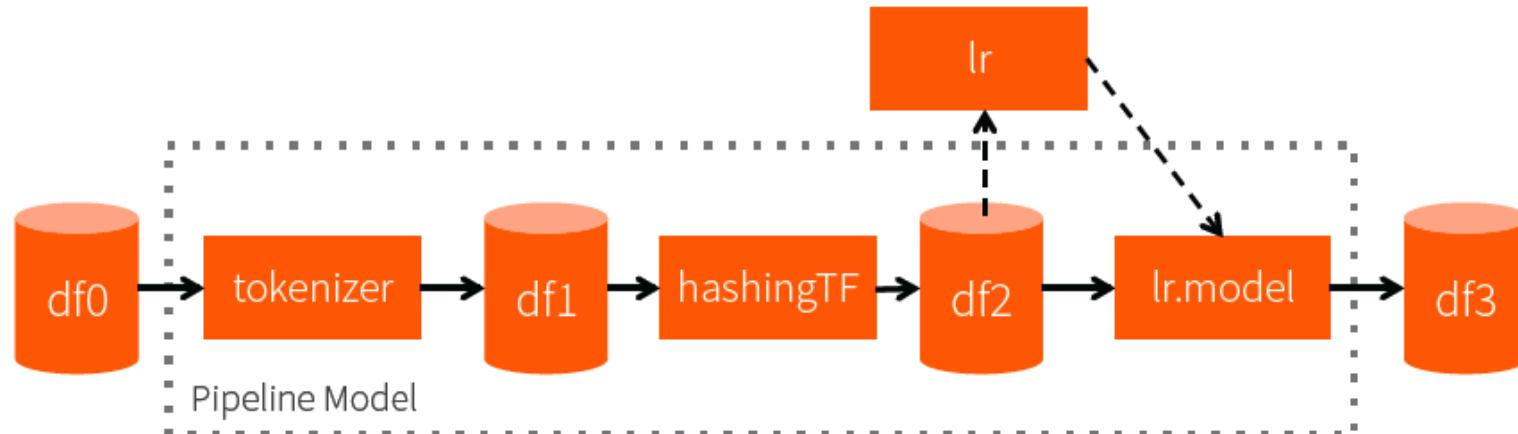
- Estimator
- Parameter grid
- Evaluator

Find best parameters

DataFrame for Machine Learning Pipelines

```
tokenizer = Tokenizer(inputCol="text", outputCol="words")
hashingTF = HashingTF(inputCol="words", outputCol="features")
lr = LogisticRegression(maxIter=10, regParam=0.01)
pipeline = Pipeline(stages=[tokenizer, hashingTF, lr])

df = sqlCtx.load("/path/to/data")
model = pipeline.fit(df)
```



SQL and Machine Learning

```
training_data_table = sql("""
    SELECT e.action, u.age, u.latitude, u.longitude
    FROM Users u
    JOIN Events e ON u.userId = e.userId""")  
  
def featurize(u):
    LabeledPoint(u.action, [u.age, u.latitude, u.longitude])  
  
// SQL results are RDDs so can be used directly in MLlib.
training_data = training_data_table.map(featurize)
model = new LogisticRegressionWithSGD.train(training_data)
```

Machine Learning Pipelines

```
// training:{eventId:String, features:Vector, label:Int}
val training = parquetFile("/path/to/training")
val lr = new LogisticRegression().fit(training)

// event: {eventId: String, features: Vector}
val event = parquetFile("/path/to/event")
val prediction =
  lr.transform(event).select('eventId, 'prediction)

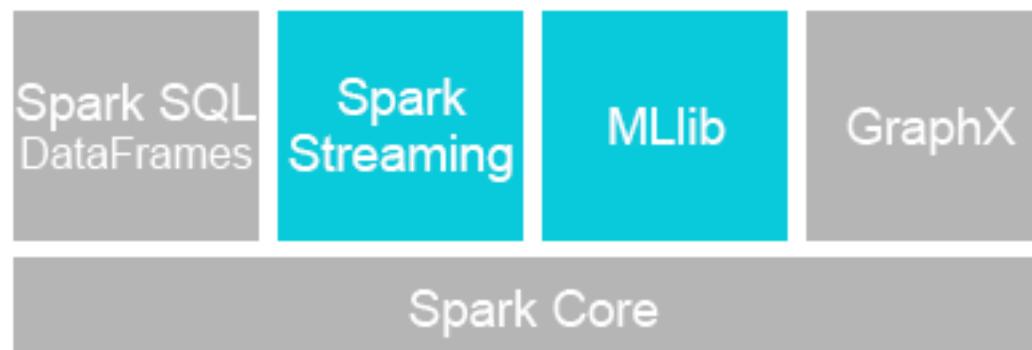
prediction.saveAsParquetFile("/path/to/prediction")
```

Combine Machine Learning with Streaming

- Learn models offline, apply them online

```
// Learn model offline  
val model = KMeans.train(dataset, ...)
```

```
// Apply model online on stream  
kafkaStream.map { event =>  
    model.predict(event.feature)  
}
```



Streaming MLlib Algorithms

Continuous learning and prediction on streaming data

StreamingLinearRegression, StreamingKMeans,
StreamingLogisticRegression

```
val model = new StreamingKMeans()  
    .setK(10).setDecayFactor(1.0).setRandomCenters(4, 0.0)  
  
model.trainOn(trainingDStream) // Train on one DStream  
  
// Predict on another DStream  
model.predictOnValues( testDStream.map { lp => (lp.label, lp.features) } )
```

<https://databricks.com/blog/2015/01/28/introducing-streaming-k-means-in-spark-1-2.html>

```
lines = KafkaUtils.createStream(  
    streamingContext, kafkaTopics, kafkaParams)  
  
counts = lines.flatMap(lambda line: line.split(" "))
```

Summary of DataFrames and ML Workflow

DataFrames

- Structured Data Schema validation
- Familiar API based on R & Python Pandas
- Distributed, Optimized Implementation User-defined Transformers & Estimators

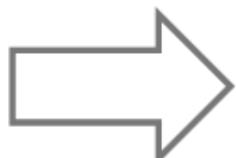
Machine Learning Pipelines

- Integration with DataFrames
- Familiar API based on scikit-learn
- Simple Parameter Tuning
 - Ongoing Research on Auto-Tuning Models

Composable & DAG Pipelines

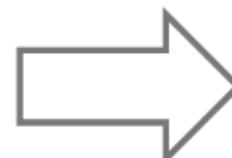
Enabling Interactive (Big) Data Science with SparkR

Spark early adopters

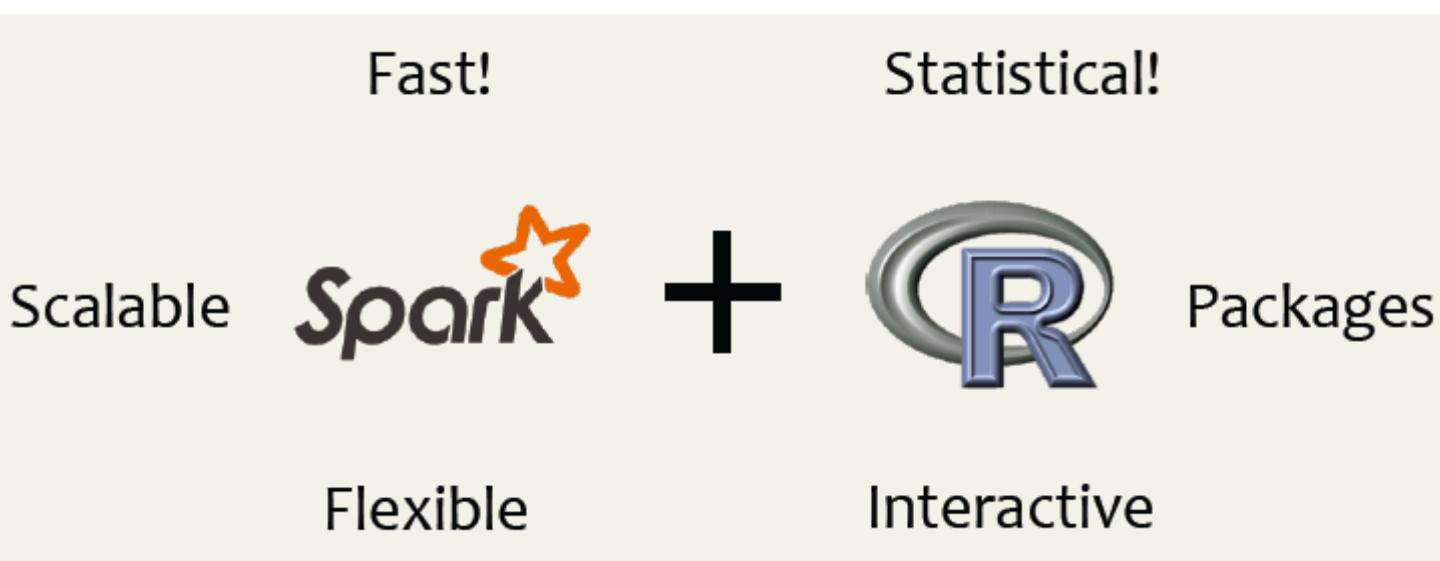


Users

Understands
MapReduce
& functional APIs



Data Engineers
Data Scientists
Statisticians
R users
PyData ...



SparkR – R package for Spark

- R Interface support via SparkR (R with RDD = R2D2)
since Spark 1.4 (released since June 2015)
 - Exposes DataFrames and MLlib in R:



```
df = jsonFile("tweets.json")  
  
summarize(  
  group_by(  
    df[df$user == "matei"],  
    "date"),  
  sum("retweets"))
```

SparkR – R package for Spark

RDD → distributed lists

SparkR

Run R on clusters

Re-use existing packages

Combine scalability & utility

Getting closer to Idiomatic R

Q: How can I use a loop to [...insert task here...] ?

A: **Don't.** Use one of the **apply** functions.

SparkR

R + RDD =
RRDD



lapply
lapplyPartition
groupByKey
reduceByKey
sampleRDD
collect
cache
...
broadcast
includePackage
textFile
parallelize

Example: Word Counting with SparkR

```
lines <- textFile(sc, "hdfs://my_text_file")  
  
words <- flatMap(lines,  
                  function(line) {  
                    strsplit(line, " ")[[1]]  
                  })  
wordCount <- lapply(words,  
                      function(word) {  
                        list(word, 1L)  
                      })  
counts <- reduceByKey(wordCount, "+", 2L)  
output <- collect(counts)
```

Example: Logistic Regression with SparkR

```
pointsRDD <- textFile(sc, "hdfs://myfile")
weights <- runif(n=D, min = -1, max = 1)

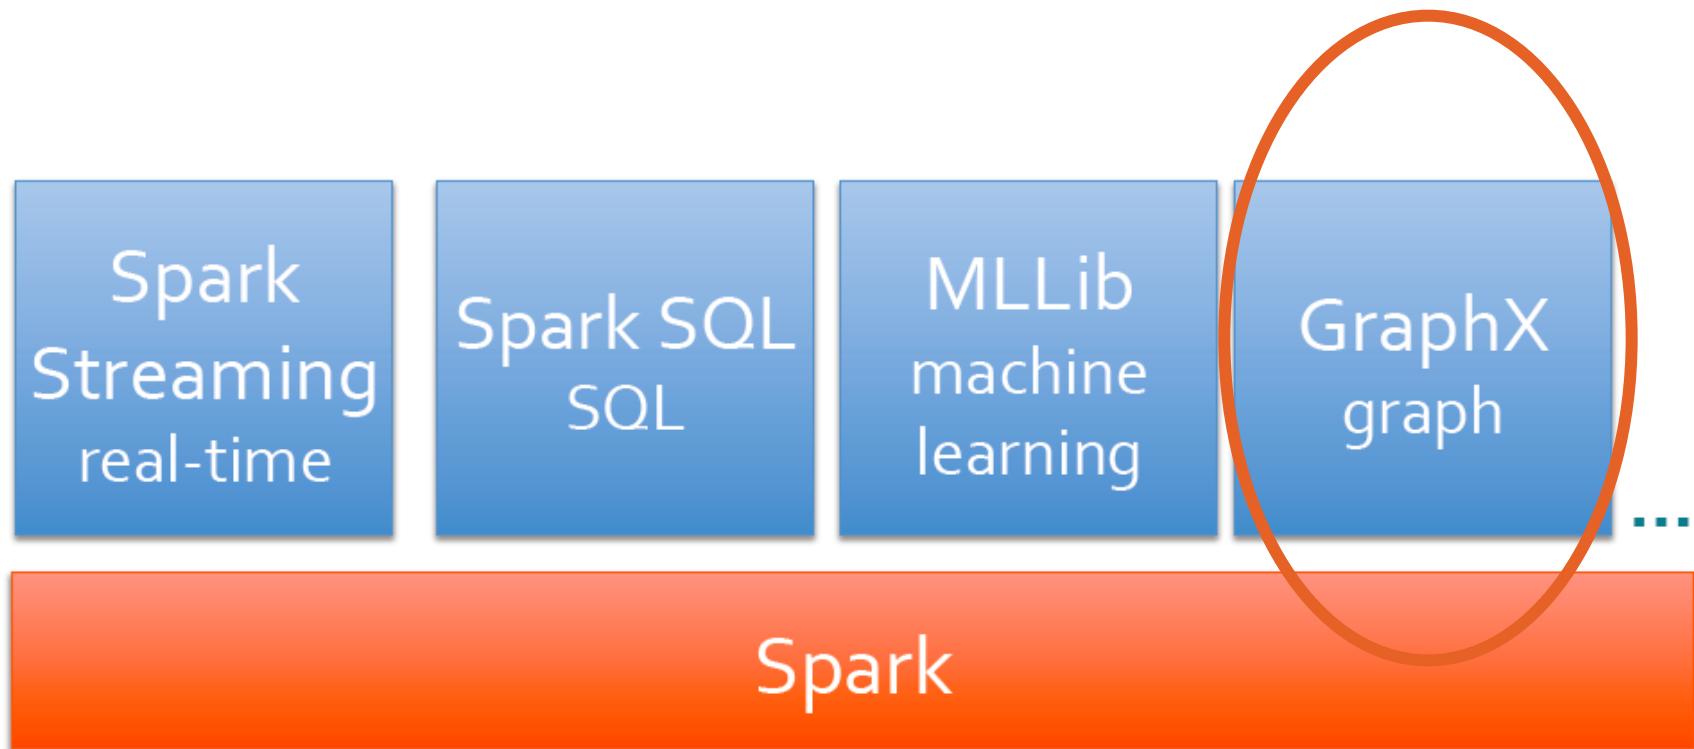
# Logistic gradient
gradient <- function(partition) {
  X <- partition[,1]; Y <- partition[,-1]
  t(X) %*% (1/(1 + exp(-Y * (X %*% weights))) - 1) * Y
}

# Iterate
weights <- weights - reduce(
  lapplyPartition(pointsRDD, gradient), "+")
```

SparkR Implementation

- Very similar to PySpark
- Relatively easy to extend Spark
 - 329 lines of Scala code
 - 2079 lines of R code
 - 693 lines of Test code in R

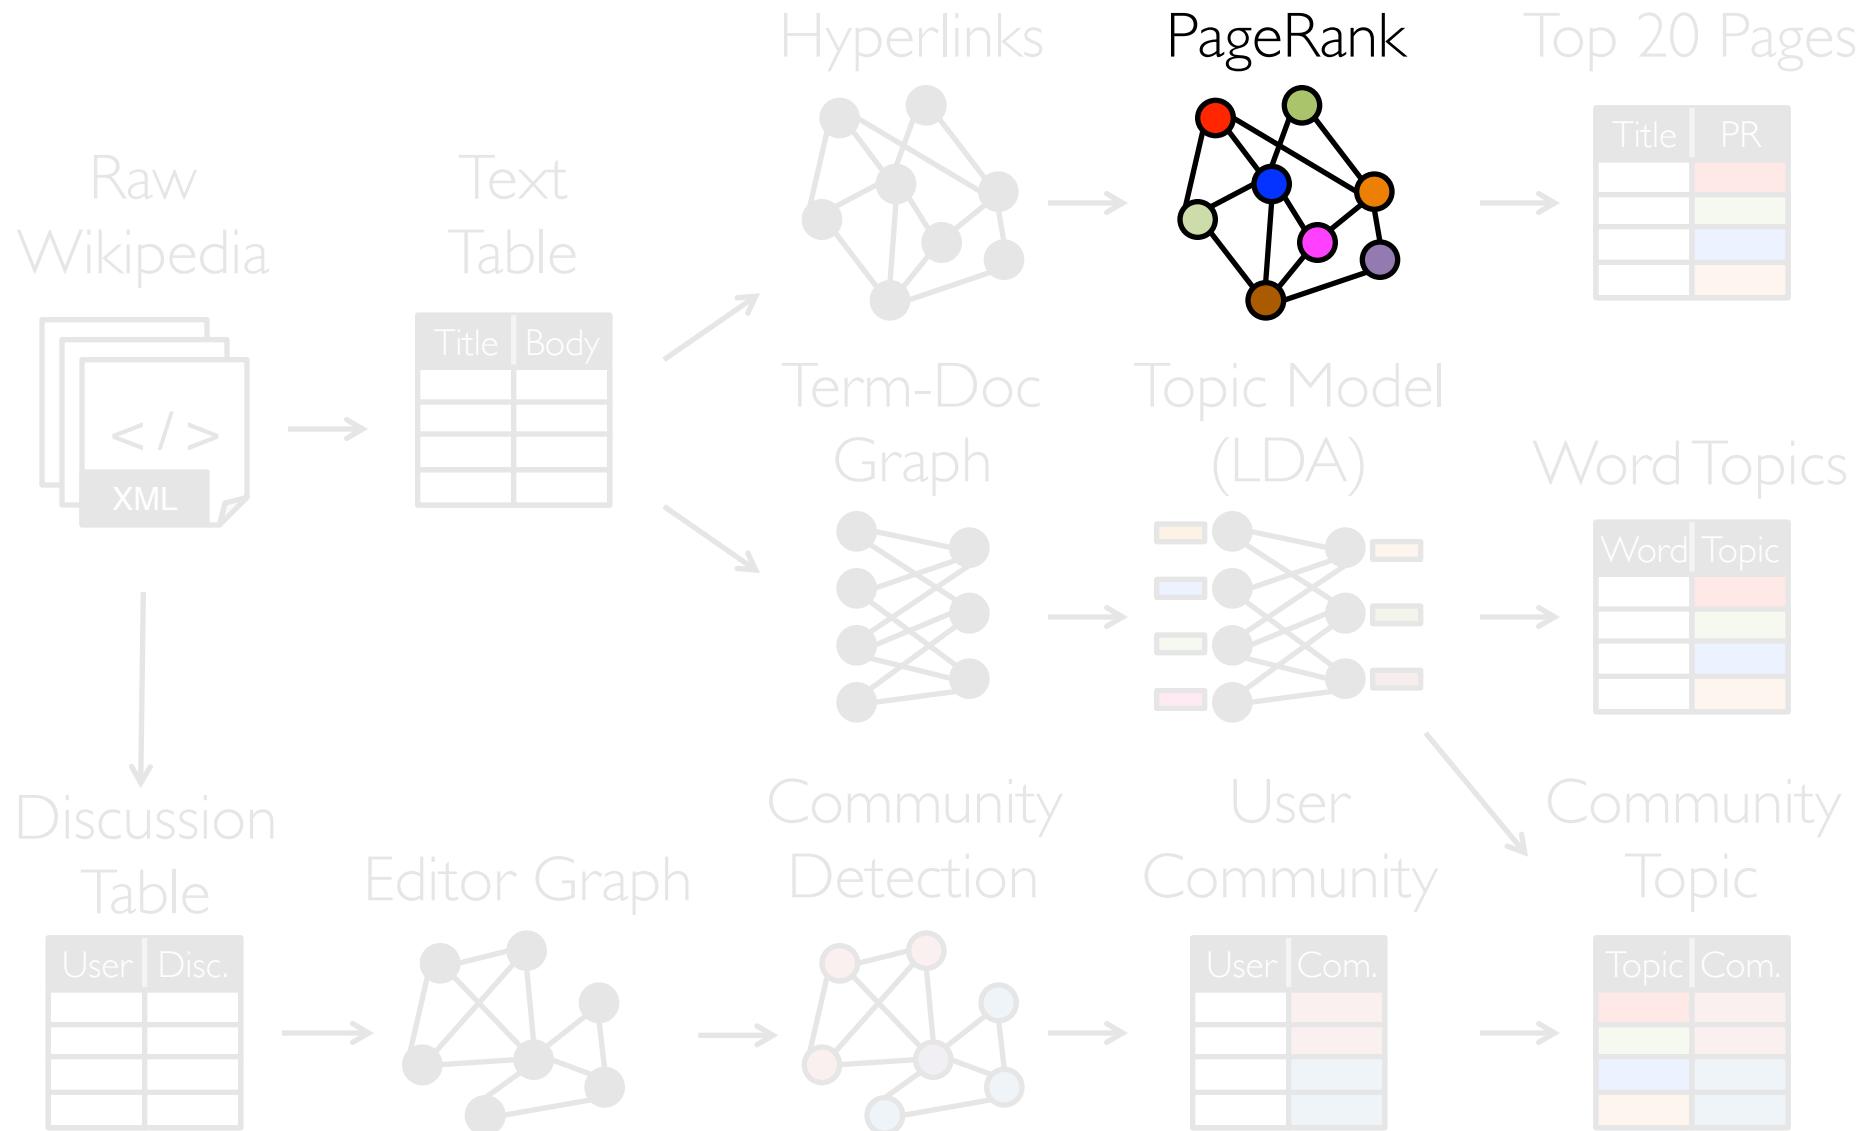
Generality of RDDs in Spark



GraphX: Full-Length Version

Unifying Data-Parallel and Graph-Parallel Analytics

Graphs are Central to Analytics



PageRank: Identifying Leaders

$$R[i] = 0.15 + \sum_{j \in \text{Nbrs}(i)} w_{ji} R[j]$$

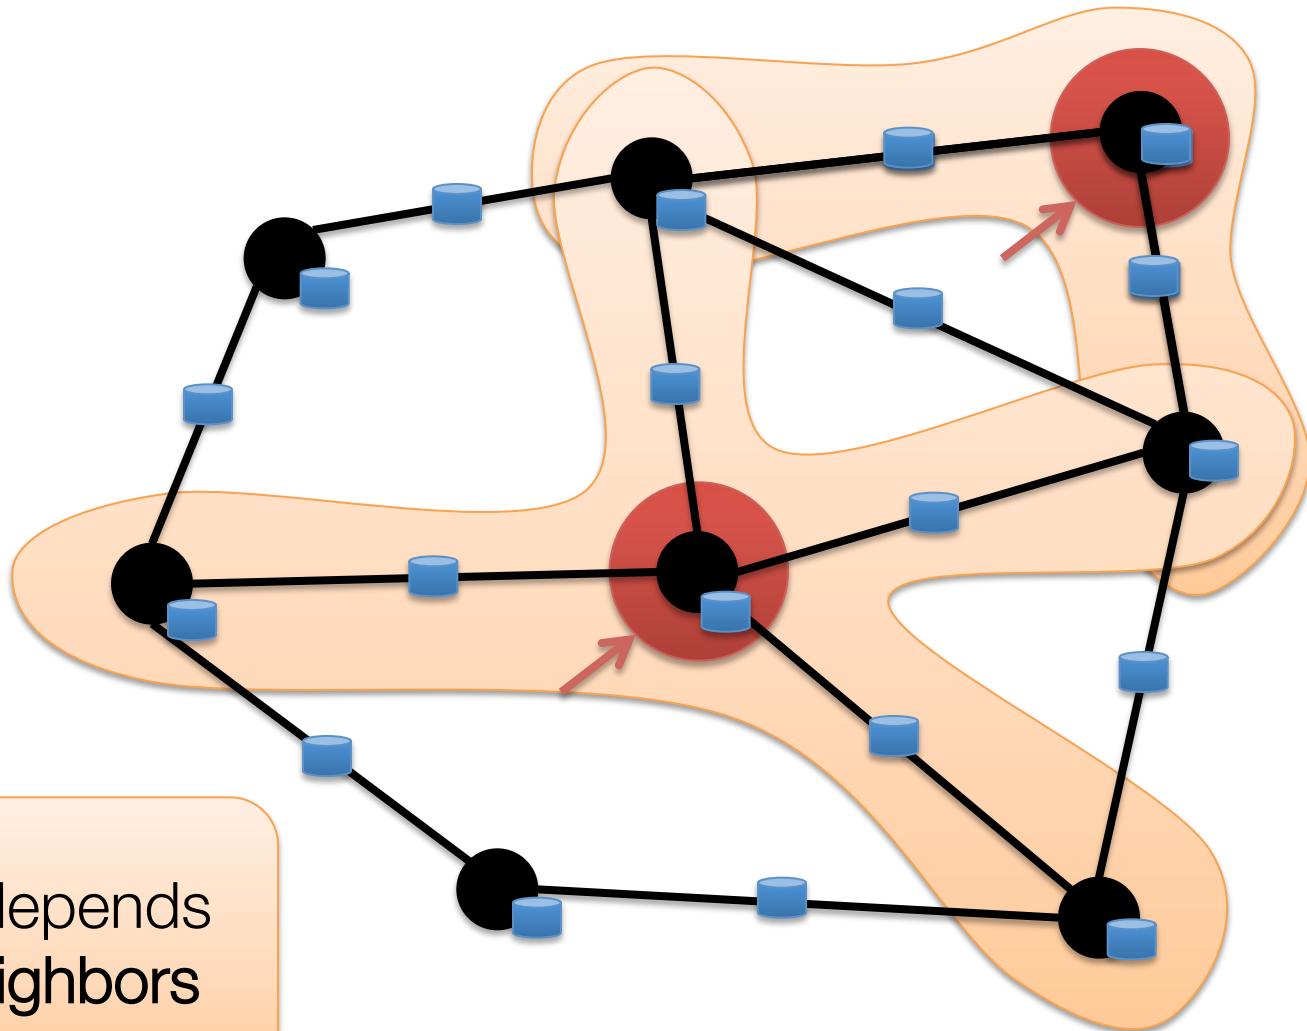
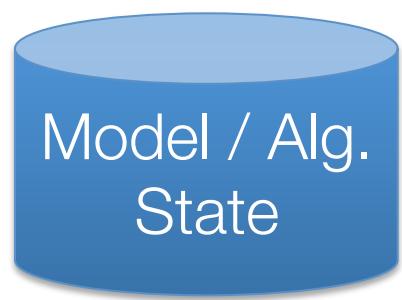
Rank of user i

Weighted sum of neighbors' ranks

Update ranks in parallel

Iterate until convergence

The Graph-Parallel Pattern



Computation depends
only on the **neighbors**

Many Graph-Parallel Algorithms

- Collaborative Filtering
 - Alternating Least Squares
 - Stochastic Gradient Descent
 - Tensor Factorization
- Structured Prediction
 - Loopy Belief Propagation
 - Max-Product Linear Programs
 - Gibbs Sampling
- Semi-supervised ML
 - Graph SSL
 - CoEM
- Community Detection
 - Triangle-Counting
 - K-core Decomposition
 - K-Truss
- Graph Analytics
 - PageRank
 - Personalized PageRank
 - Shortest Path
 - Graph Coloring
- Classification
 - Neural Networks

Graph-Parallel Systems

Pregel
oo^{gle}

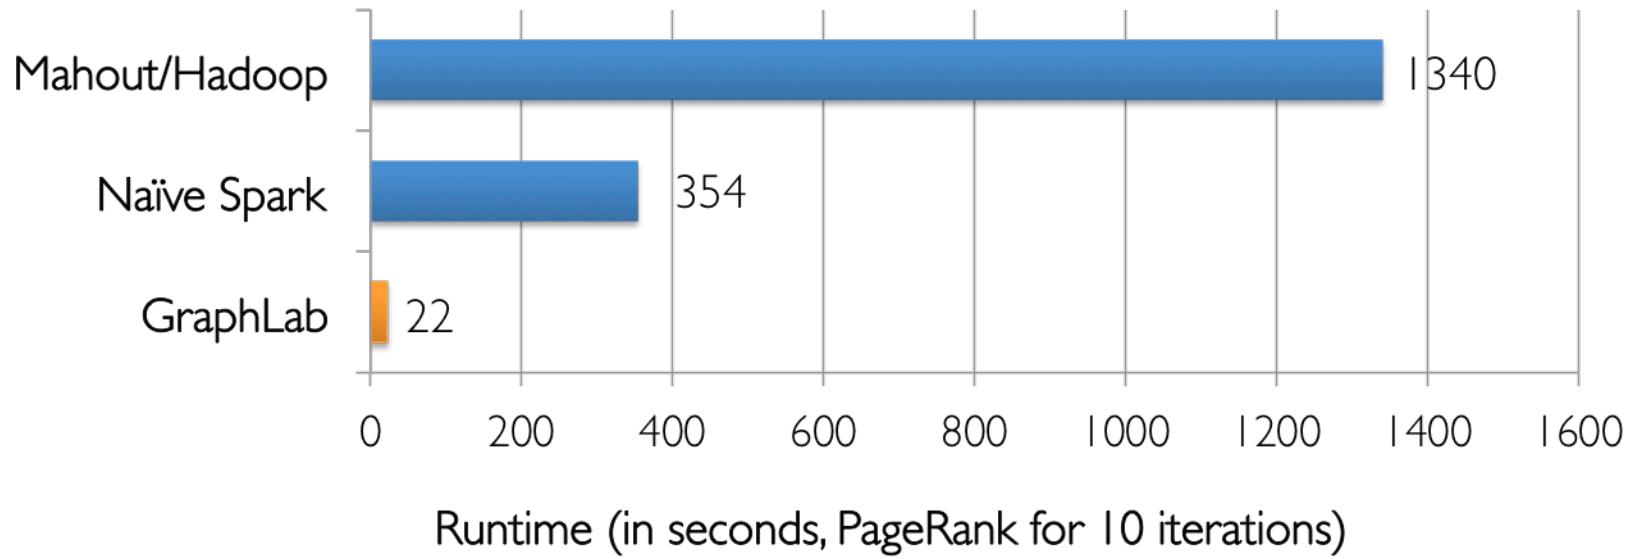


GraphLab

Expose *specialized APIs* to simplify graph programming.

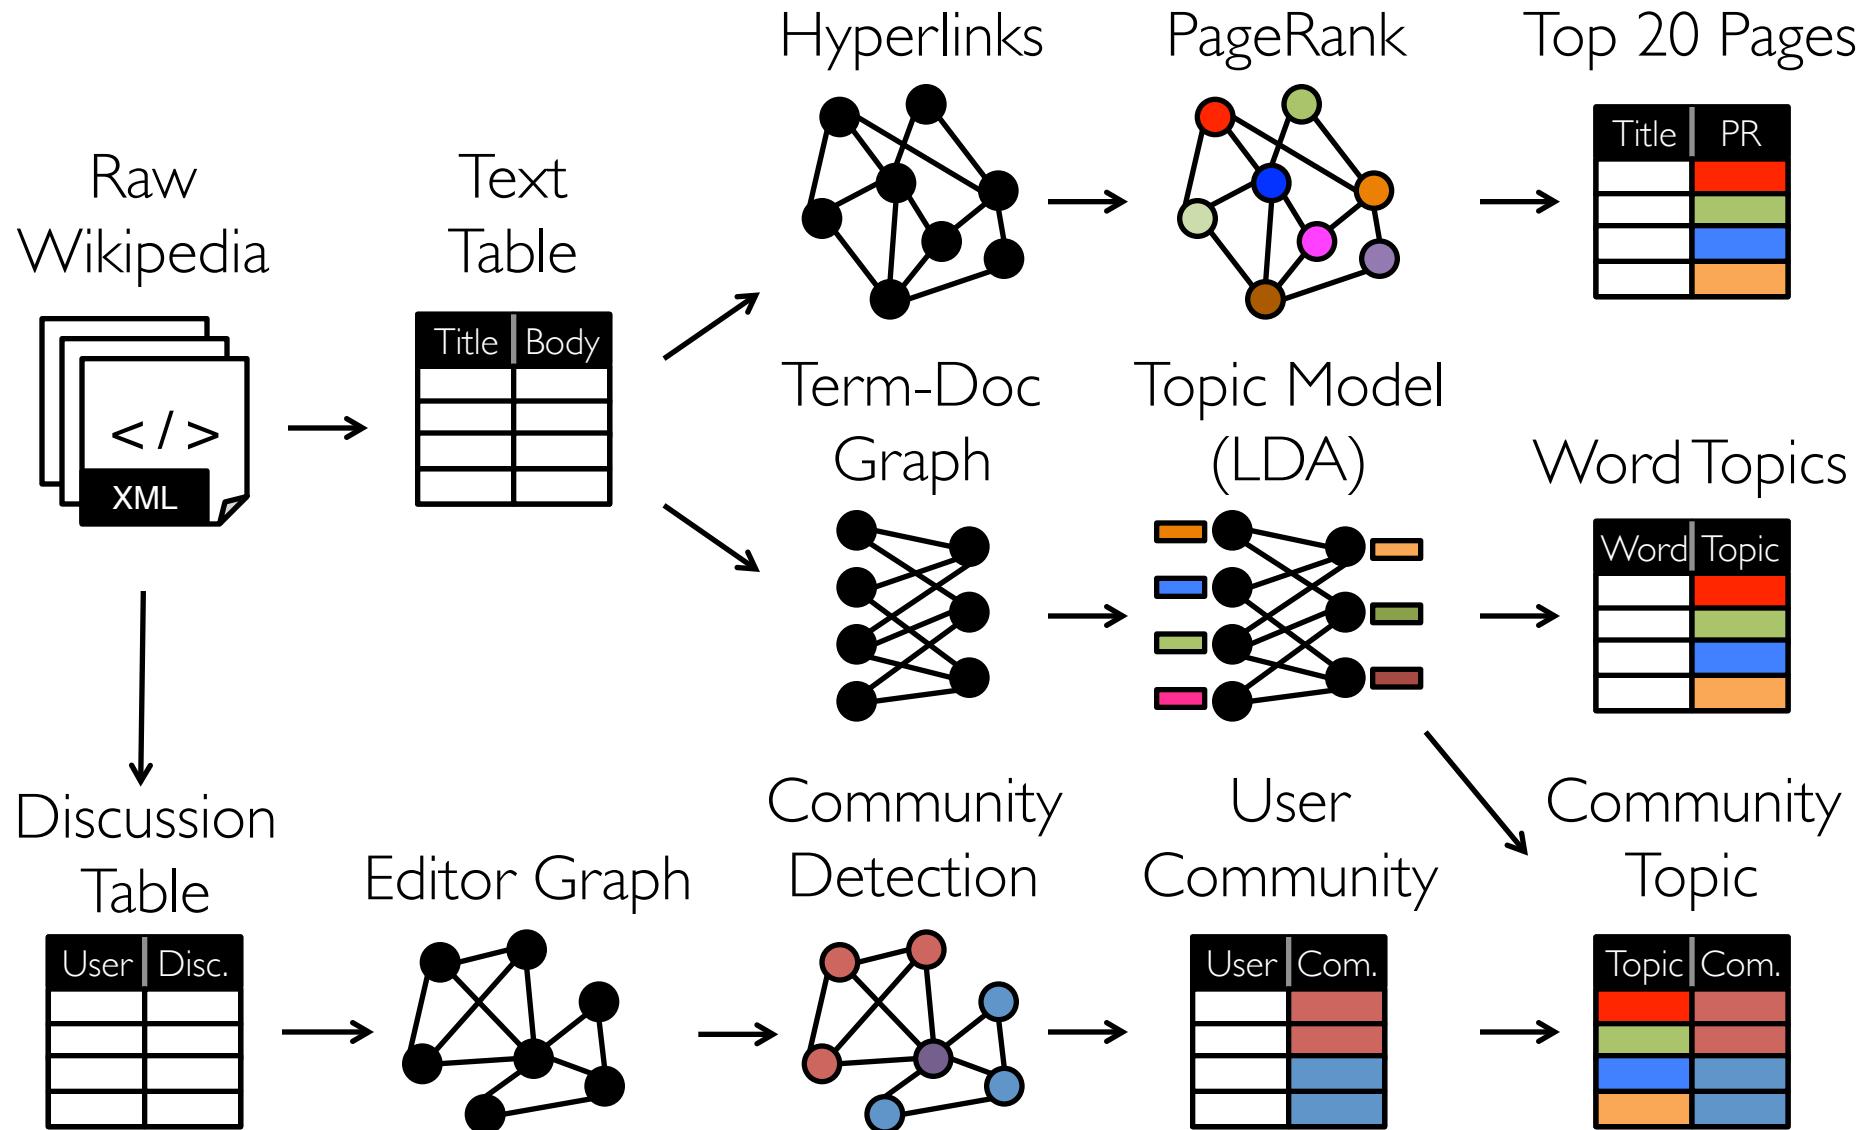
Exploit graph structure to achieve *orders-of-magnitude performance gains* over more general data-parallel systems.

PageRank on the Live-Journal Graph



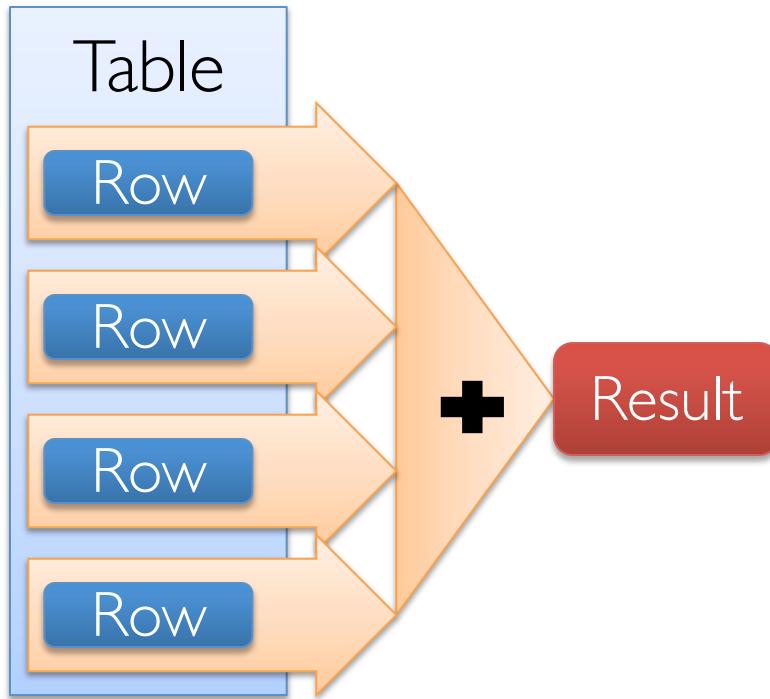
GraphLab is *60x faster* than Hadoop
GraphLab is *16x faster* than Spark

Graphs are Central to Analytics



Separate Systems to Support Each View

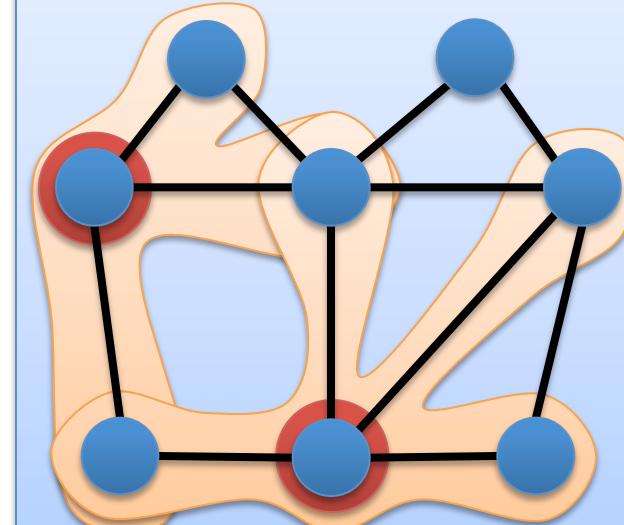
Table View



Graph View



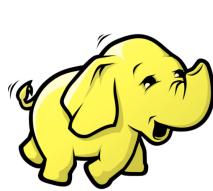
Dependency Graph



*Having separate systems
for each view is
difficult to use and inefficient*

Difficult to Program and Use

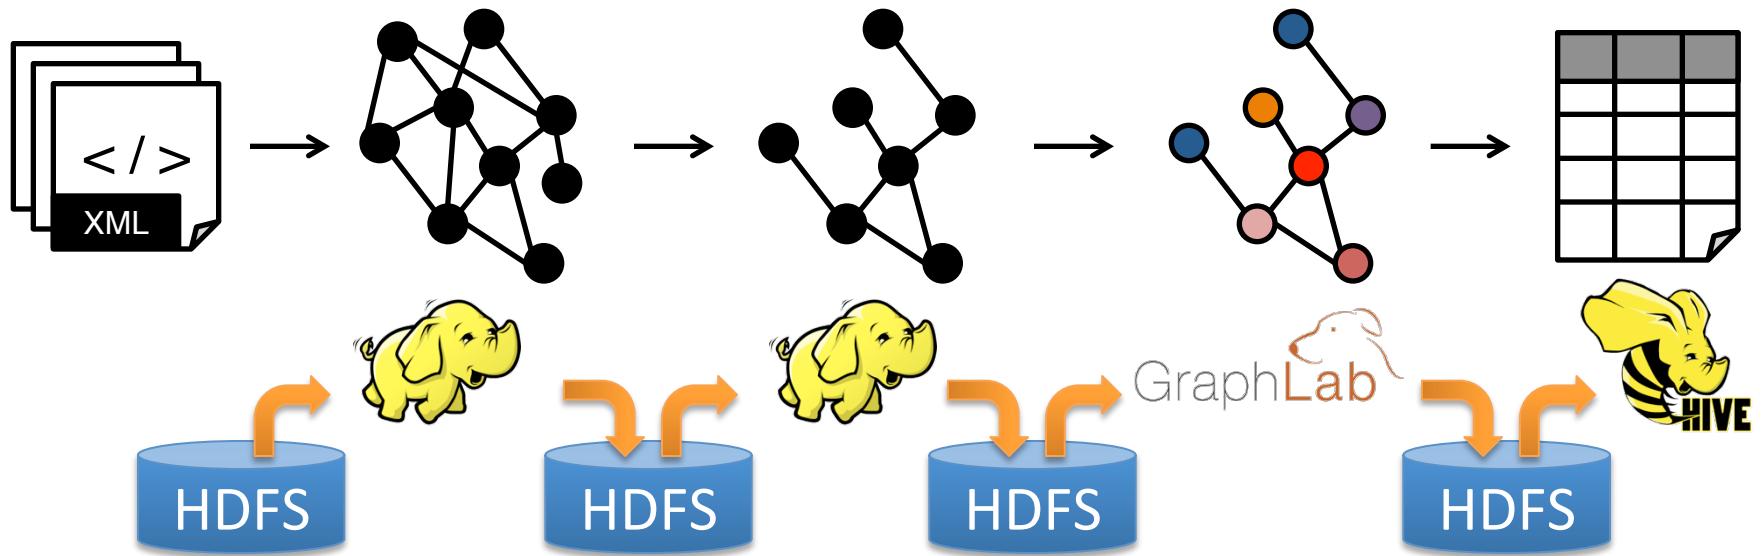
Users must *Learn, Deploy, and Manage* multiple systems



Leads to brittle and often
complex interfaces

Inefficient

Extensive **data movement** and **duplication** across
the network and file system

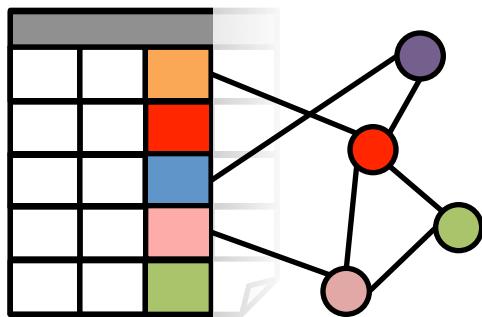


Limited reuse internal data-structures
across stages

*Solution:*The GraphX Unified Approach

New API

*Blurs the distinction between
Tables and Graphs*



New System

*Combines Data-Parallel
Graph-Parallel Systems*



Enabling users to **easily** and **efficiently**
express the entire graph analytics pipeline

Tables and Graphs are composable views of the same *physical* data

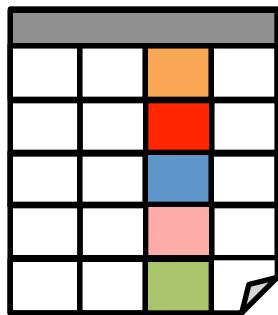
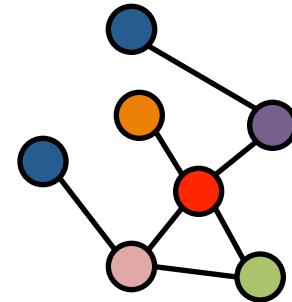
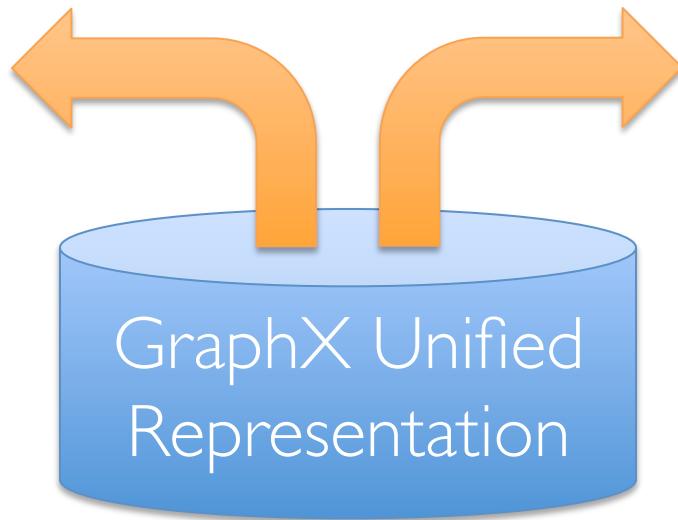


Table View

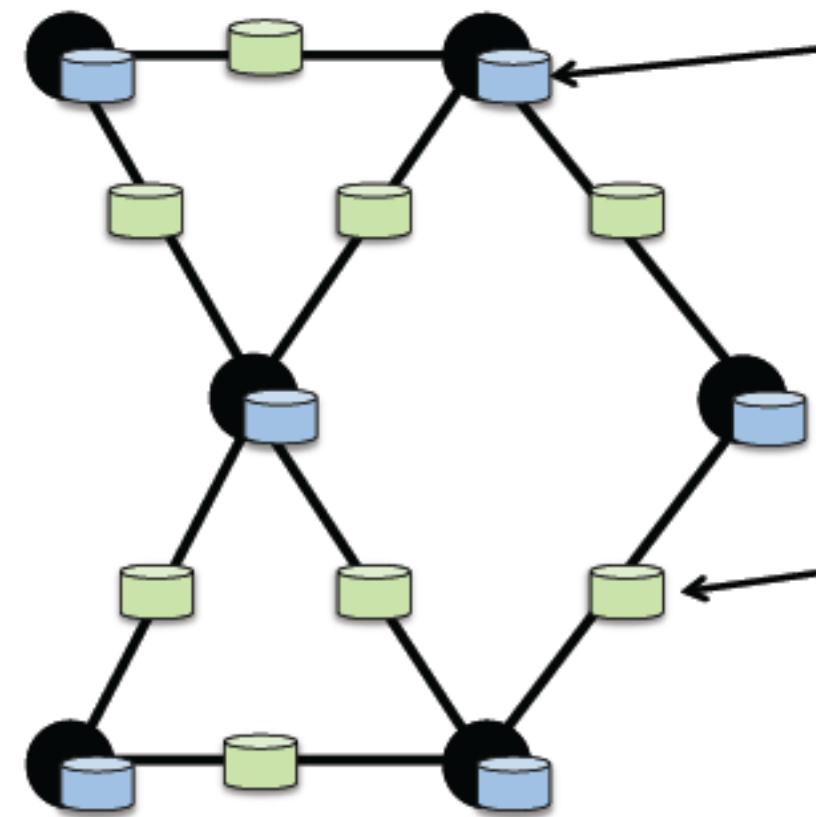


Graph View

Each view has its own operators that
exploit the semantics of the view
to achieve efficient execution

The GraphX API

Property Graphs



Vertex Property:

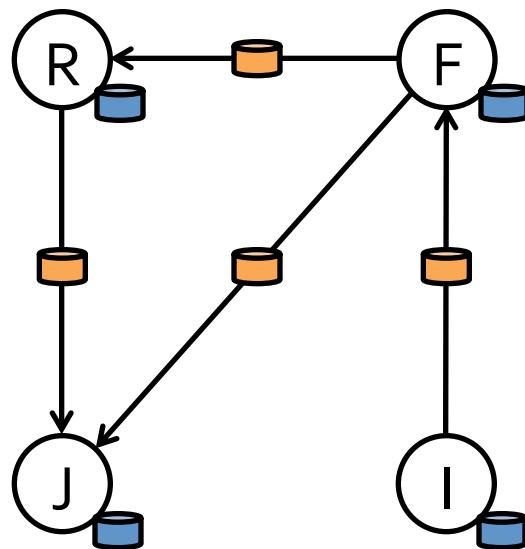
- User Profile
- Current PageRank Value

Edge Property:

- Weights
- Relationships
- Timestamps

View a Graph as a Table

Property Graph



Vertex Property Table

Id	Property (V)
Rxin	(Stu., Berk.)
Jegonzal	(PstDoc, Berk.)
Franklin	(Prof., Berk)
Istoica	(Prof., Berk)

Edge Property Table

SrcId	DstId	Property (E)
rxin	jegonzal	Friend
franklin	rxin	Advisor
istoica	franklin	Coworker
franklin	jegonzal	PI

Table Operators

Table (RDD) operators are inherited from Spark:

map	reduce	sample
filter	count	take
groupBy	fold	first
sort	reduceByKey	partitionBy
union	groupByKey	mapwith
join	cogroup	pipe
leftOuterJoin	cross	save
rightOuterJoin	zip	...

Creating a Graph (Scala)

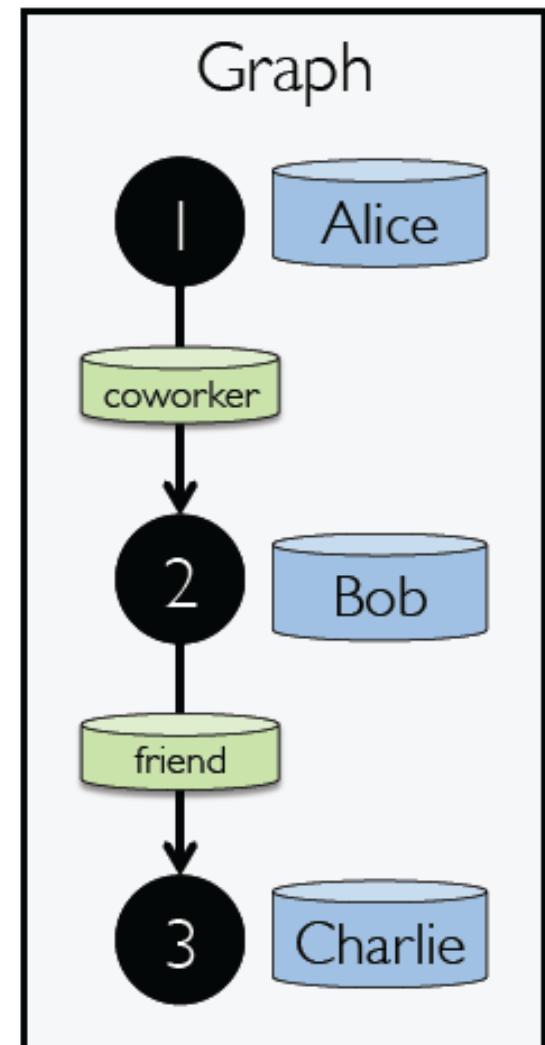
```
type VertexId = Long

val vertices: RDD[(VertexId, String)] =
  sc.parallelize(List(
    (1L, "Alice"),
    (2L, "Bob"),
    (3L, "Charlie")))

class Edge[ED](
  val srcId: VertexId,
  val dstId: VertexId,
  val attr: ED)

val edges: RDD[Edge[String]] =
  sc.parallelize(List(
    Edge(1L, 2L, "coworker"),
    Edge(2L, 3L, "friend")))

val graph = Graph(vertices, edges)
```



Graph Operations (Scala)

```
/** Summary of the functionality in the property graph */
class Graph[VD, ED] {
    // Information about the Graph =====
    =====
    val numEdges: Long
    val numVertices: Long
    val inDegrees: VertexRDD[Int]
    val outDegrees: VertexRDD[Int]
    val degrees: VertexRDD[Int]
    // Views of the graph as collections =====
    =====
    val vertices: VertexRDD[VD]
    val edges: EdgeRDD[ED]
    val triplets: RDD[EdgeTriplet[VD, ED]]
    // Functions for caching graphs =====
    =====
    def persist(newLevel: StorageLevel = StorageLevel.MEMORY_ONLY): Graph[VD, ED]
    def cache(): Graph[VD, ED]
    def unpersistVertices(blocking: Boolean = true): Graph[VD, ED]
    // Change the partitioning heuristic =====
    =====
    def partitionBy(partitionStrategy: PartitionStrategy): Graph[VD, ED]
```

Graph Operations (Scala)

```
// Transform vertex and edge attributes =====
=====
def mapVertices[VD2](map: (VertexID, VD) => VD2): Graph[VD2, ED]
def mapEdges[ED2](map: Edge[ED] => ED2): Graph[VD, ED2]
def mapEdges[ED2](map: (PartitionID, Iterator[Edge[ED]]) => Iterator[ED2]): Graph[VD
, ED2]
def mapTriplets[ED2](map: EdgeTriplet[VD, ED] => ED2): Graph[VD, ED2]
def mapTriplets[ED2](map: (PartitionID, Iterator[EdgeTriplet[VD, ED]]) => Iterator[E
D2])
: Graph[VD, ED2]
// Modify the graph structure =====
=====
def reverse: Graph[VD, ED]
def subgraph(
  epred: EdgeTriplet[VD, ED] => Boolean = (x => true),
  vpred: (VertexID, VD) => Boolean = ((v, d) => true))
: Graph[VD, ED]
def mask[VD2, ED2](other: Graph[VD2, ED2]): Graph[VD, ED]
def groupEdges(merge: (ED, ED) => ED): Graph[VD, ED]
```

```

// Join RDDs with the graph =====
=====
def joinVertices[U](table: RDD[(VertexID, U)])(mapFunc: (VertexID, VD, U) => VD): Graph[VD, ED]

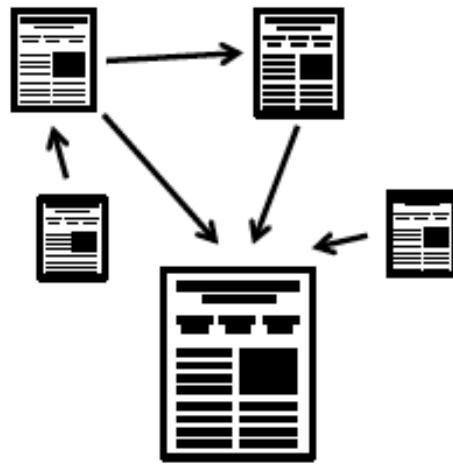
def outerJoinVertices[U, VD2](other: RDD[(VertexID, U)])
  (mapFunc: (VertexID, VD, Option[U]) => VD2)
  : Graph[VD2, ED]
// Aggregate information about adjacent triplets =====
=====
def collectNeighborIds(edgeDirection: EdgeDirection): VertexRDD[Array[VertexID]]
def collectNeighbors(edgeDirection: EdgeDirection): VertexRDD[Array[(VertexID, VD)]]
def aggregateMessages[Msg: ClassTag](
  sendMsg: EdgeContext[VD, ED, Msg] => Unit,
  mergeMsg: (Msg, Msg) => Msg,
  tripletFields: TripletFields = TripletFields.All)
  : VertexRDD[A]
// Iterative graph-parallel computation =====
=====
def pregel[A](initialMsg: A, maxIterations: Int, activeDirection: EdgeDirection)(
  vprog: (VertexID, VD, A) => VD,
  sendMsg: EdgeTriplet[VD, ED] => Iterator[(VertexID,A)],
  mergeMsg: (A, A) => A)
  : Graph[VD, ED]
// Basic graph algorithms =====
=====
def pageRank(tol: Double, resetProb: Double = 0.15): Graph[Double, Double]
def connectedComponents(): Graph[VertexID, ED]
def triangleCount(): Graph[Int, ED]
def stronglyConnectedComponents(numIter: Int): Graph[VertexID, ED]
}

```

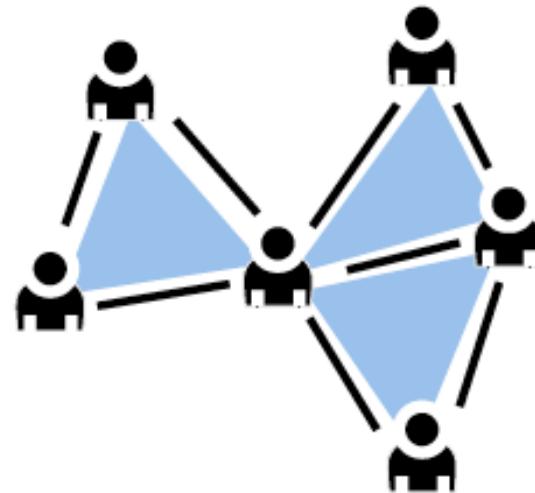
Built-in Algorithms (Scala)

```
// Continued from previous slide
def pageRank(tol: Double): Graph[Double, Double]
def triangleCount(): Graph[Int, ED]
def connectedComponents(): Graph[VertexId, ED]
// ...and more: org.apache.spark.graphx.lib
}
```

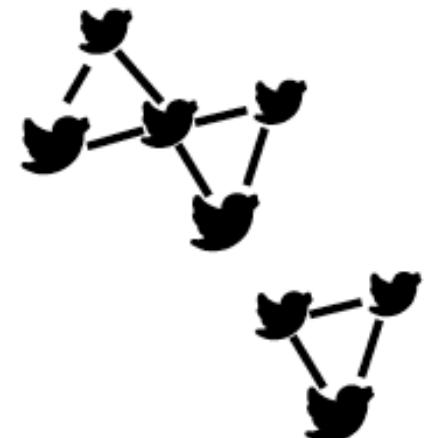
PageRank



Triangle Count



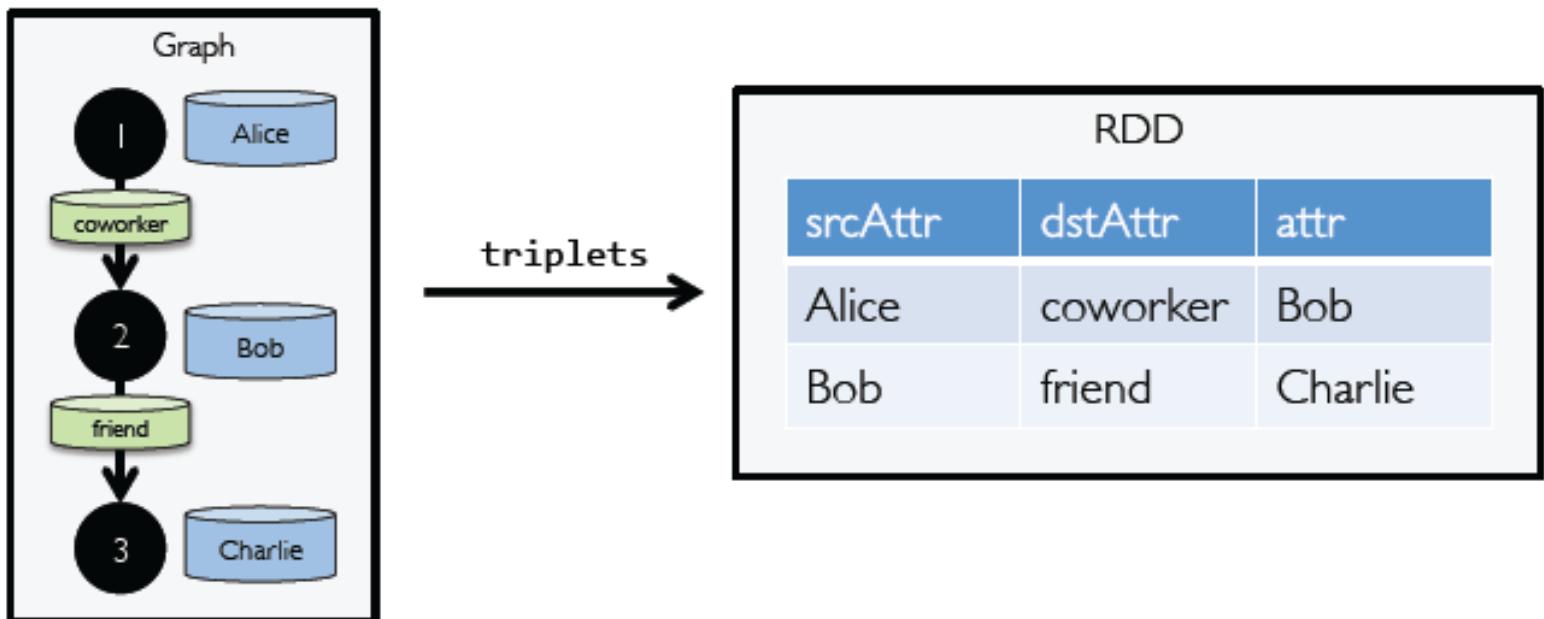
Connected Components



The “triplets” view

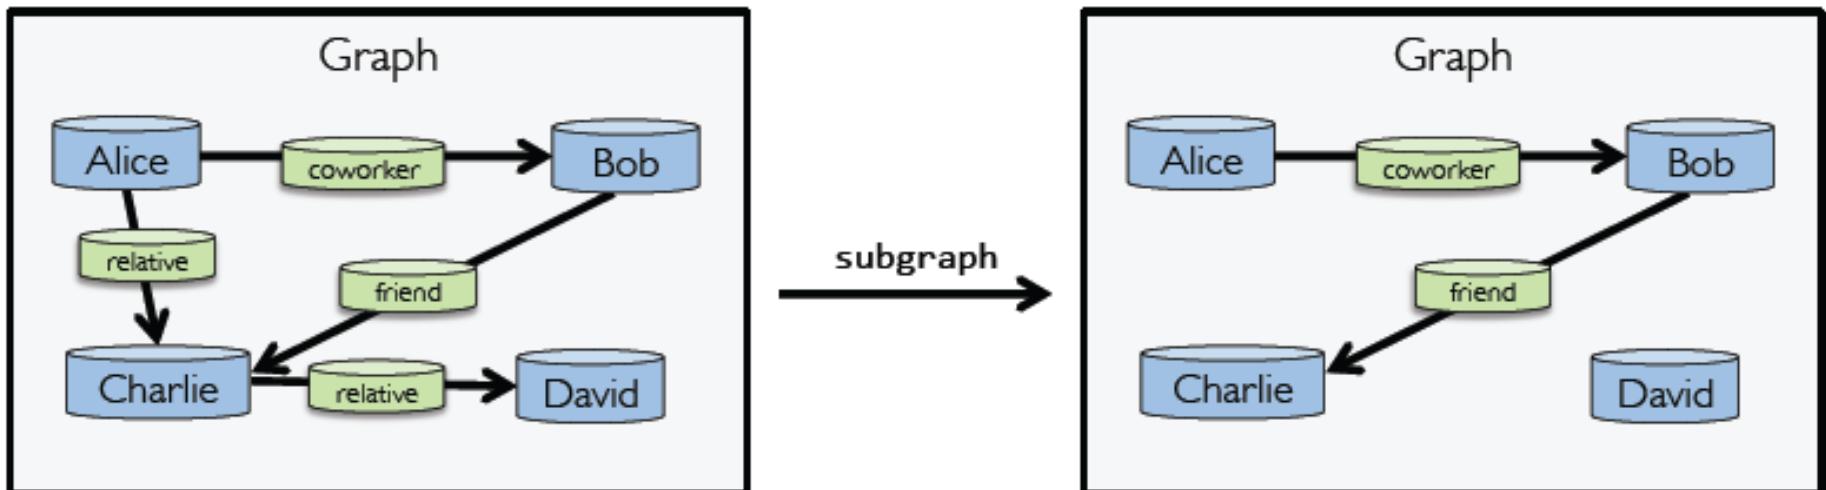
```
class Graph[VD, ED] {  
    def triplets: RDD[EdgeTriplet[VD, ED]]  
}
```

```
class EdgeTriplet[VD, ED](  
    val srcId: VertexId, val dstId: VertexId, val attr: ED,  
    val srcAttr: VD, val dstAttr: VD)
```



The subgraph transformation

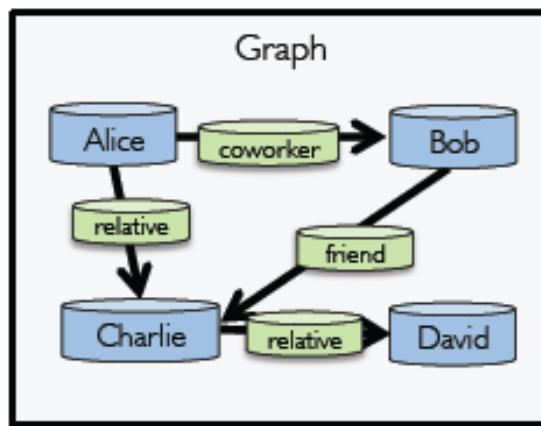
```
class Graph[VD, ED] {  
    def subgraph(epred: EdgeTriplet[VD, ED] => Boolean,  
                vpred: (VertexId, VD) => Boolean): Graph[VD, ED]  
}  
  
graph.subgraph(epred = (edge) => edge.attr != "relative")
```



Computation w/ mapReduceTriplets

```
class Graph[VD, ED] {  
    def mapReduceTriplets[A](  
        sending: Function[(VertexId, VD), RDD[(VertexId, A)]],  
        merging: (A, A) => A) RDD[(VertexId, A)]  
    }  
  
    ! upgrade to aggregateMessages  
    in Spark 1.2.0
```

```
graph.mapReduceTriplets(  
    edge => Iterator(  
        (edge.srcId, 1),  
        (edge.dstId, 1)),  
    _ + _)
```



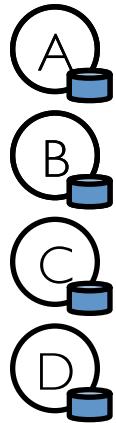
mapReduceTriplets →

RDD	
vertex id	degree
Alice	2
Bob	2
Charlie	3
David	1

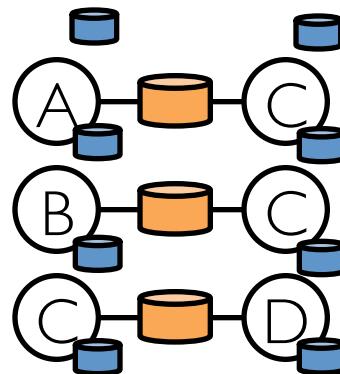
Triplets Join Vertices and Edges

The *triplets* operator joins vertices and edges:

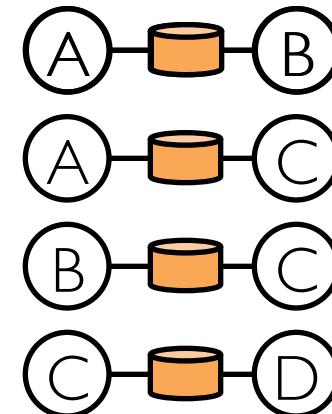
Vertices



Triplets



Edges



The *mapreduceTriplets* operator sums adjacent triplets.

```
SELECT t.dstId, reduceUDF( mapUDF(t) ) AS sum  
FROM triplets AS t GROUPBY t.dstId
```

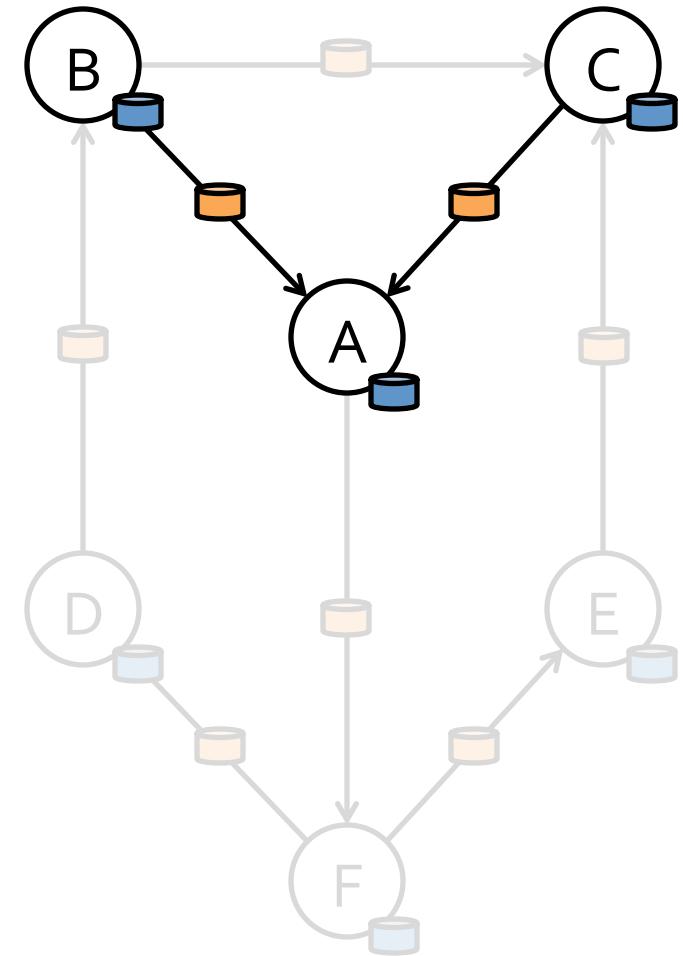
Map Reduce Triplets

Map-Reduce for each vertex

mapF($(A \leftarrow B)$) $\rightarrow A_1$

mapF($(A \leftarrow C)$) $\rightarrow A_2$

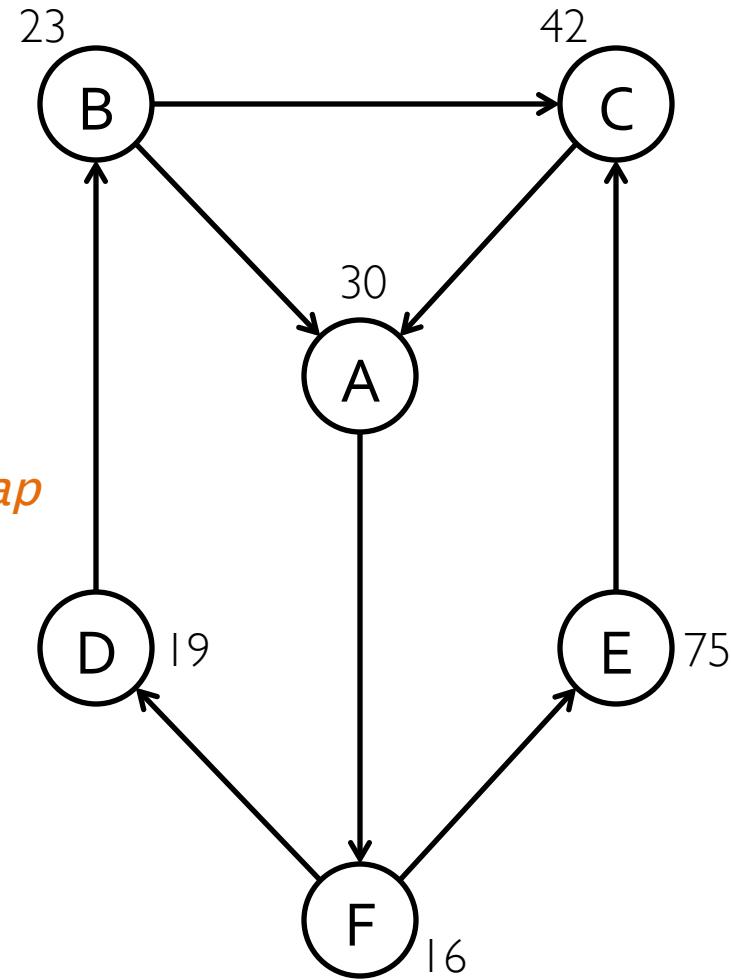
reduceF((A_1, A_2)) $\rightarrow A$



Example: Oldest Follower

What is the age of the oldest follower for each user?

```
val oldestFollowerAge = graph
  .mapreduceTriplets(
    e=> (e.dst.id, e.src.age), //Map
    (a,b)=> max(a, b) //Reduce
  )
  .vertices
```



Have Expressed the Pregel and GraphLab abstractions using the GraphX operators in less than 50 lines of code!

By composing these operators we can construct entire graph-analytics pipelines.

Re-implementation of the Pregel abstraction using the GraphX API

```
def Pregel(g: Graph[V, E],
           vprog: (Id, V, M) => V,
           sendMsg: (Triplet) => M,
           gather: (M, M) => M): Collection[V] = {
    // Set all vertices as active
    g = g.mapV((id, v) => (v, halt=false))
    // Loop until convergence
    while (g.vertices.exists(v => !v.halt)) {
        // Compute the messages
        val msgs: Collection[(Id, M)] =
            // Restrict to edges with active source
            g.subgraph(ePred=(s, d, sP, eP, dP) => !sP.halt)
            // Compute messages
            .mrTriplets(sendMsg, gather)
        // Receive messages and run vertex program
        g = g.leftJoinV(msgs).mapV(vprog)
    }
    return g.vertices
}
```

Finding Connected Components using the GraphX variant of Pregel

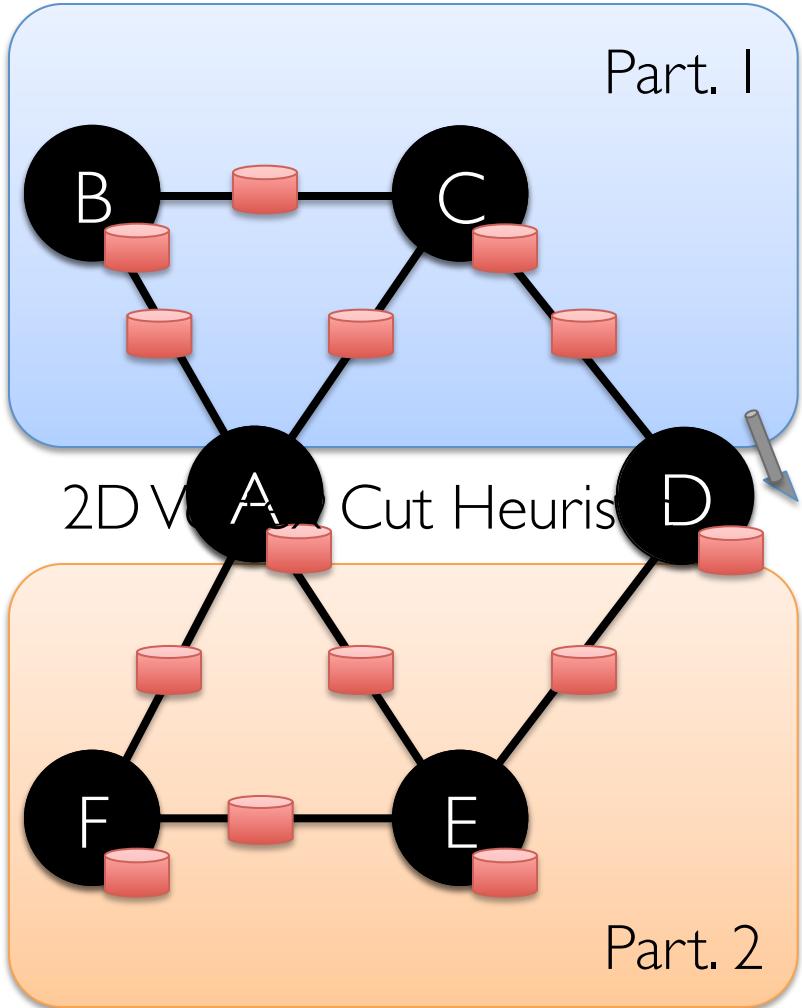
```
def ConnectedComp(g: Graph[V, E]) = {
    g = g.mapV(v => v.id) // Initialize vertices
    def vProg(v: Id, m: Id): Id = {
        if (v == m) voteToHalt(v)
        return min(v, m)
    }
    def sendMsg(t: Triplet): Id =
        if (t.src.cc < t.dst.cc) t.src.cc
        else None // No message required
    def gatherMsg(a: Id, b: Id): Id = min(a, b)
    return Pregel(g, vProg, sendMsg, gatherMsg)
}
```

Listing 6: **Connected Components:** For each vertex we compute the lowest reachable vertex id using Pregel.

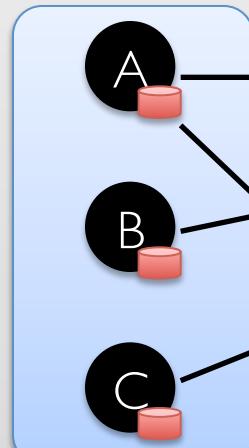
GraphX System Design

Distributed Graphs as Tables (RDDs)

Property Graph



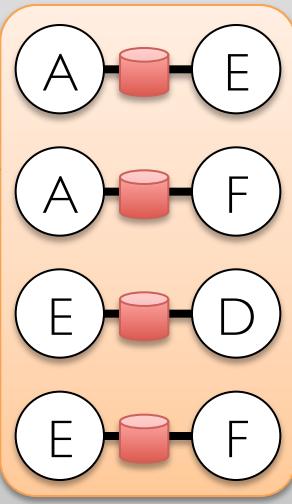
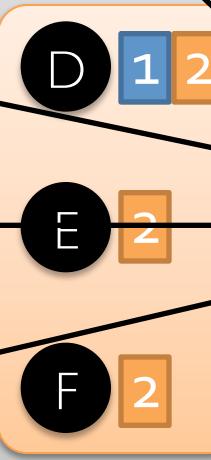
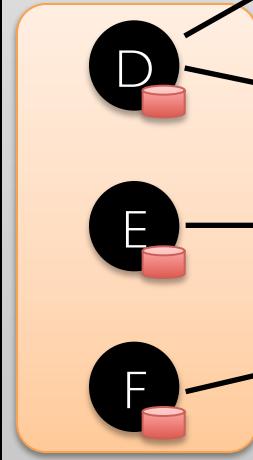
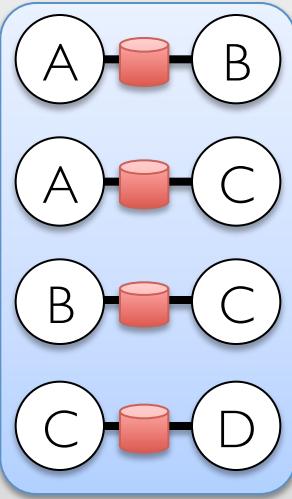
Vertex Table (RDD)



Routing Table (RDD)

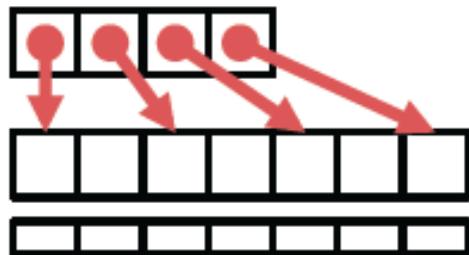


Edge Table (RDD)

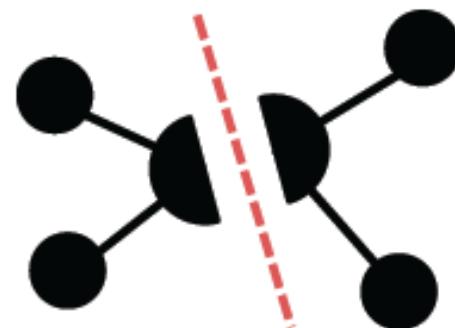


Graph System Optimizations

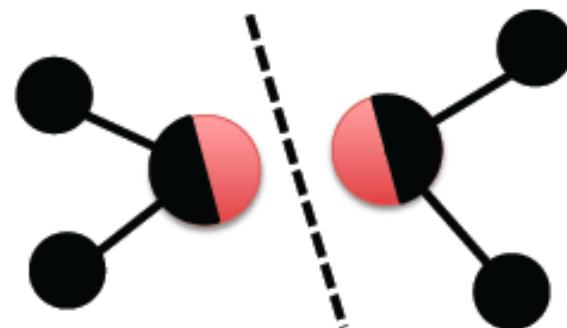
Specialized
Data-Structures



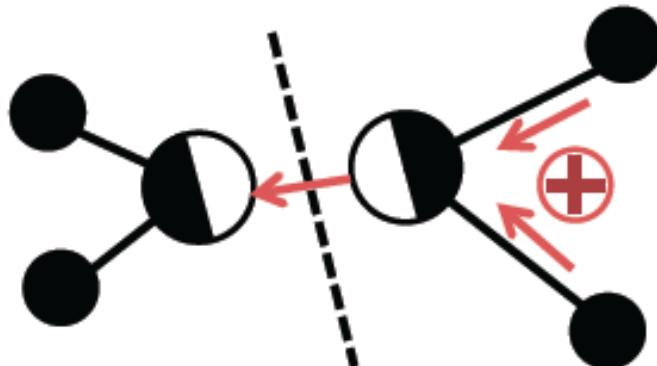
Vertex-Cuts
Partitioning



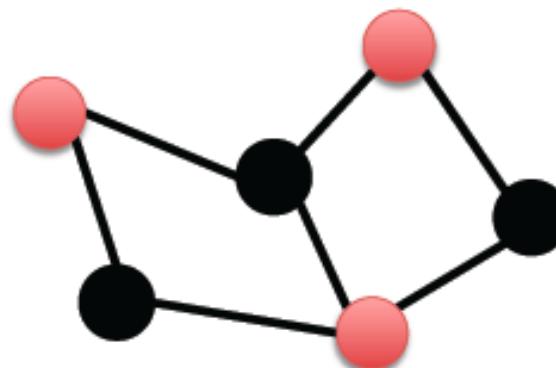
Remote
Caching / Mirroring



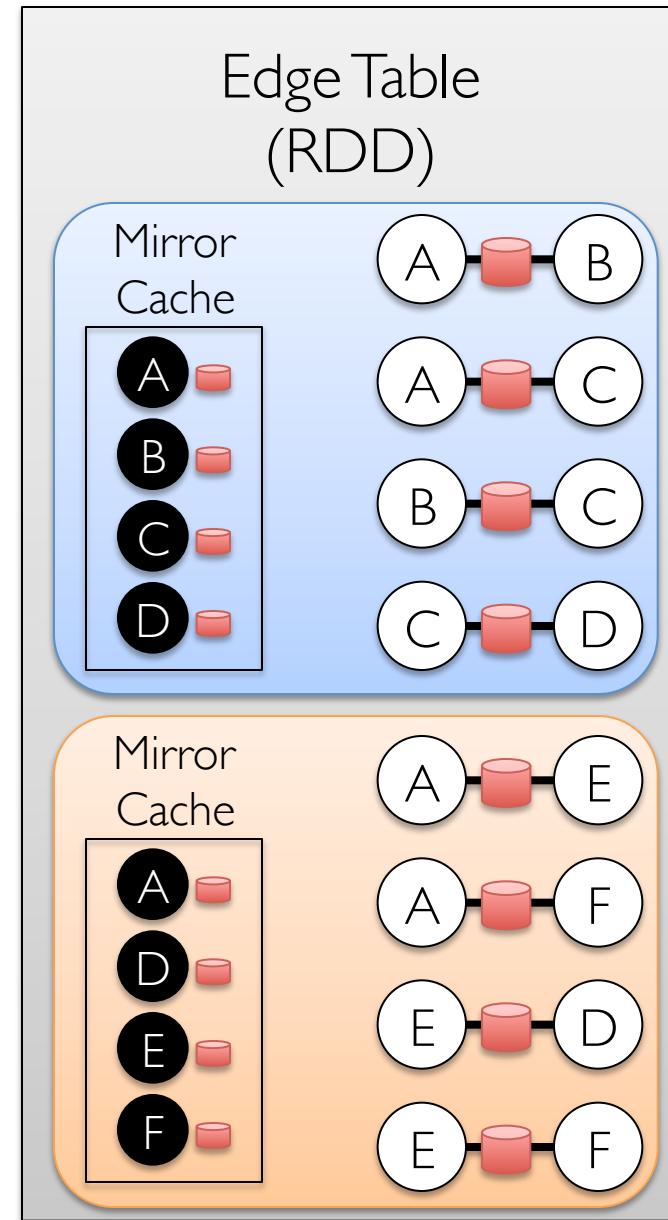
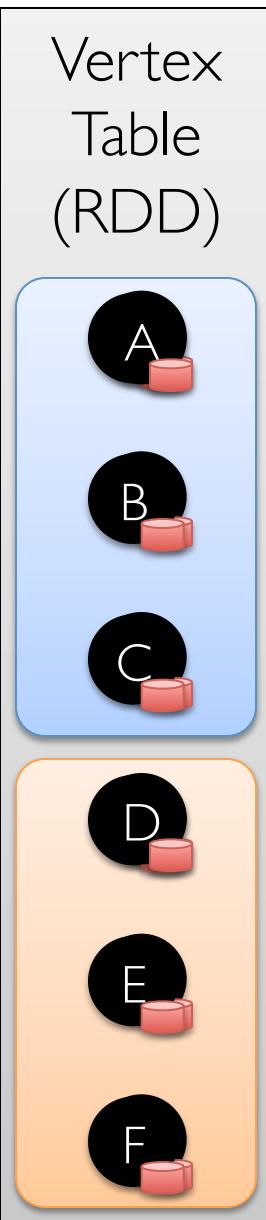
Message Combiners



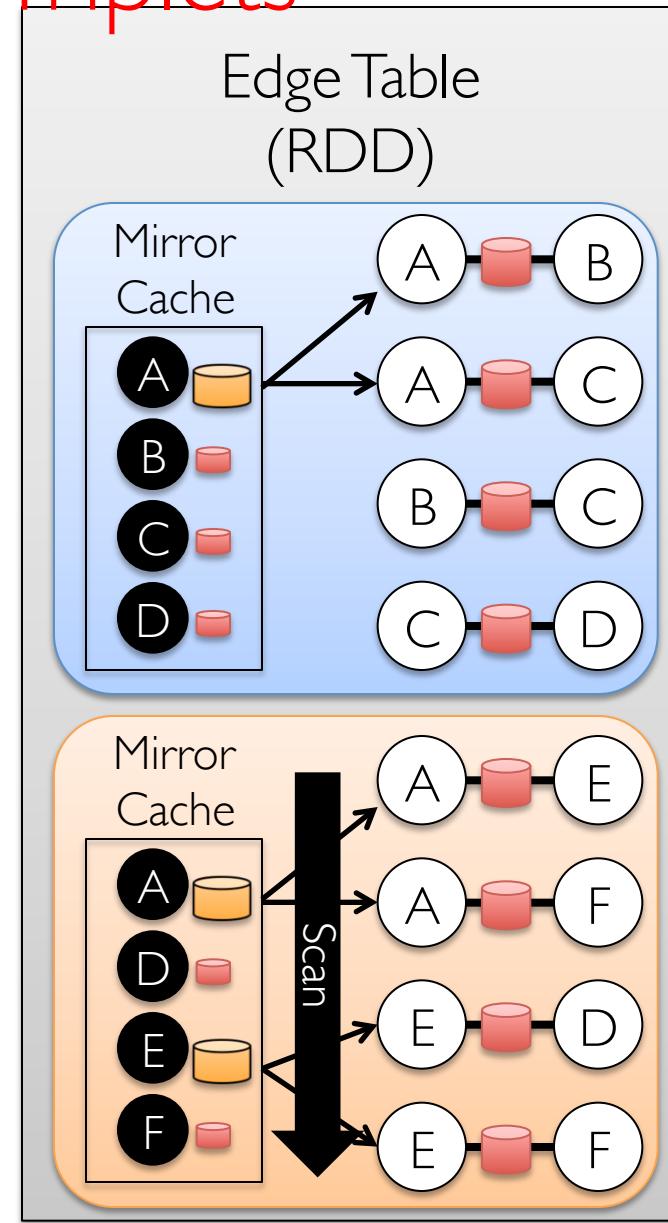
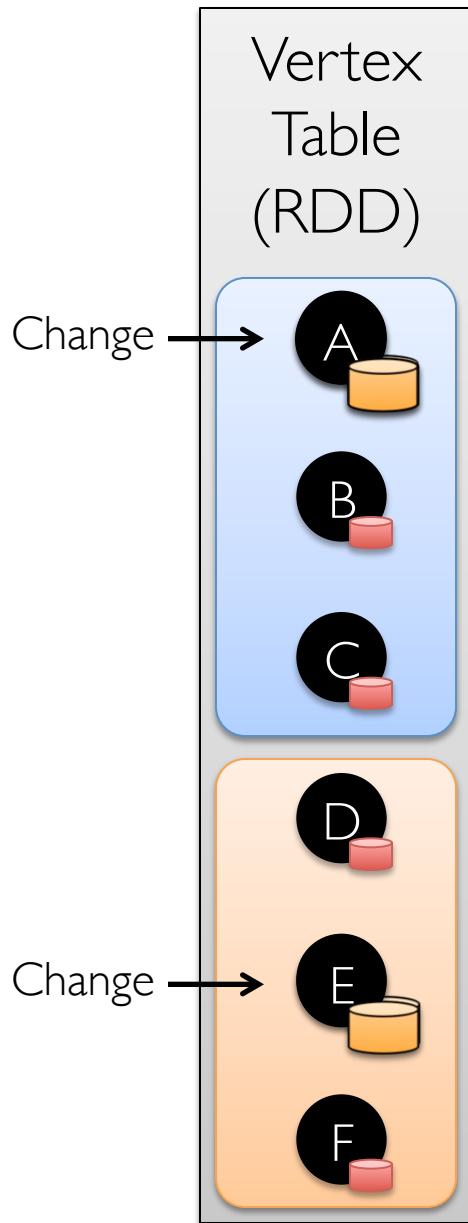
Active Set Tracking



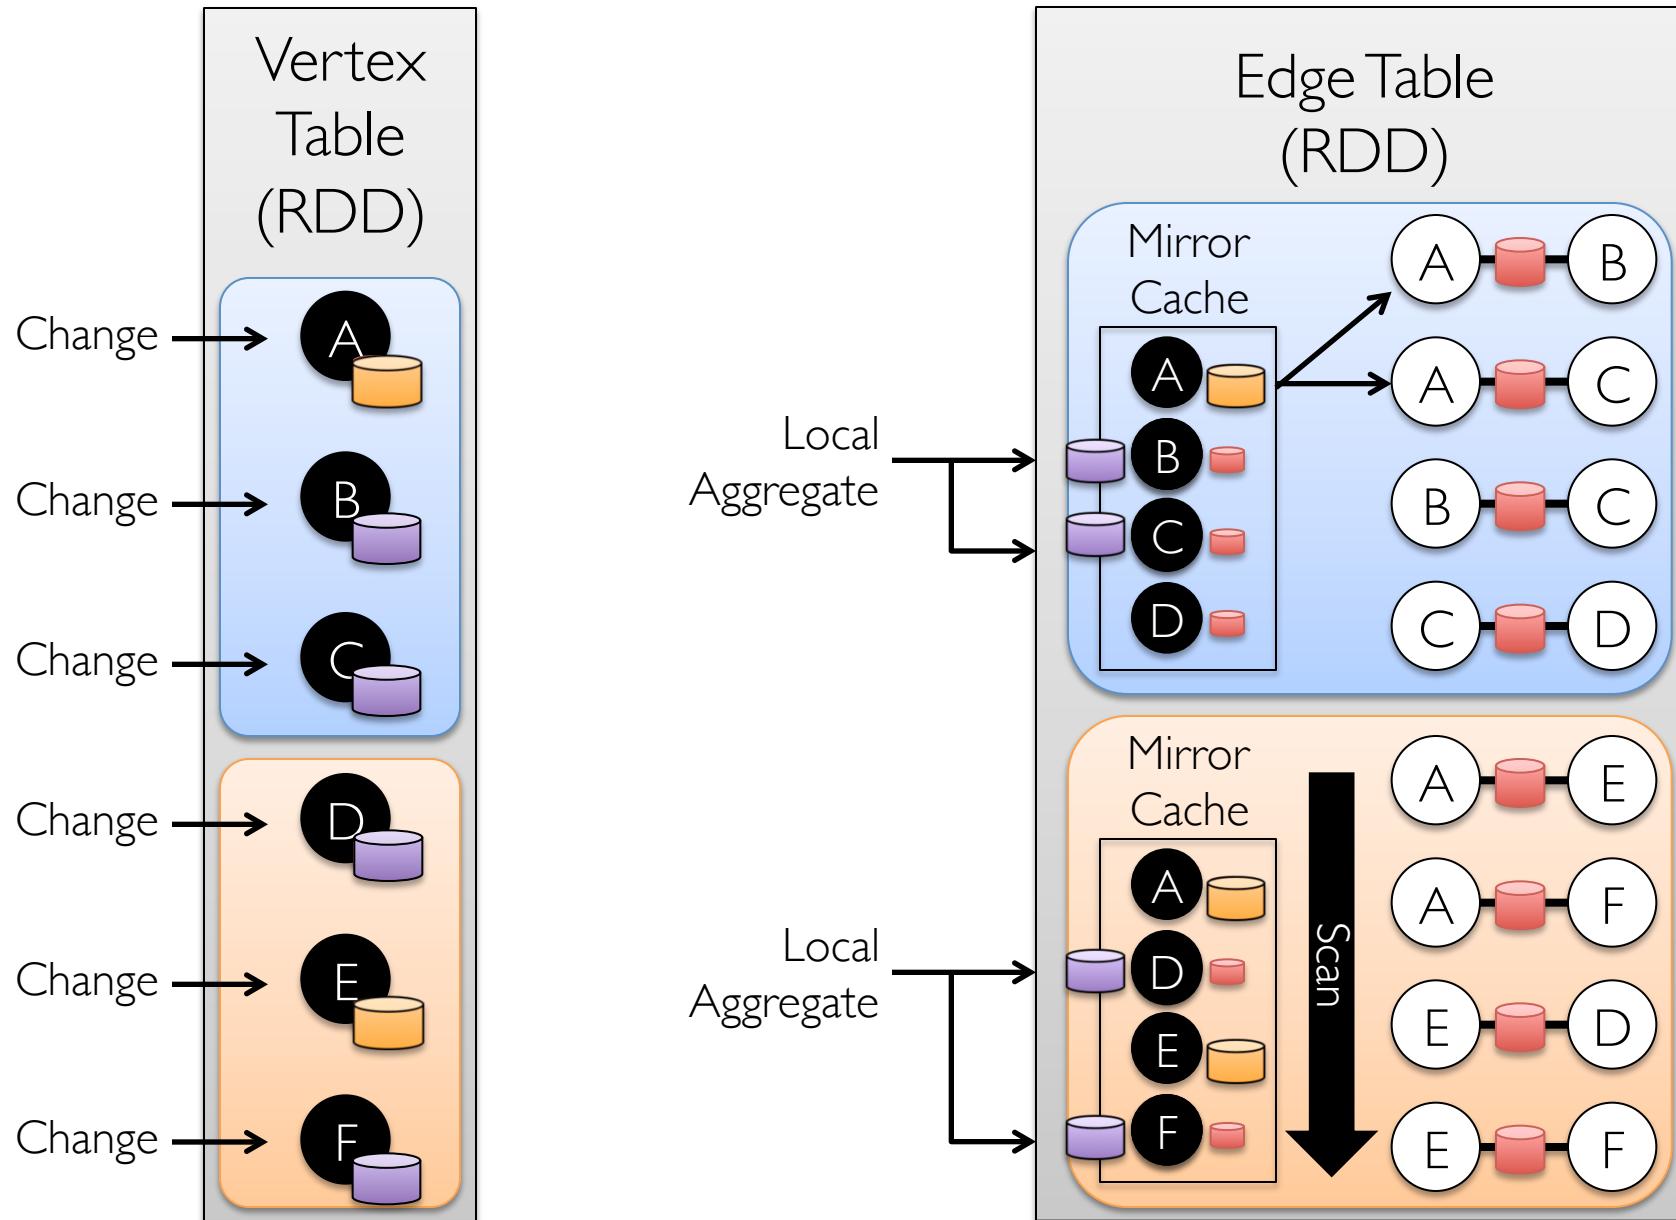
Caching for Iterative mapreduceTriplets



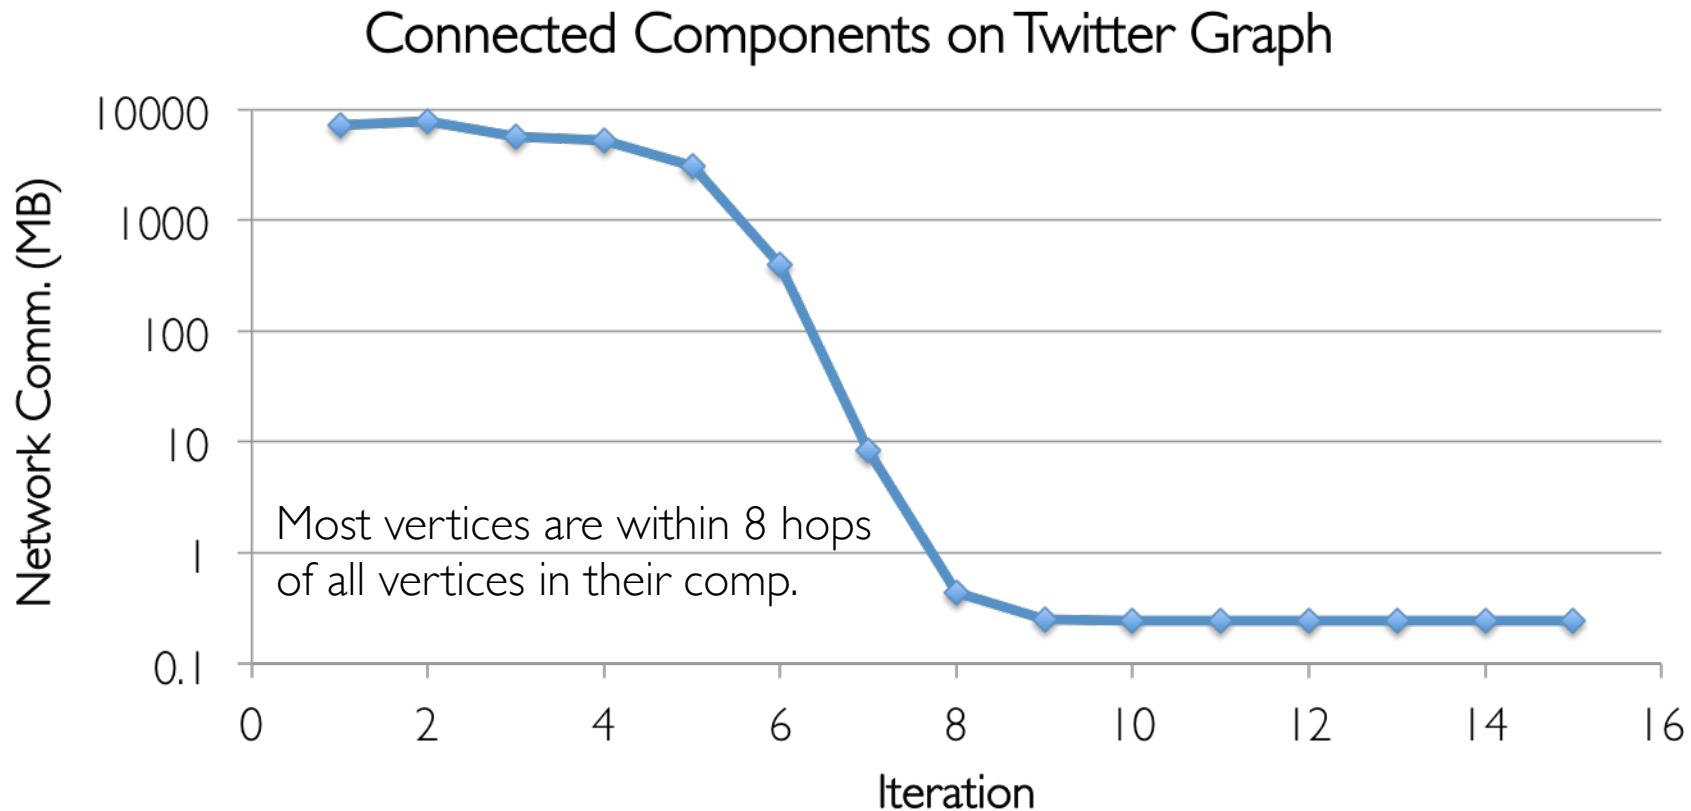
Incremental Updates for Iterative mapreduceTriplets



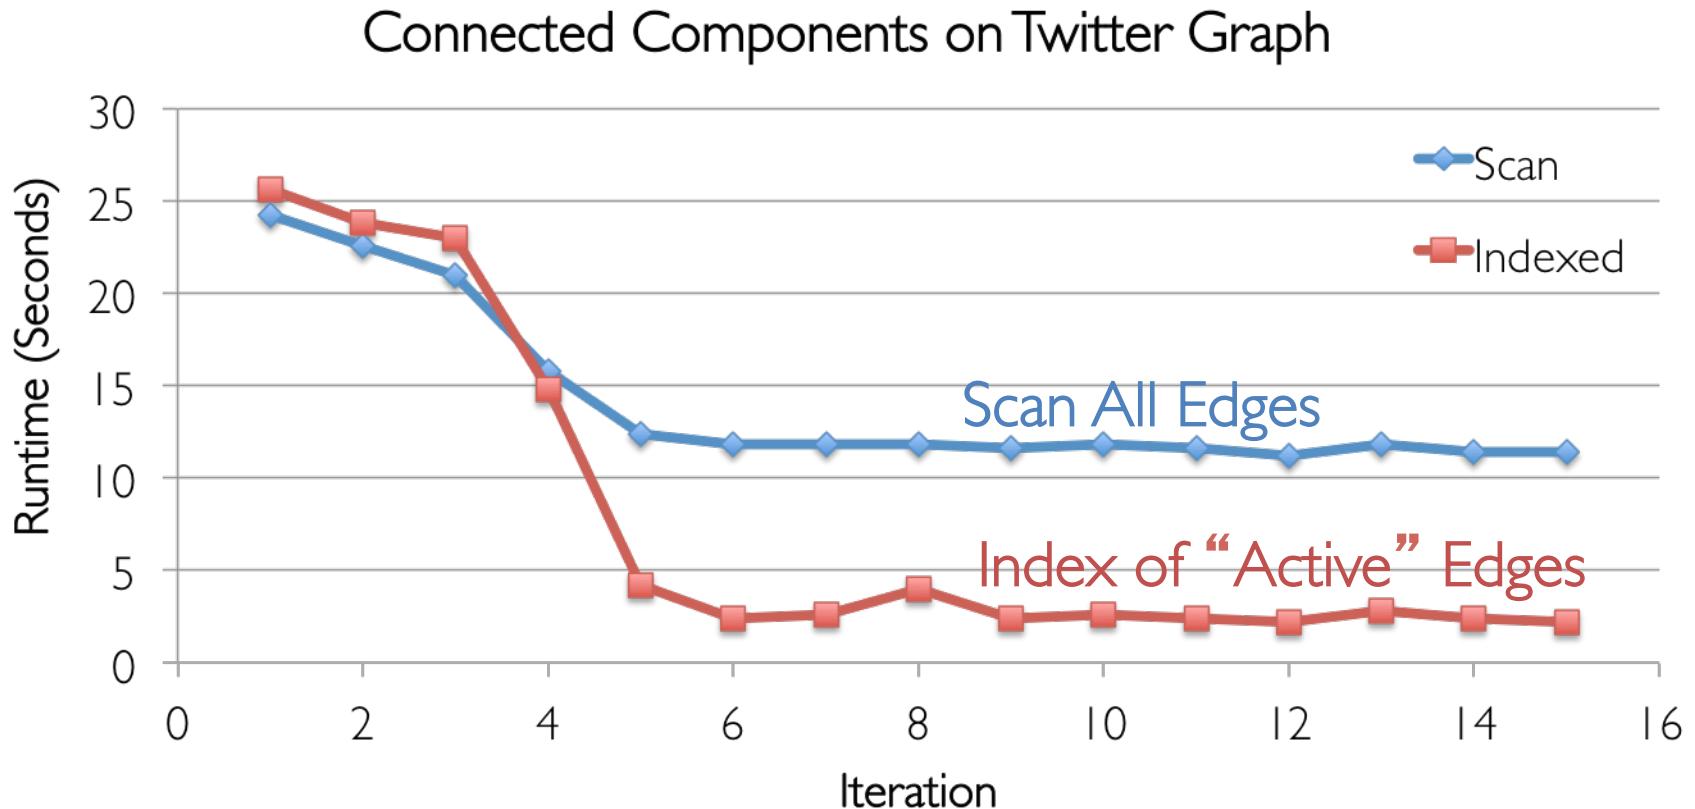
Aggregation for Iterative mapreduceTriplets



Reduction in Communication Due to Cached Updates

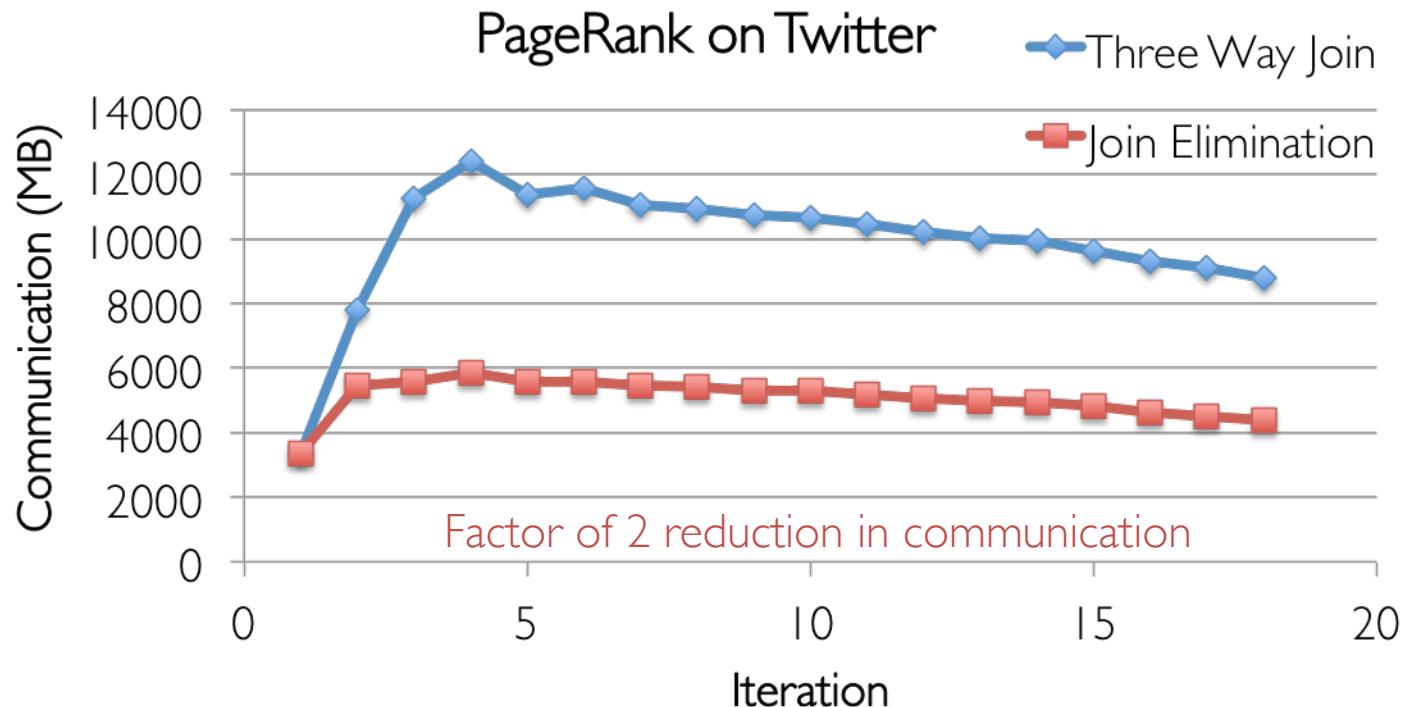


Benefit of Indexing Active Edges



Join Elimination

Identify and bypass joins for unused triplets fields
»Example: PageRank only accesses source attribute



Additional Query Optimizations

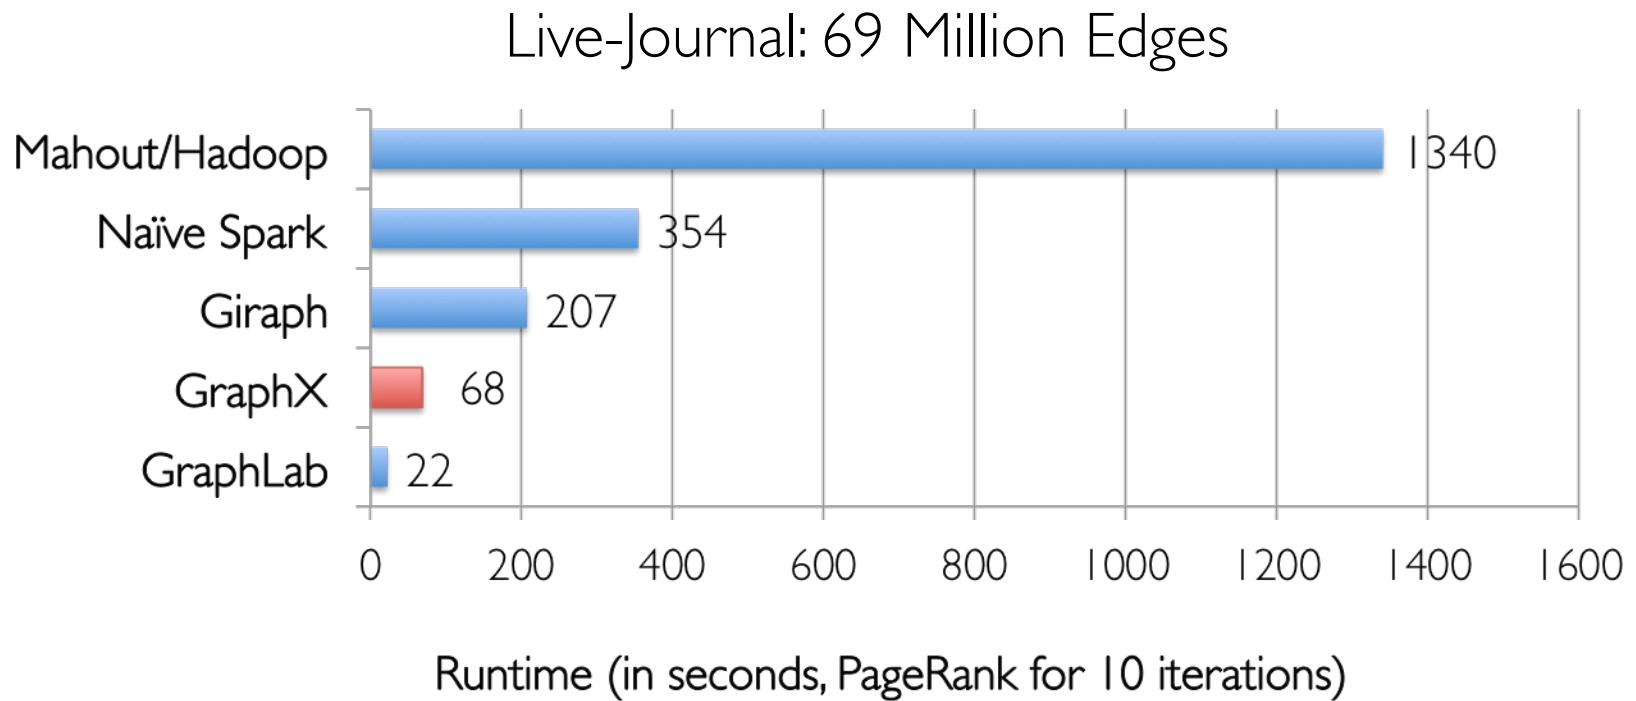
Indexing and Bitmaps:

- » To accelerate joins across graphs
- » To efficiently construct sub-graphs

Substantial Index and Data Reuse:

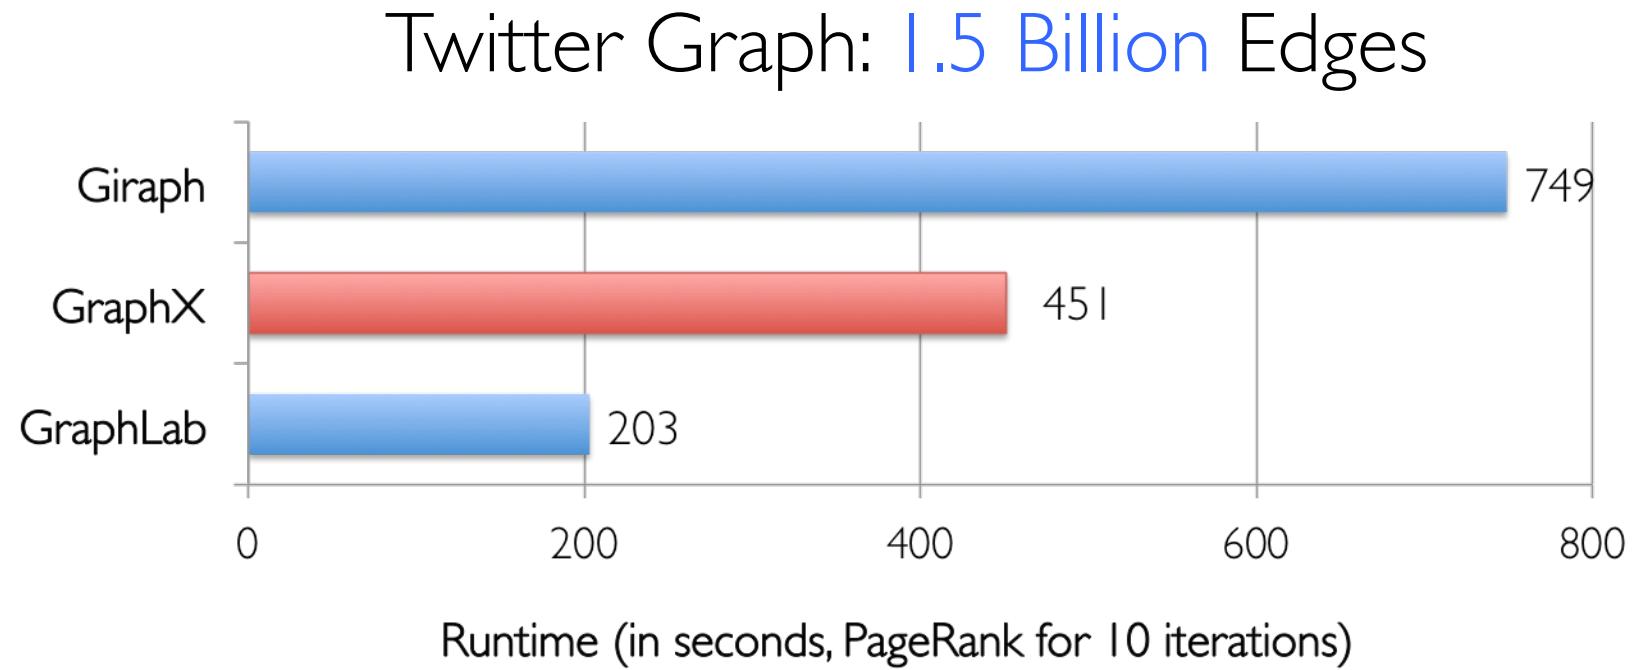
- » Reuse routing tables across graphs and sub-graphs
- » Reuse edge adjacency information and indices

Performance Comparisons



GraphX is roughly *3x slower* than GraphLab

GraphX scales to larger graphs



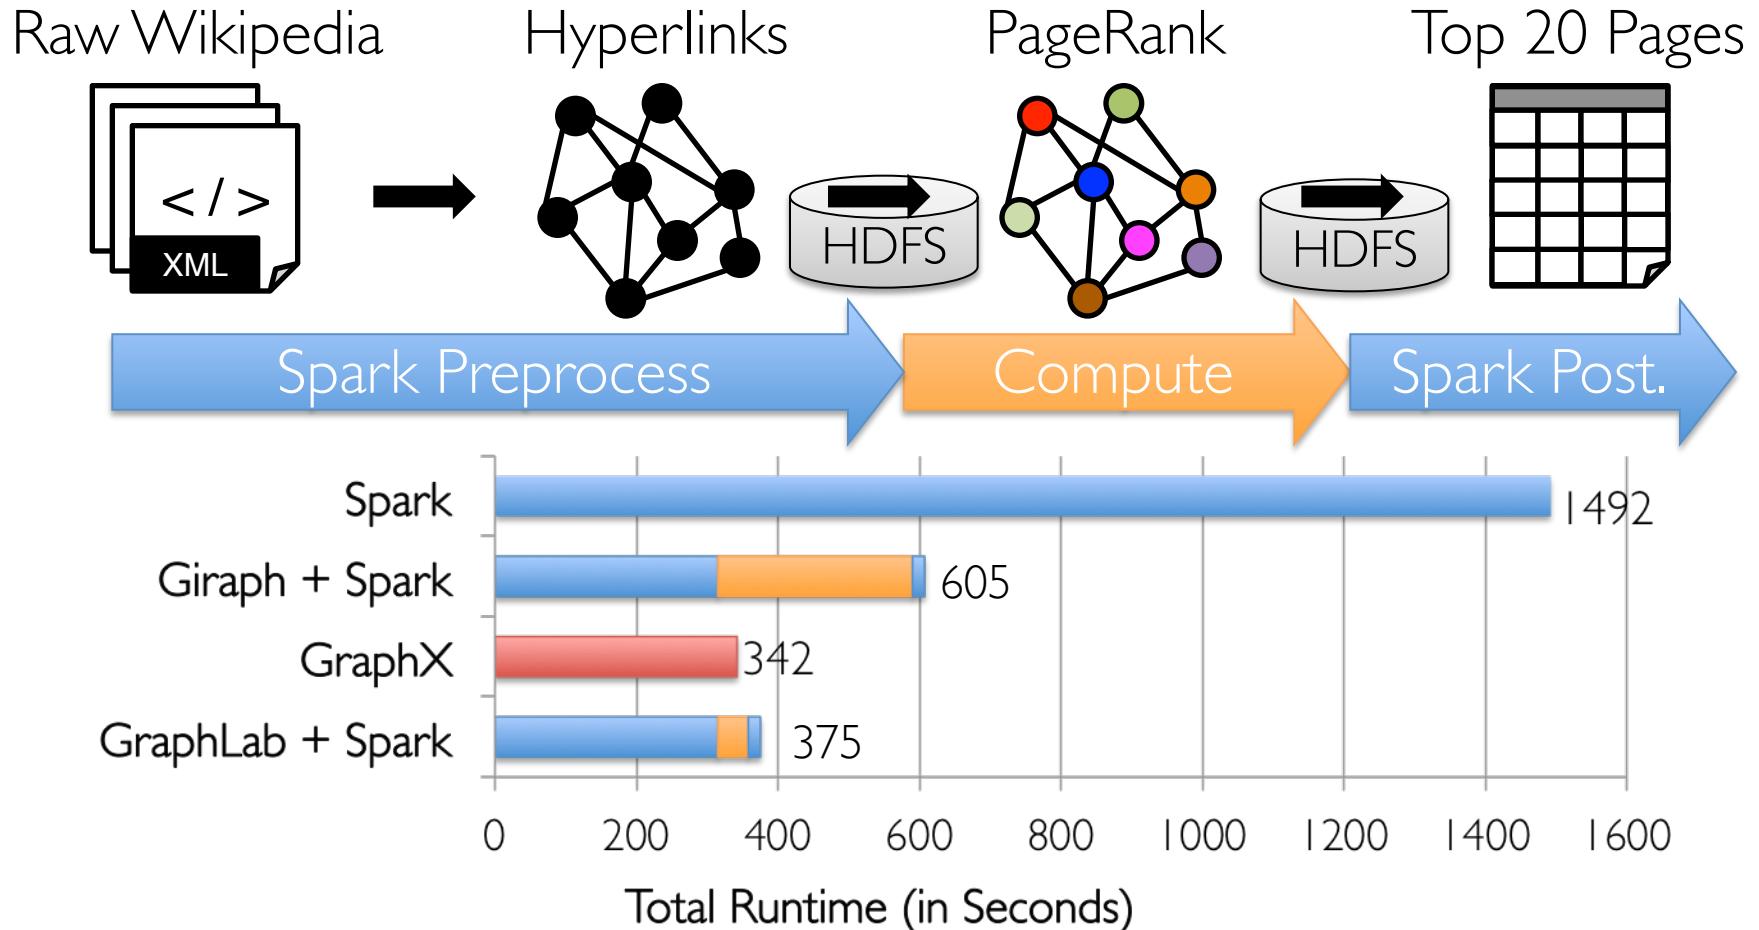
GraphX is roughly *2x slower* than GraphLab

- » Scala + Java overhead: Lambdas, GC time, ...
- » No shared memory parallelism: *2x increase* in comm.

PageRank is just one stage....

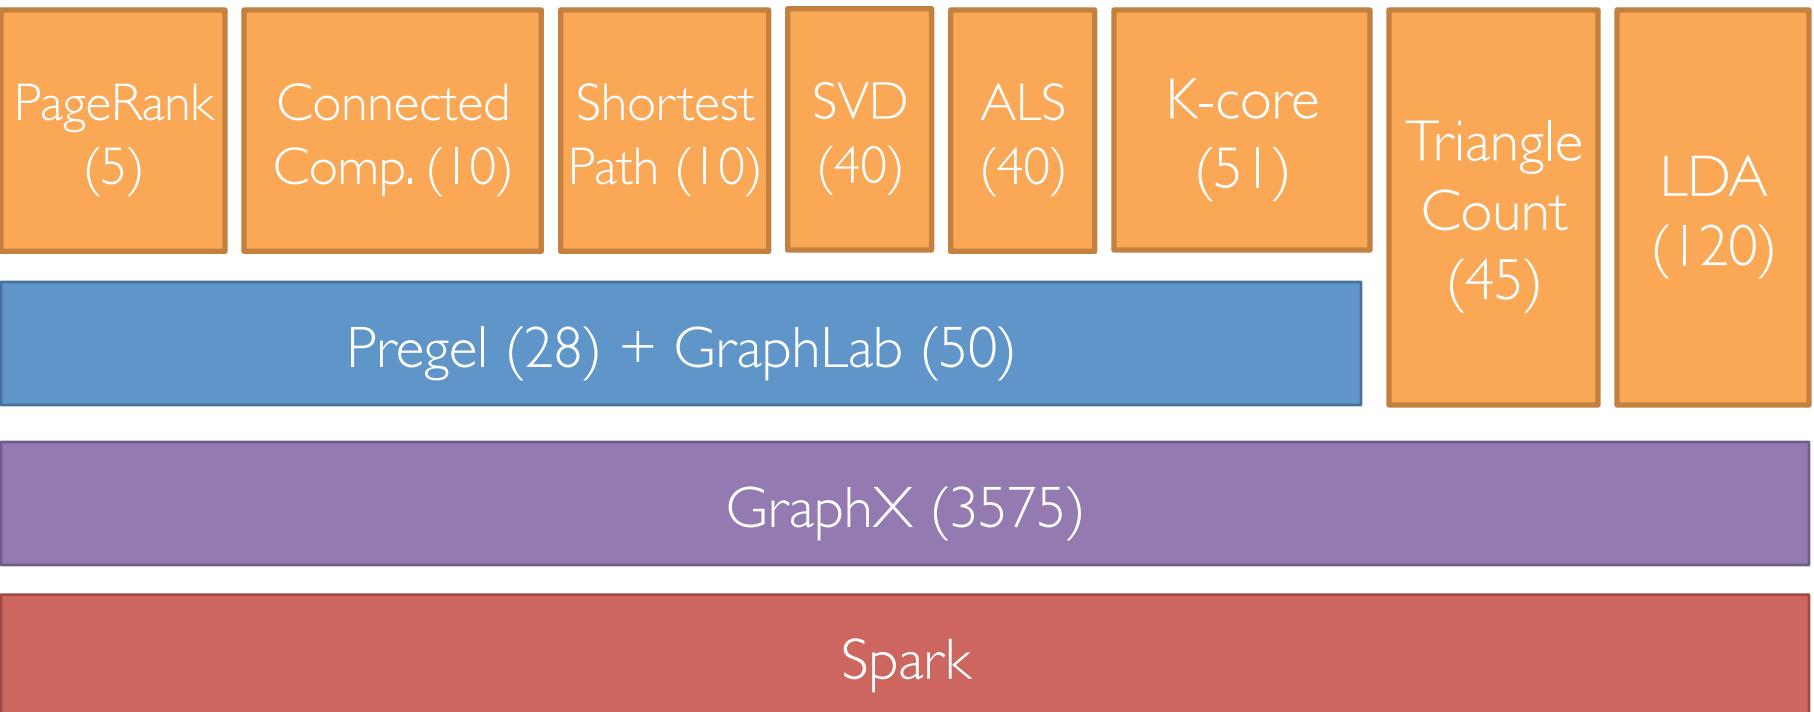
What about a pipeline?

A Small Pipeline in GraphX



Timed end-to-end GraphX is *faster* than GraphLab

The GraphX Stack (Lines of Code)



Conclusion and Observations

Domain specific views: *Tables and Graphs*

- » tables and graphs are first-class composable objects
- » specialized operators which exploit view semantics

Single system that efficiently spans the pipeline

- » minimize data movement and duplication
- » eliminates need to learn and manage multiple systems

Graphs through the lens of database systems

- » Graph-Parallel Pattern → Triplet joins in relational alg.
- » Graph Systems → Distributed join optimizations

Active Research

Static Data → Dynamic Data

- » Apply GraphX unified approach to time evolving data
- » Model and analyze relationships over time

Serving Graph Structured Data

- » Allow external systems to interact with GraphX
- » Unify distributed graph databases with relational database technology

More Info at:

<http://amplab.github.io/graphx/>

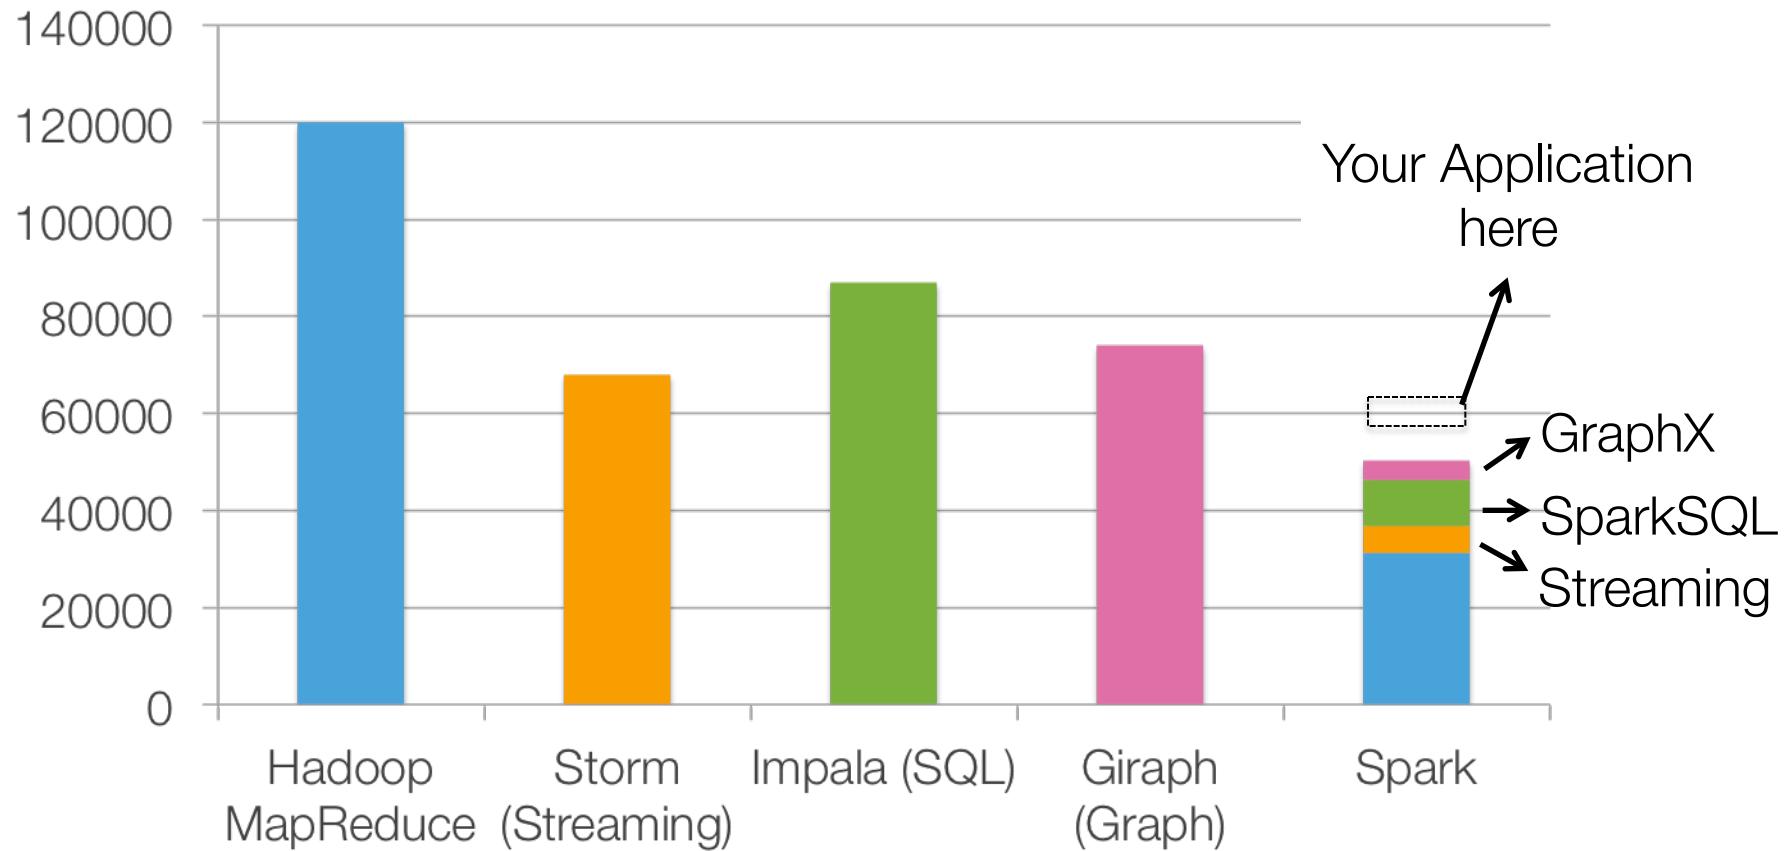
ankurd@eecs.berkeley.edu

crankshaw@eecs.berkeley.edu

rxin@eecs.berkeley.edu

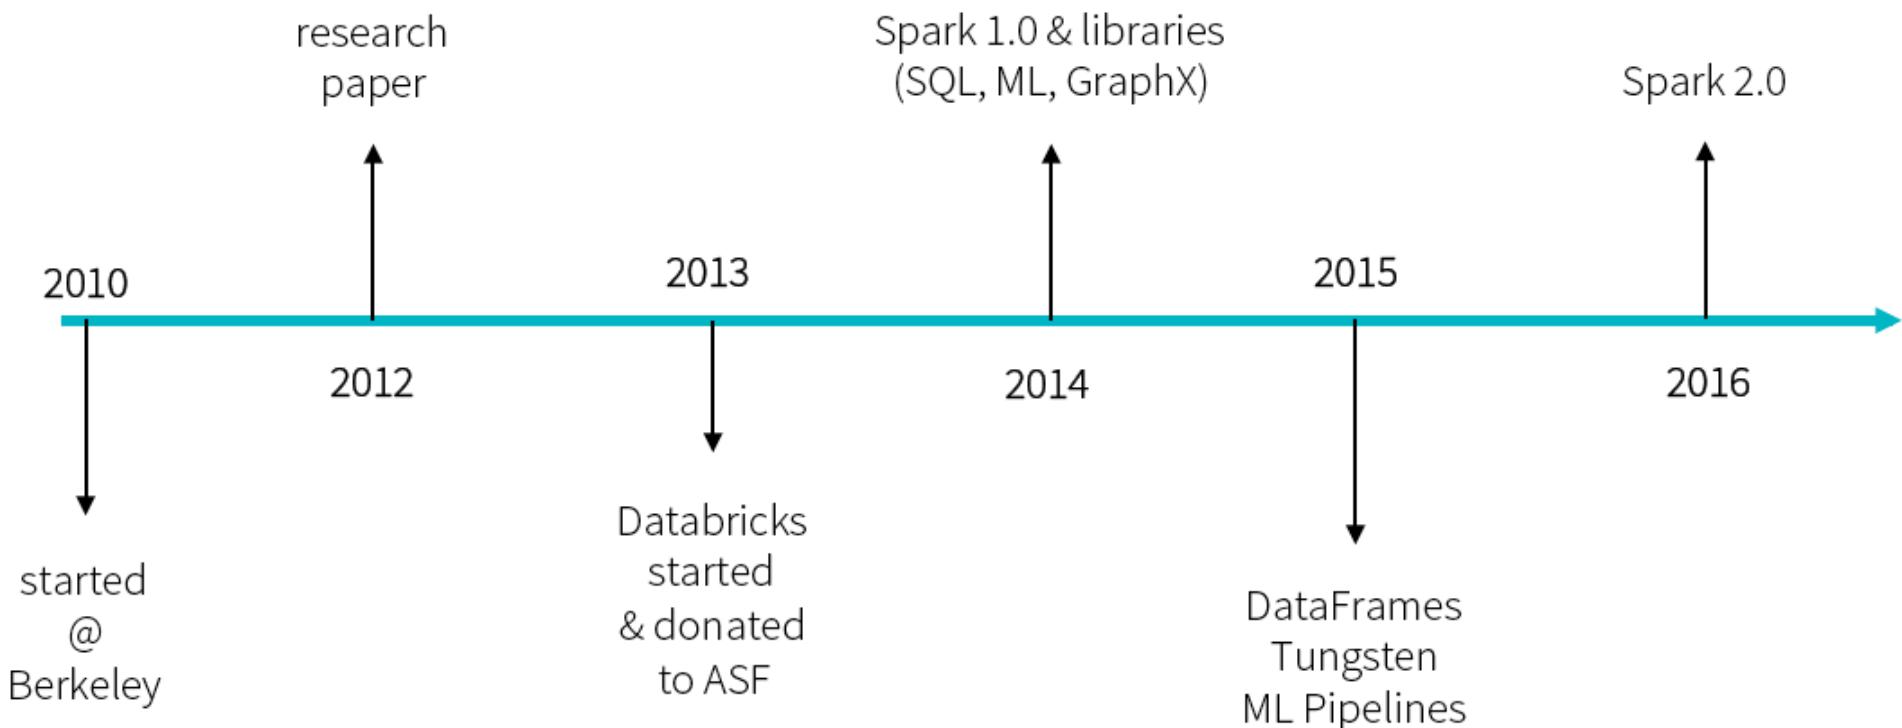
jegonzal@eecs.berkeley.edu

Powerful Stack – Agile Development



non-test, non-example source lines

Spark/ BDAS Timeline so far...



Major Features in Spark 2.0



Tungsten Phase 2
speedups of 5-10x



Structured Streaming
real-time engine
on SQL/DataFrames



Unifying Datasets
and DataFrames

Boosting Spark Performance via Project Tungsten

Goal

To Overcome JVM Performance limitations and bring Spark performance closer to Bare Metal via:

- **Native Memory Management and Binary Processing:** leveraging application semantics to manage memory explicitly and eliminate the overhead of JVM object model and garbage collection
- **Cache-aware computation:** algorithms and data structures to exploit memory hierarchy
- **Runtime Code generation:** using code generation to exploit modern compilers and CPUs

Project Tungsten: Key areas of Optimization

Data
representations

Code Generation

Sorting

Inspired by traditional database systems

Aggregation

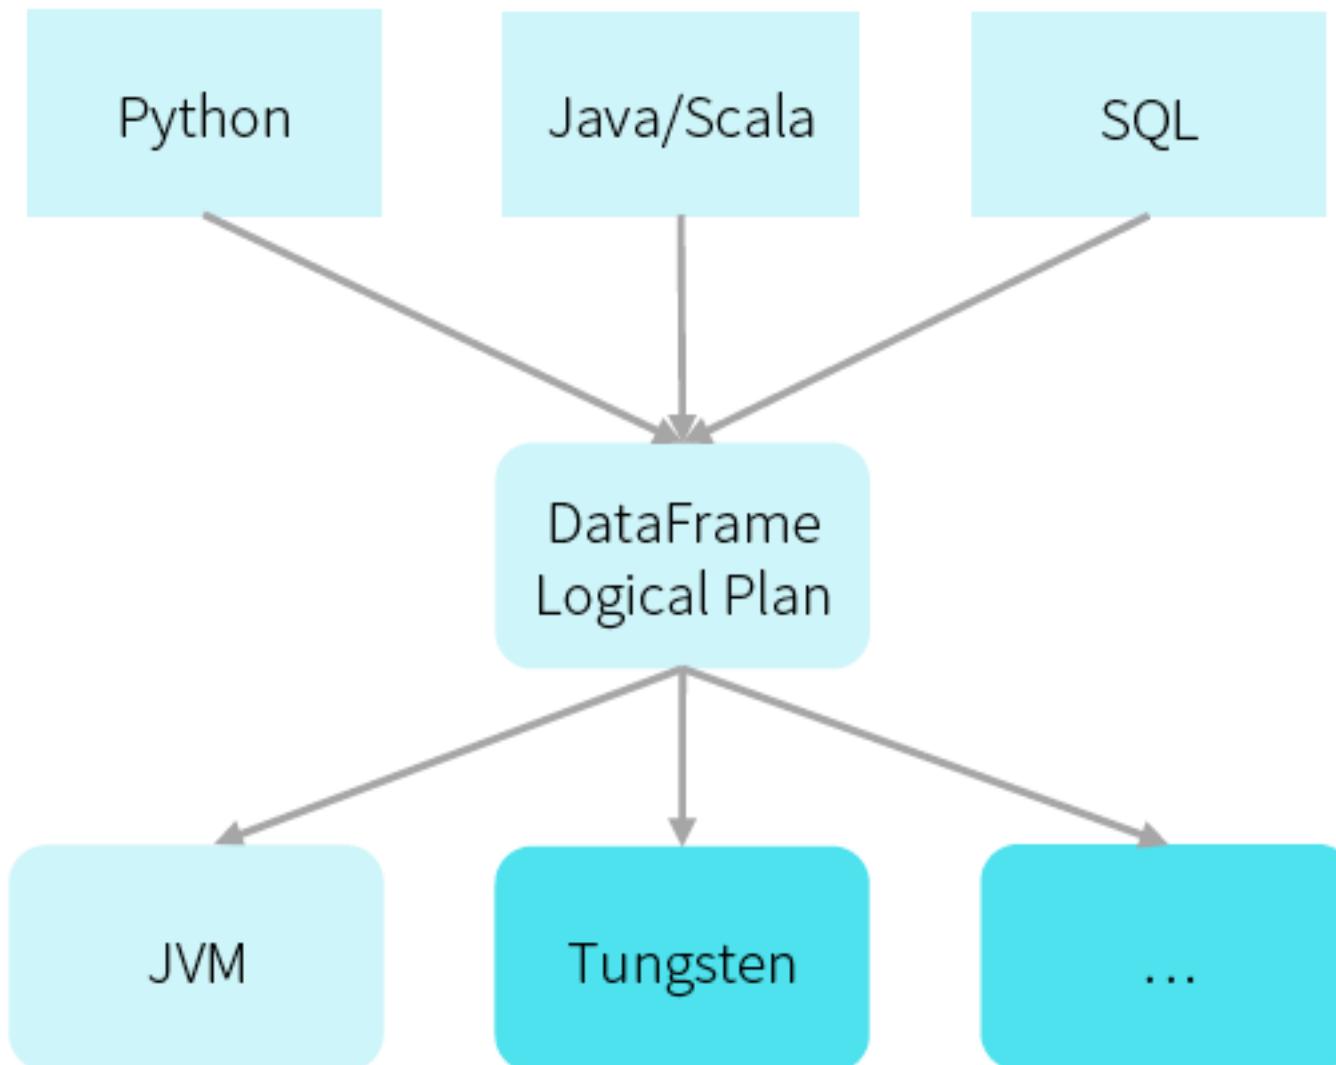
Joins

Broadcasting
and Shuffling

Optimized Data Representations

- Java Objects have two downsides:
 - Space overheads
 - Garbage collection overheads
- Tungsten sidesteps these problems by performing its own manual memory management

Further Performance Optimization via Project Tungsten



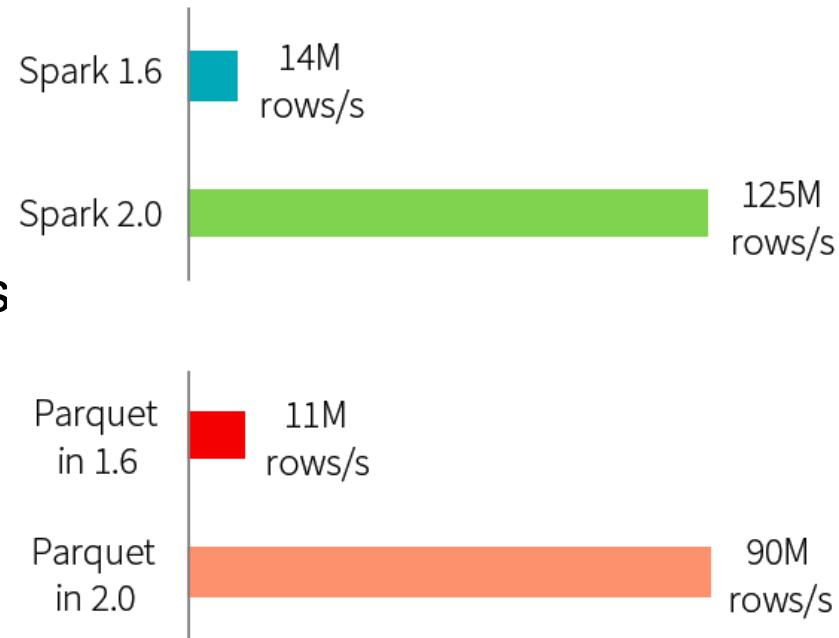
Status of Tungsten (as of Feb 2016)

In Spark 1.4-1.6

- Added Binary Storage and Basic Code Generation
- DataFrame + Dataset APIs enable Tungsten in User Programs
- Tungsten also being used under SparkSQL + parts of MLlib

For Spark 2.0 (1H2016 Release)

- Whole-stage Code Generation
 - Remove expensive Iterator calls
 - Fuse across multiple operators
- Optimized Input/Output
 - Parquet + Built-in Cache



Automatically applies to SQL, DataFrames, Datasets

Datasets and DataFrames

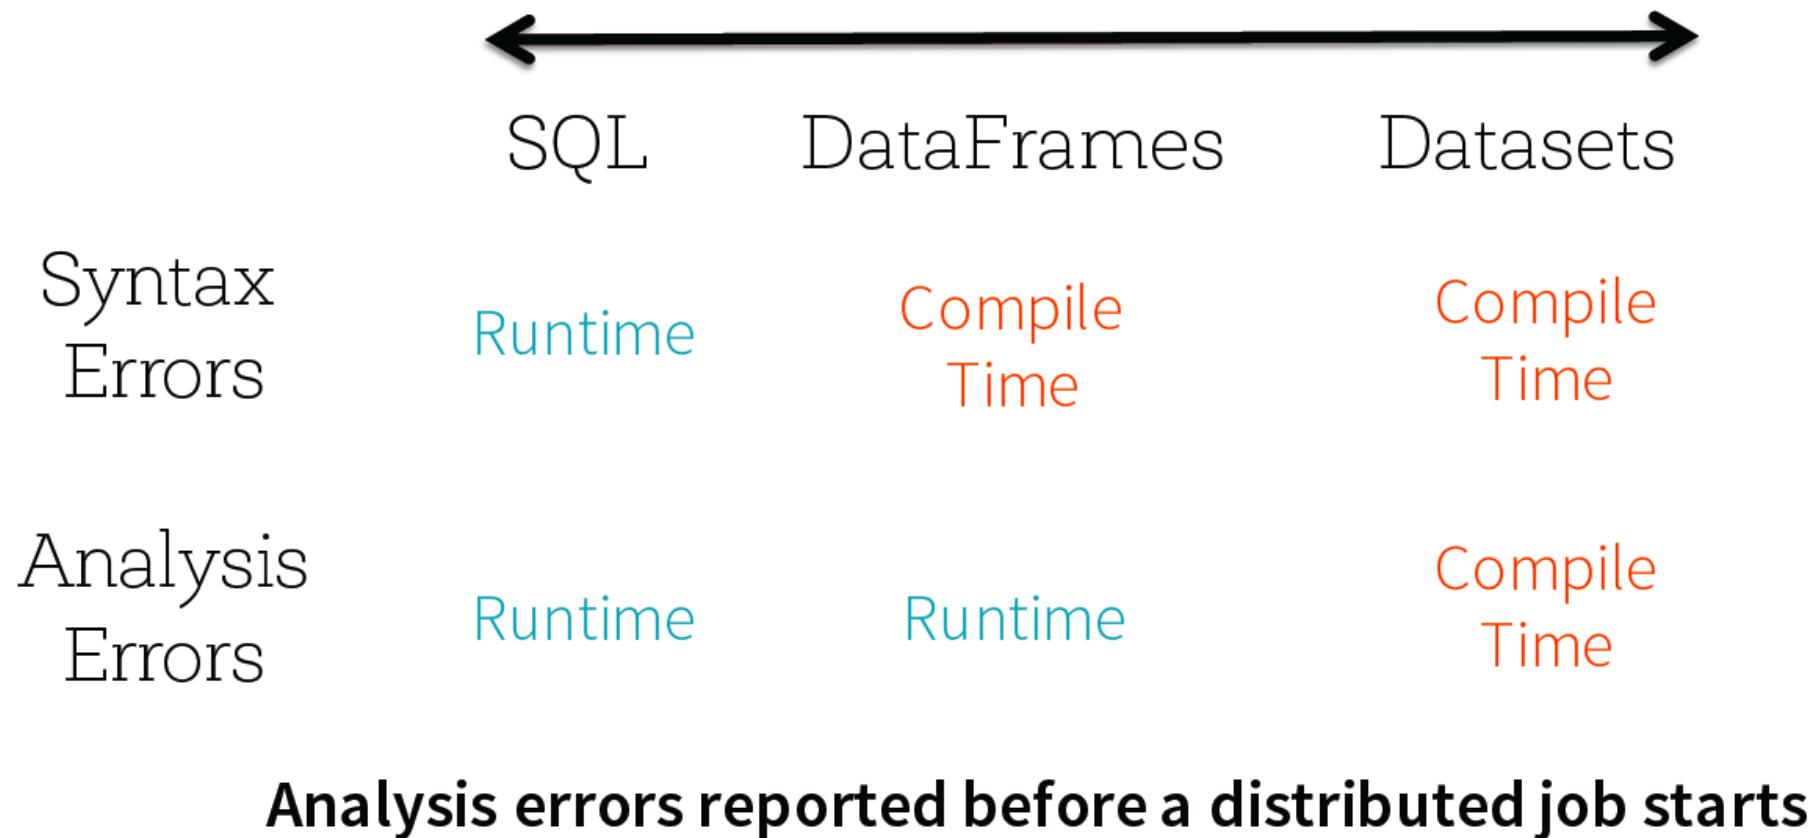
- DataFrames are collections of rows with a schema
- Datasets add static types, e.g. Dataset[Person]
- Both run on Tungsten
- Spark 2.0 will merge these APIs:

Dataframe = Dataset [Row]

Benefits of Merging

- Simpler to understand
 - Only kept Dataset separate to keep binary compatibility in Spark 1.x
- Libraries can take data of both forms
- With Streaming, same API will also work on streams

Datasets: Another new Structured Data Abstraction and its API in Spark



Example for Datasets and DataFrames

```
case class User(name: String, id: Int)
case class Message(user: User, text: String)

dataframe = sqlContext.read.json("log.json")           // DataFrame, i.e. Dataset[Row]
messages = dataframe.as[Message]                      // Dataset[Message]

users = messages.filter(m => m.text.contains("Spark"))
               .map(m => m.user)                         // Dataset[User]

pipeline.train(users)                                // MLlib takes either DataFrames or Datasets
```

Datasets API

- Type-safe: Operate on domain objects with compiled lambda functions

```
val df = ctx.read.json("people.json")

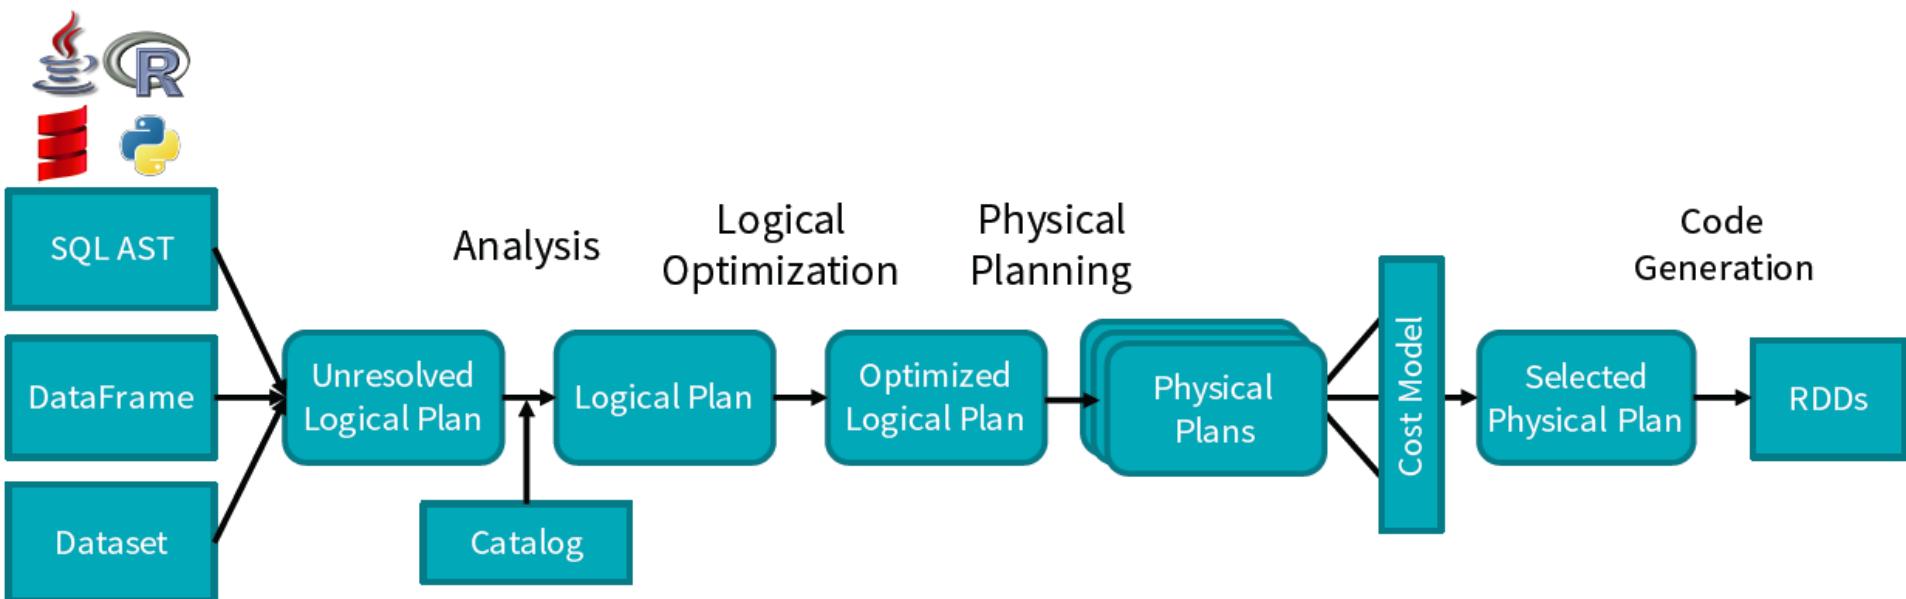
// Convert data to domain objects.
case class Person(name: String, age: Int)
val ds: Dataset[Person] = df.as[Person]
ds.filter(_.age > 30)

// Compute histogram of age by name.
val hist = ds.groupBy(_.name).mapGroups {
  case (name, people: Iter[Person]) =>
    val buckets = new Array[Int](10)
    people.map(_.age).foreach { a =>
      buckets(a / 10) += 1
    }
    (name, buckets)
}
```

Long-Term Direction

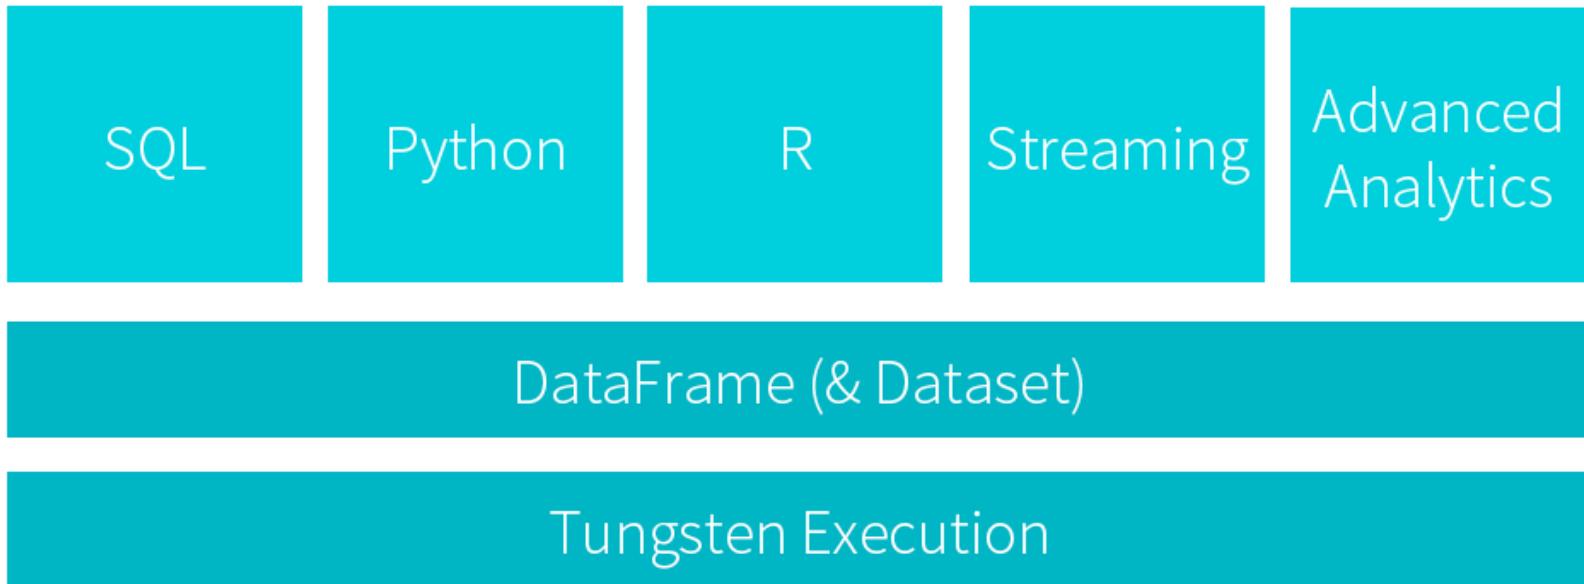
- RDD will remain the low-level API in Spark
- Datasets and DataFrames give richer semantics and optimizations
 - New libraries will increasingly use these as interchange format, e.g. Structured Streaming, MLlib and GraphFrames

Shared Optimization and Execution



DataFrames, Datasets and SQL
share the same optimization/execution pipeline

Execution Architecture for Spark Going Forward (Ver. 2.0 and Beyond)

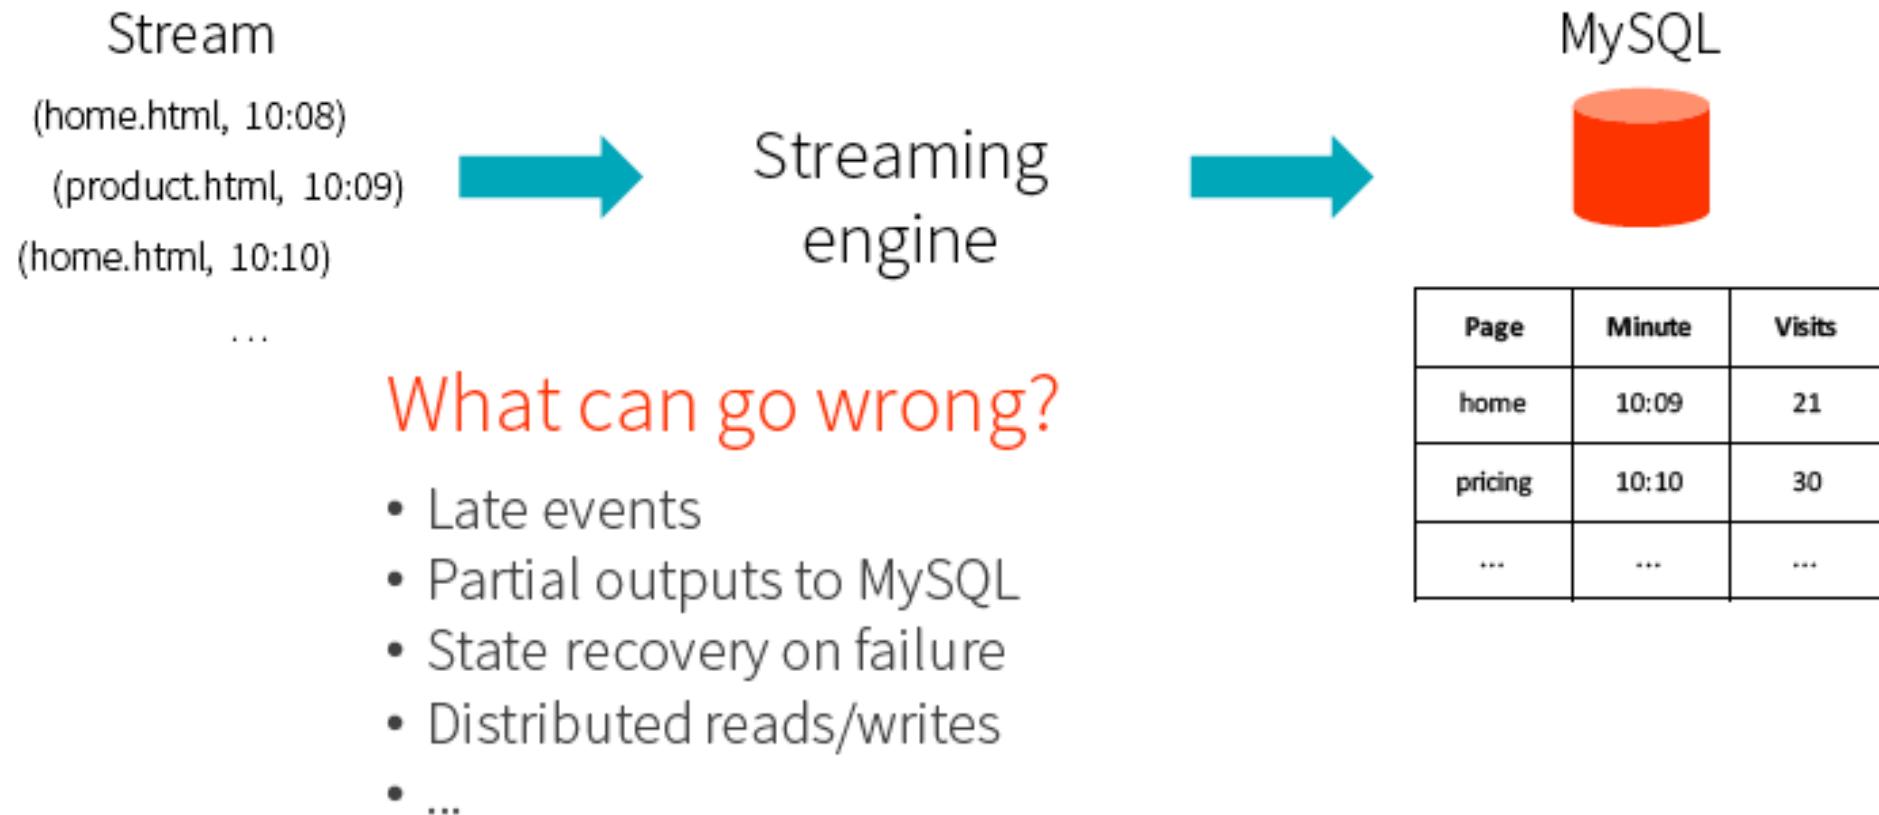


Motivation for Structured Streaming

- Real-time processing is increasingly important
- Most applications need to combine it with Batch & Interactive queries, e.g.
 - Track state using a stream, then run SQL queries
 - Train an Machine Learning model offline, then update it with new, online data

Not just streaming, but
“continuous applications”

Challenges of Integrating Streaming into a Real-world Application Infrastructure



Complex Programming Models

Data

Late arrival, varying distribution over time, ...

Processing

Business logic change & new ops
(windows, sessions)

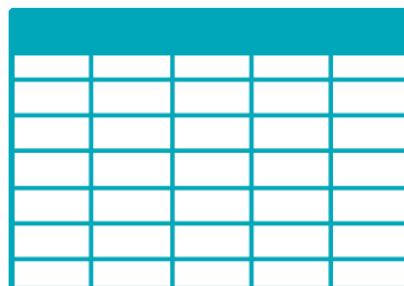
Output

How do we define
output over time & correctness?

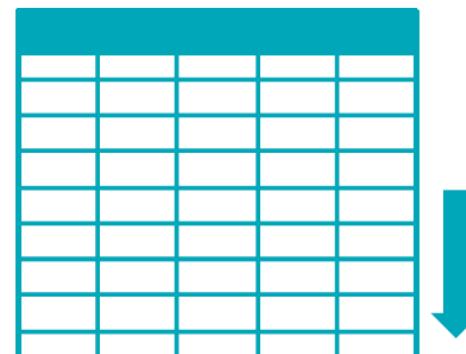
Structured Streaming

- High-level Streaming API built on Spark SQL engine
 - Run the same queries on DataFrames
 - Event-time, windowing, sessions, source and sinks
- Unify Streaming, Interactive and Batch Queries
 - Aggregate data in a stream, then serve using JDBC
 - Change queries at runtime
 - Build and Apply Machine Learning models

Spark 1.3
Static DataFrames



Spark 2.0
Continuous DataFrames



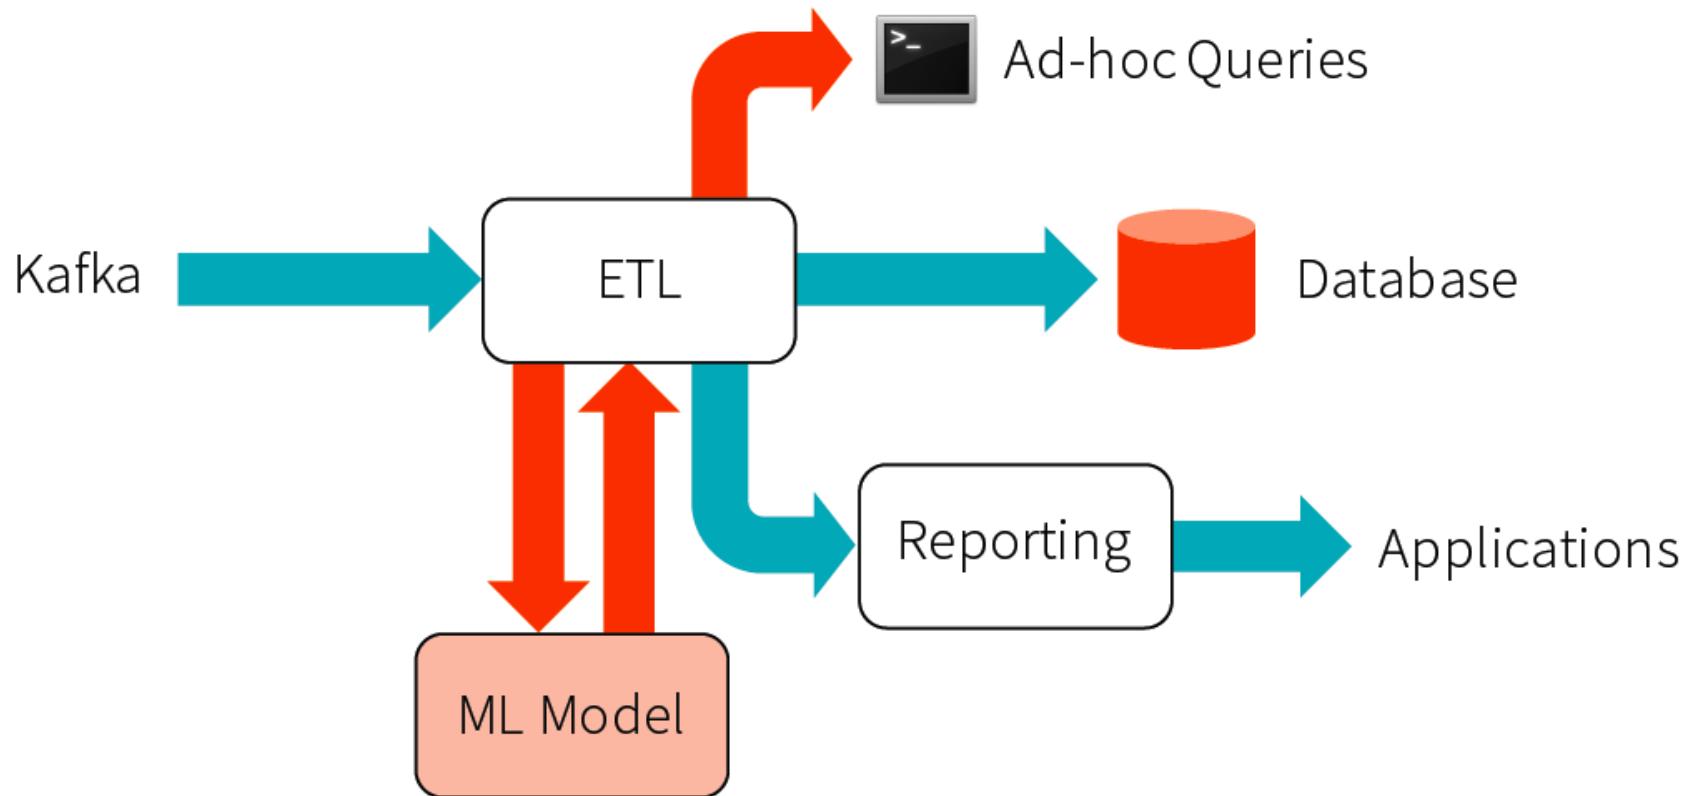
Not just streaming, but
“continuous applications”



Single API !

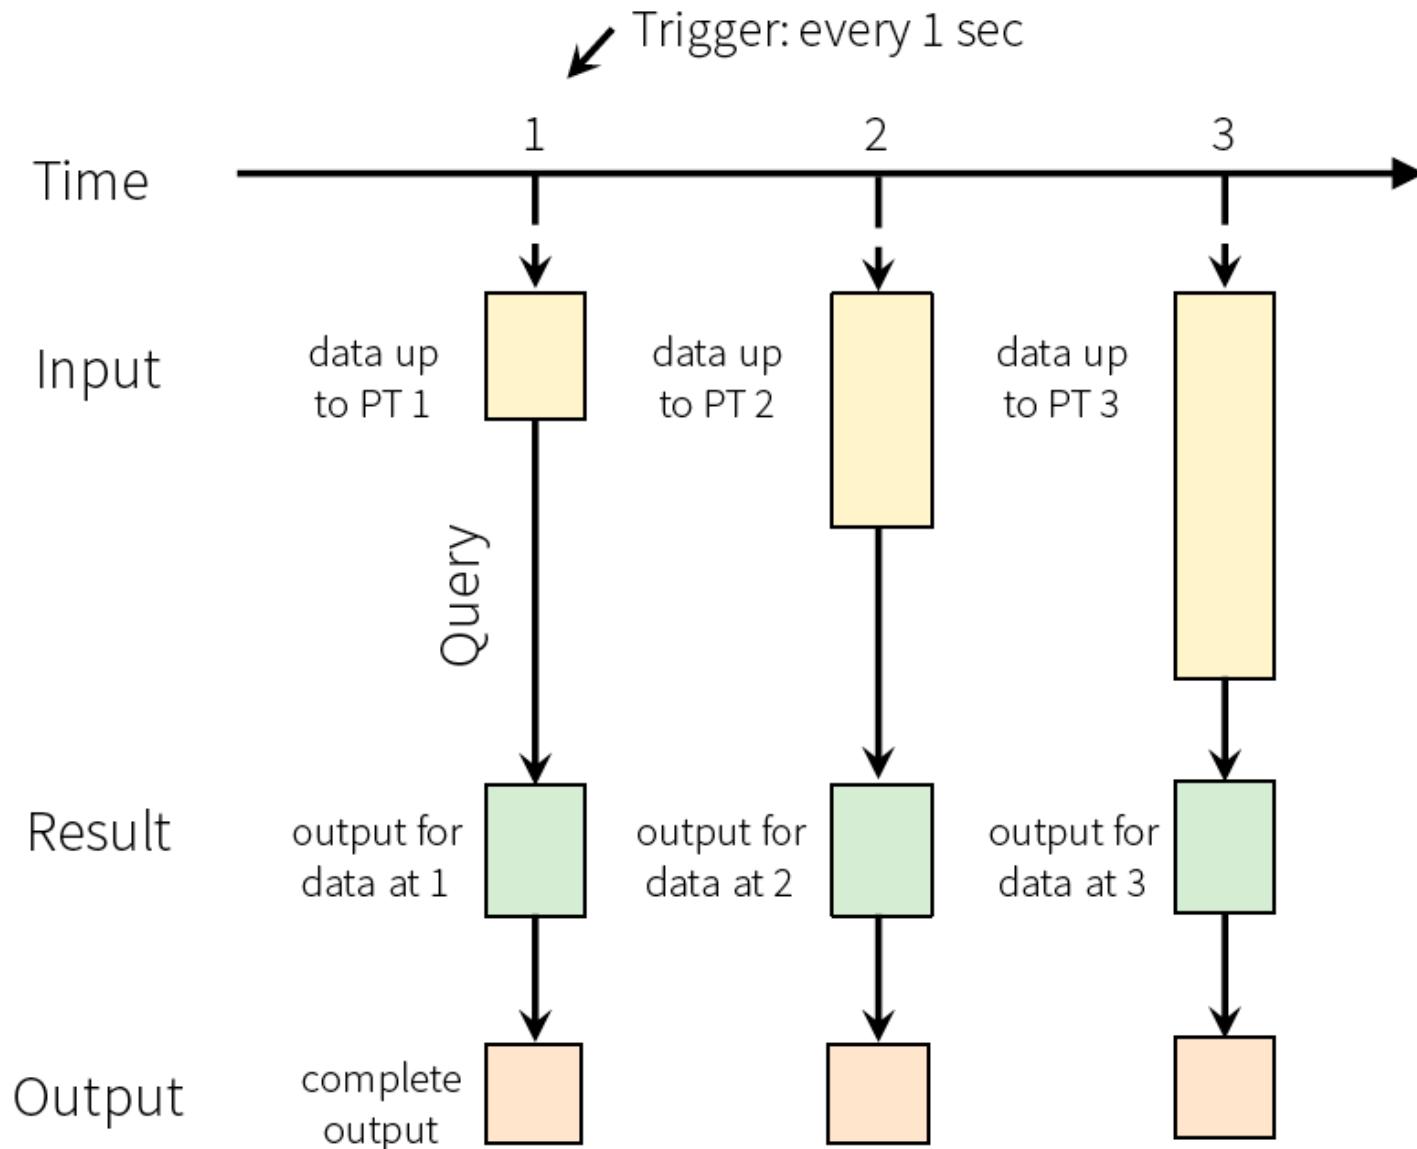
An Example

- Traditional streaming
- Other processing types

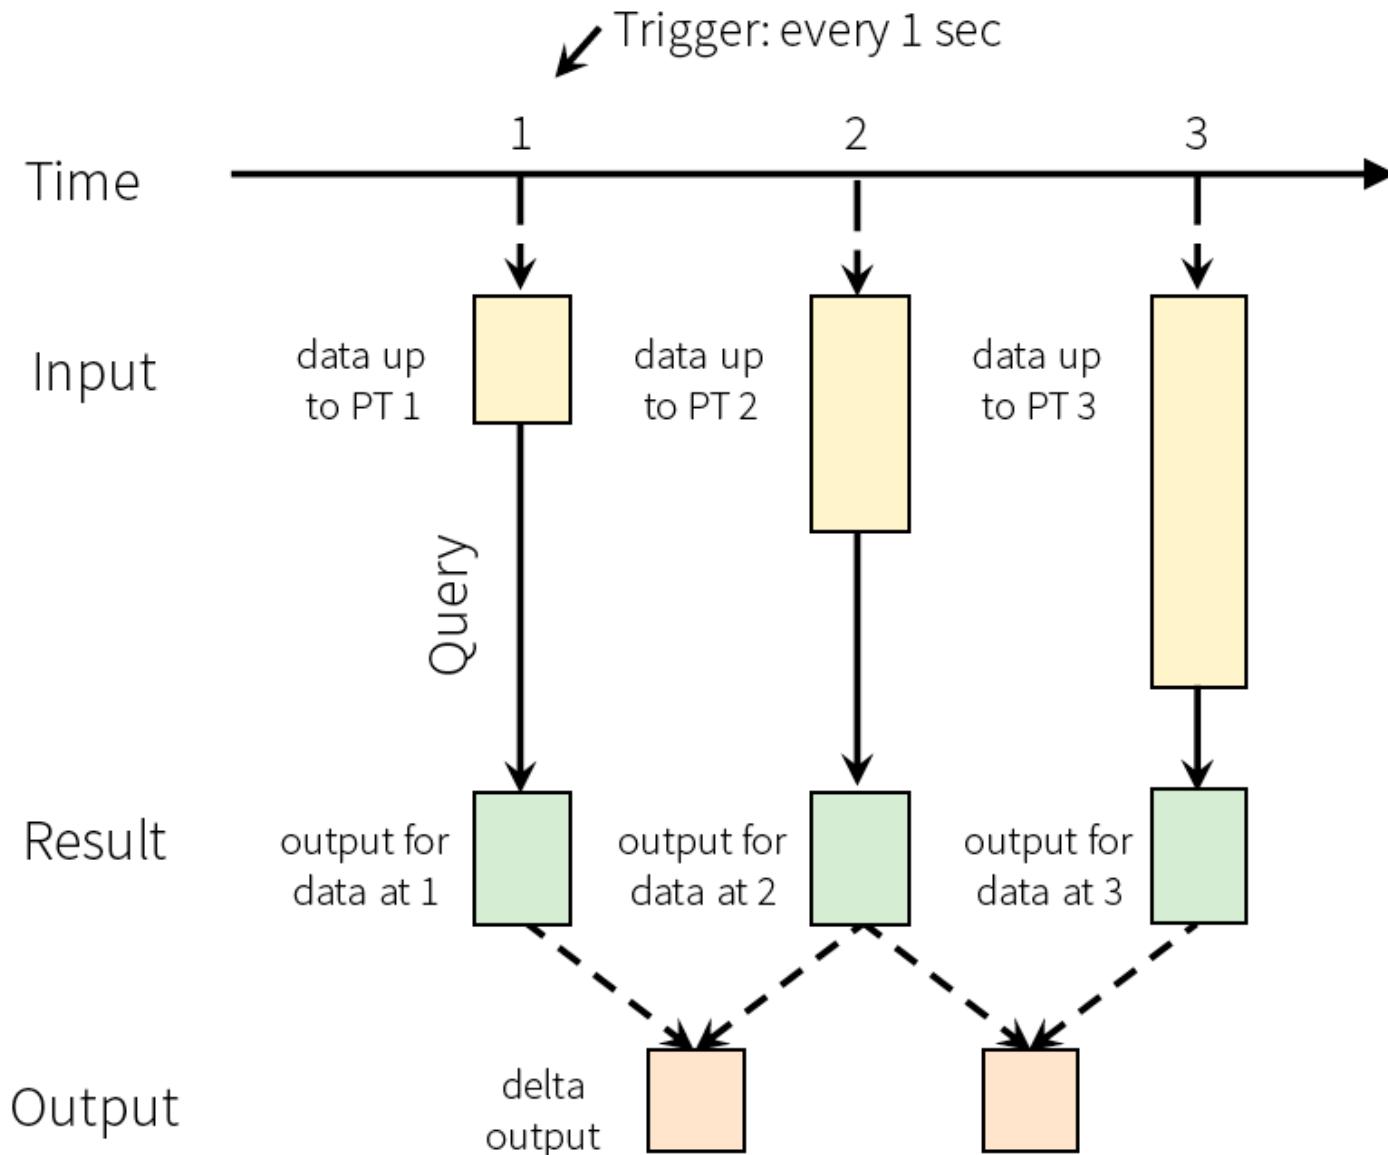


Goal: end-to-end continuous applications

Model for Structured Streaming



Model for Structured Streaming



Model Details for Structured Streaming

- **Input Sources:** Append-Only Tables
- **Queries:** New operators for Windowing, Sessions, etc
- **Triggers:** based on time (e.g. every 1 sec)
- **Output modes:** Complete, Deltas, Update-in-Place

An ETL Example

- **Input:** files in S3
- **Query:** map (transform each record)
- **Trigger:** “every 5 sec”
- **Output mode:** “new records”, into S3 sink

A Page View Count Example

- **Input:** records in Kafka
- **Query:** select count (*) group by page, minute(evtime)
- **Trigger:** “every 5 sec”
- **Output mode:** “update-in-place”, into MySQL sink

NOTE: This will automatically update “old” records on late data

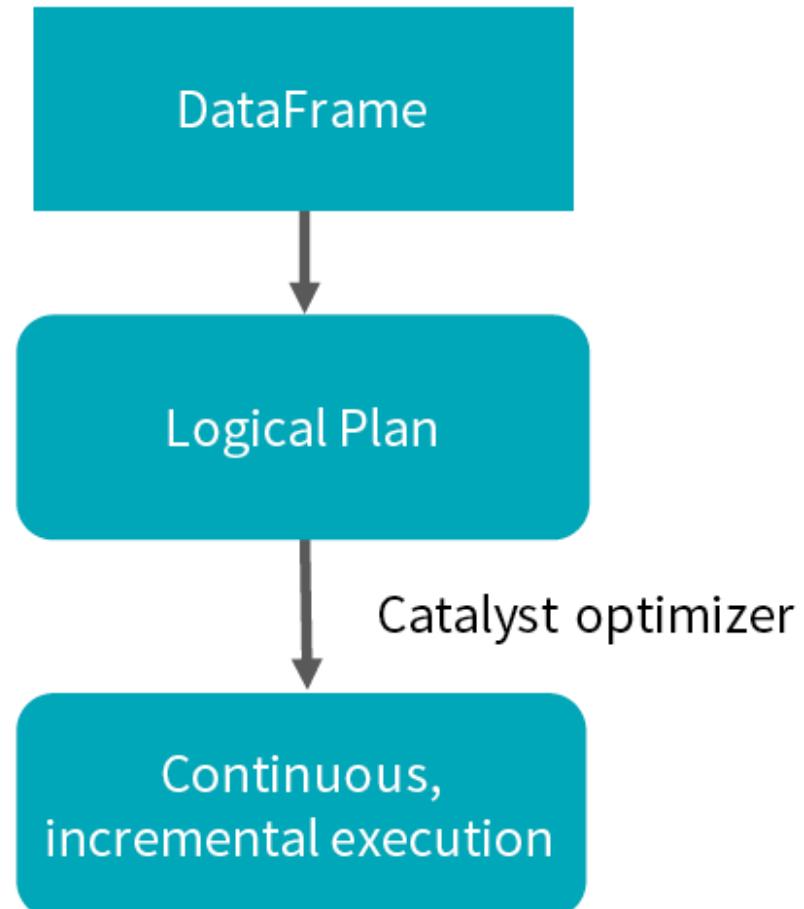
Execution for Structured Streaming

Logically:

- DataFrame operations on static data (i.e. as easy to understand as batch)

Physically:

- Spark automatically runs the query in Streaming fashion (i.e. incrementally and continuously)



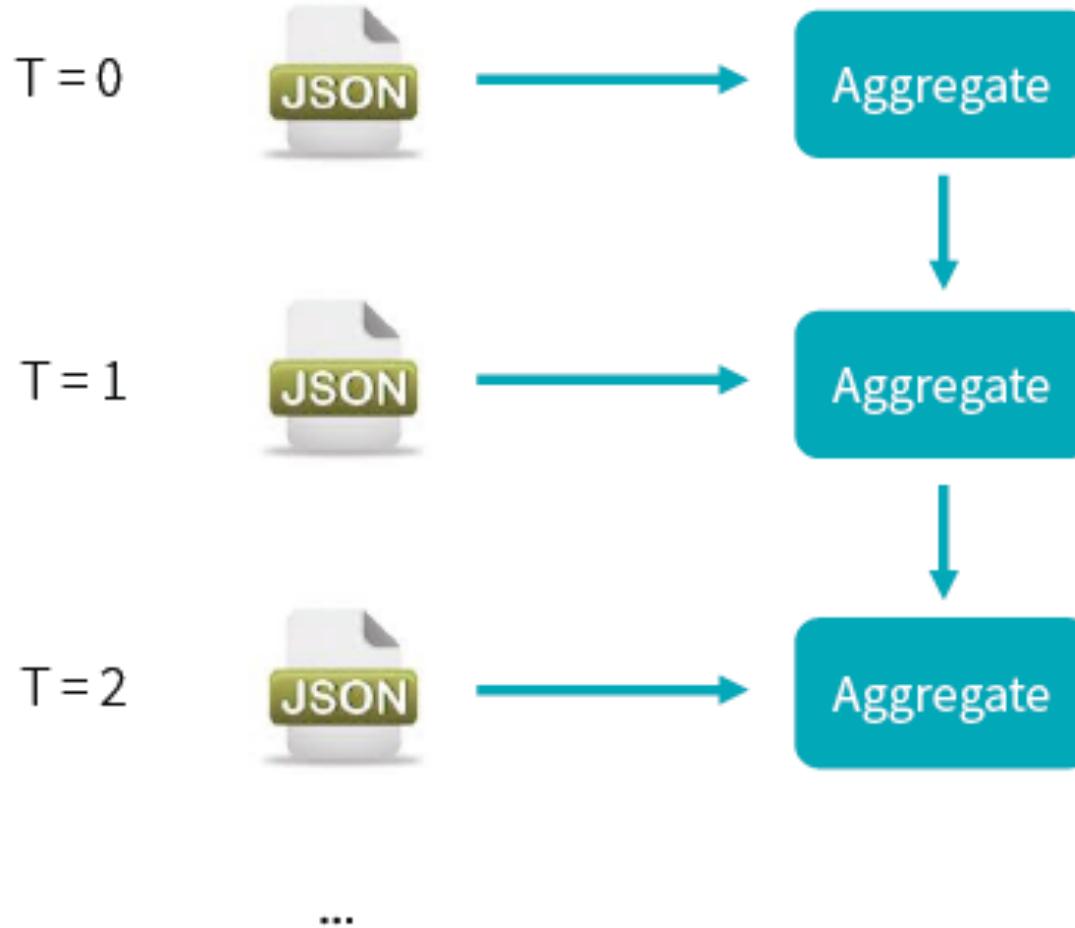
Example: Batch Aggregation

```
logs = ctx.read.format("json").open("s3://logs")  
  
logs.groupBy(logs.user_id).agg(sum(logs.time))  
    .write.format("jdbc")  
    .save("jdbc:mysql//...")
```

Example: Continuous Aggregation

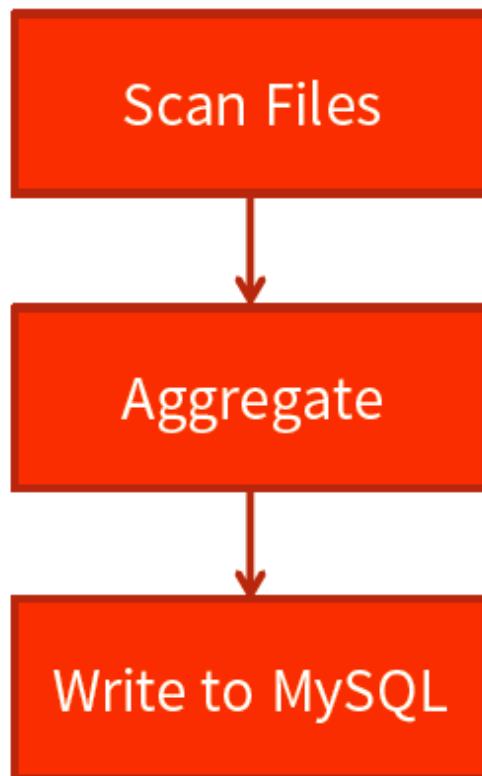
```
logs = ctx.read.format("json").stream("s3://logs")  
  
logs.groupBy(logs.user_id).agg(sum(logs.time))  
.write.format("jdbc")  
.stream("jdbc:mysql//...")
```

Automatic Incremental Execution



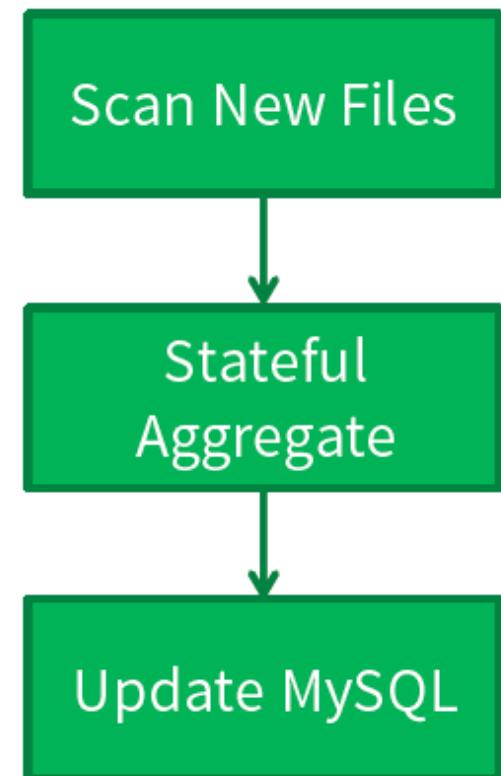
Incrementalized by Spark

Batch



Transformation
requires
information
about the
structure

Continuous



Rest of Spark will follow

- Interactive queries should just work
- Spark's data source API will be updated to support seamless streaming integration
 - Exactly once semantics **end-to-end**
 - Different Output modes (complete, delta, update-in-place)
- Machine Learning algorithms will be updated according to this new model

Going Forward: What to expect from Spark ?

- Spark 2.0
 - Unification of the APIs
 - Basic Streaming API
 - Event-time Aggregations
- Spark 2.1+
 - Other Streaming sources/sinks
 - Machine Learning
- Structure in other libraries, e.g.
 - MLlib
 - GraphFrames